



CHULA ENGINEERING
Foundation toward Innovation

COMPUTER



Introduction to Deep Learning (CNN)

Prof. Peerapon Vateekul, Ph.D.

Department of Computer Engineering,
Faculty of Engineering, Chulalongkorn University
Peerapon.v@chula.ac.th

www.cp.eng.chula.ac.th/~peerapon/



Outline

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN) [Imaging Task]
- Image Classification
- Object Detection
- Semantic Segmentation



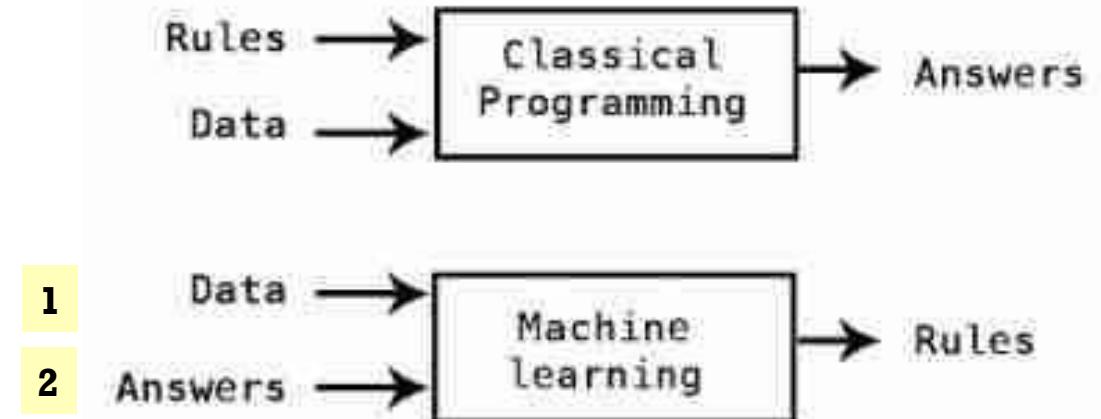
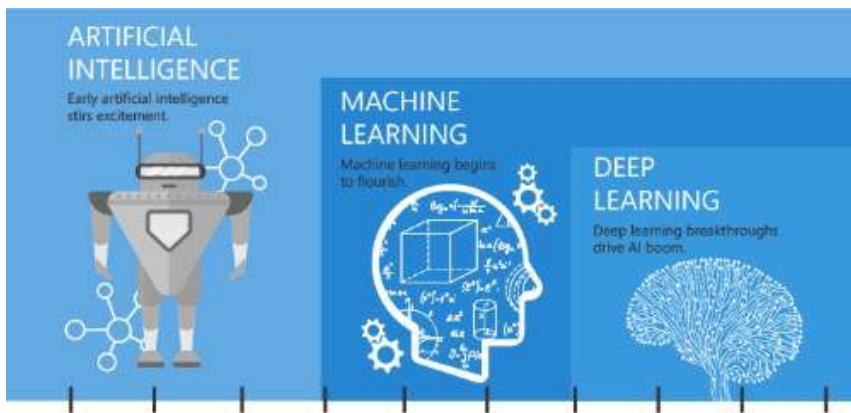


Introduction to Deep Learning



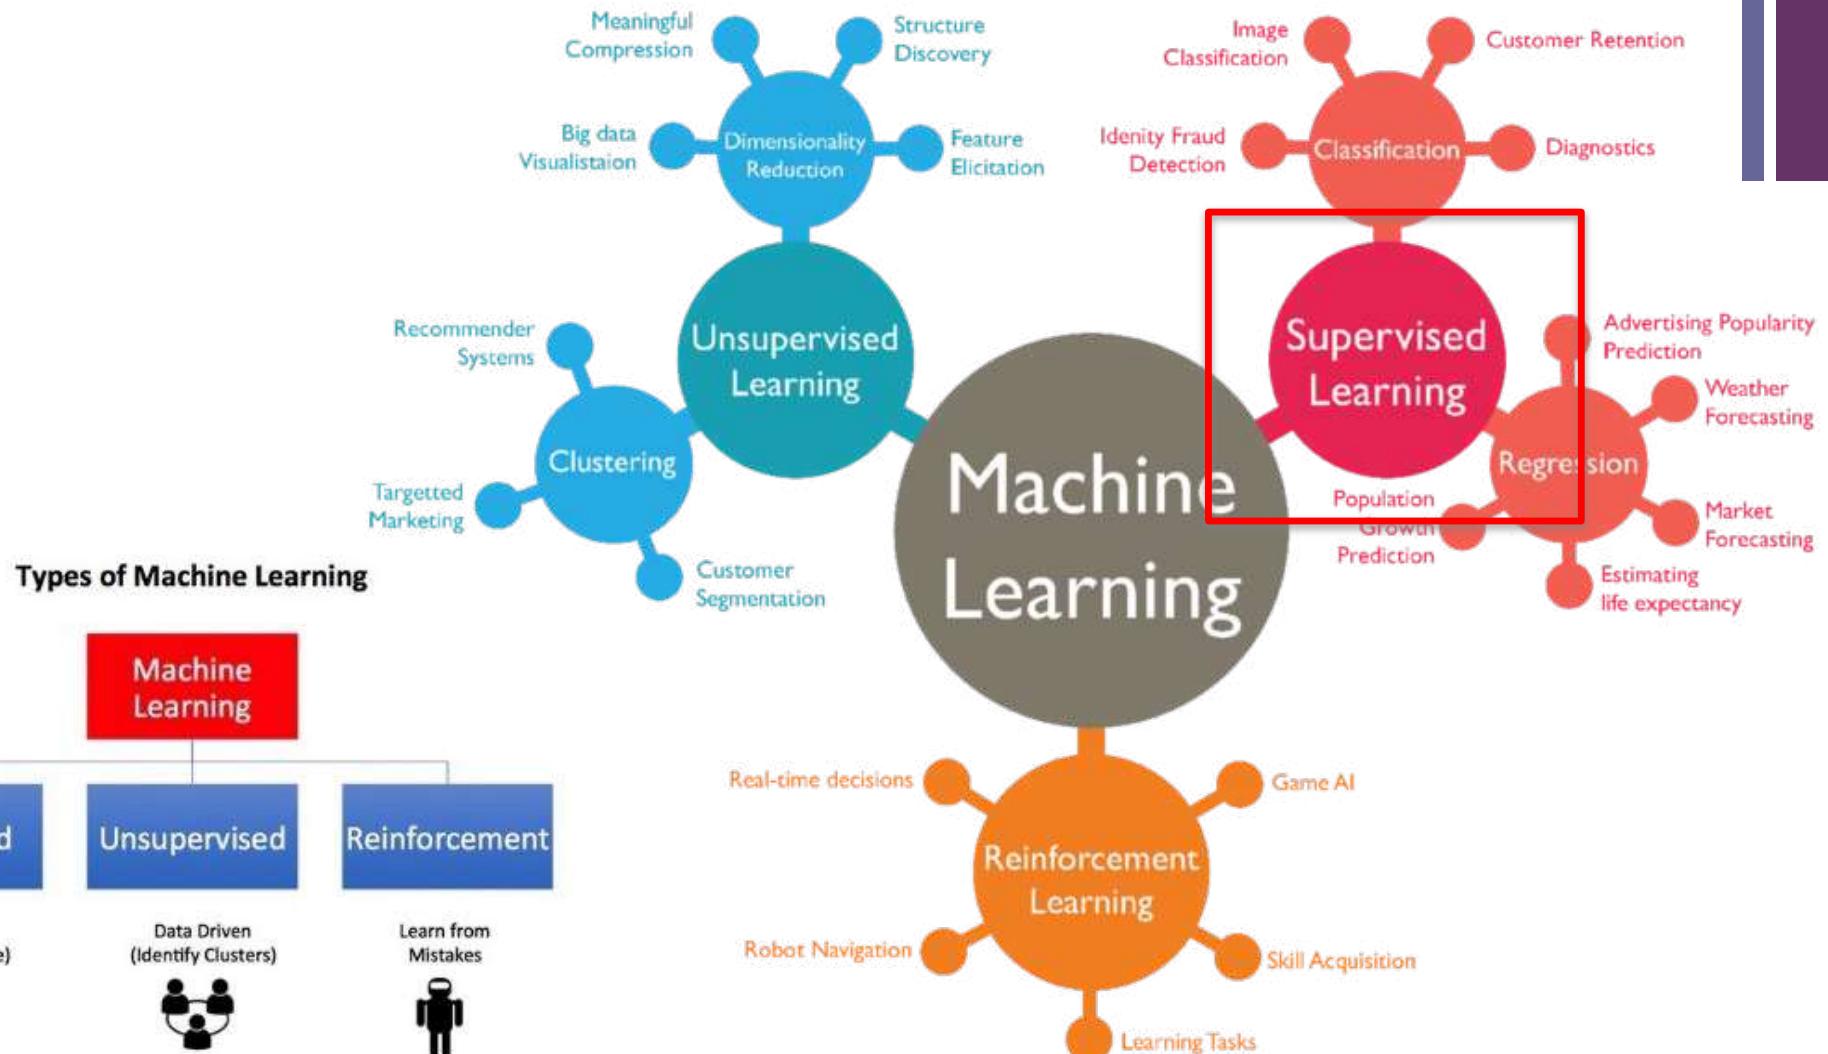
AI = Automation

- 1) Rule-based AI (Symbolic AI)
- 2) Machine Learning



<https://mc.ai/machine-learning-basics-artificial-intelligence-machine-learning-and-deep-learning/>

+ Machine Learning





Task1: Supervised learning

Handcrafted features

Training Data



inputs					target
Age	Gender	BodyTemp	Cough	Corona	
12	Female	37	Yes	Yes	
35	Female	39	No	Yes	
32	Male	38	Yes	No	

Testing Data



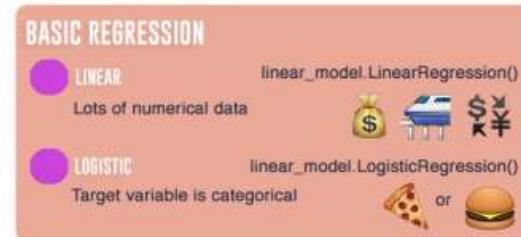
Age	Gender	BodyTemp	Cough	Corona
25	Male	40	No	?

Application: Corona Prediction



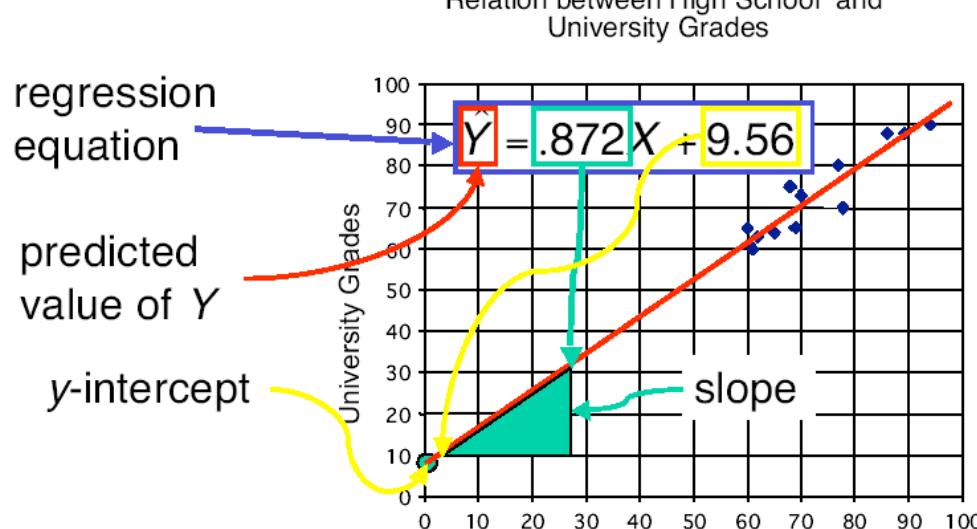
Prediction algorithms

- Decision Tree
- (Logistic) Regression
- kNN
- Support Vector Machine
- Neural Networks (NN)
- Deep Learning





Regression – Linear Relationship



$$\hat{y} = \hat{w}_0 + \hat{w}_1 x_1 + \hat{w}_2 x_2$$

target intercept input

- The least square method aims to minimize the following term

$$\sum_{\text{training data}} (y_i - \hat{y}_i)^2$$

+

Logistic Regression (cont.)

Linear Relationship

Training Data



inputs					target
Age	Gender	BodyTemp	Cough	Corona	
12	Female (0)	37	Yes (1)	Yes	
35	Female (0)	39	No (0)	Yes	
32	Male (1)	38	Yes (1)	No	

$$\text{Logit_score} = w_0 + w_1 * \text{Age} + w_2 * \text{Gender} + w_3 * \text{Temp} + w_4 * \text{Cough}$$

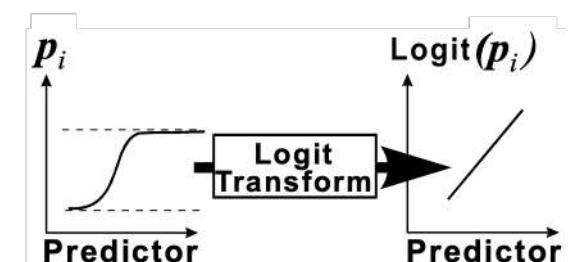
$$\text{Logit_score} = 0.01 - 0.3 * \text{Age} + 0.2 * \text{Gender} + 0.2 * \text{Temp} + 0.9 * \text{Cough}$$

Example

$$\text{Logit_score} = 0.01 - 0.3 * 12 + 0.2 * 0 + 0.2 * 37 + 0.9 * 1 = 4.71$$

prob = 0.9911 → "Yes"

Application: Corona Prediction

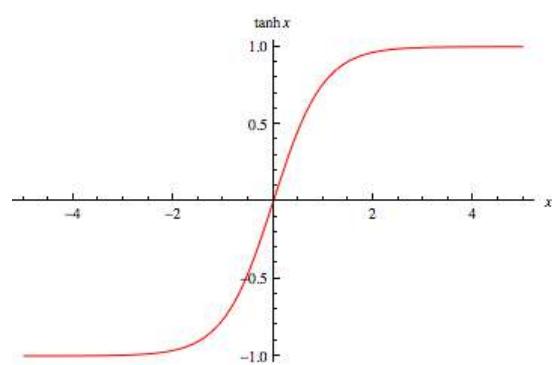
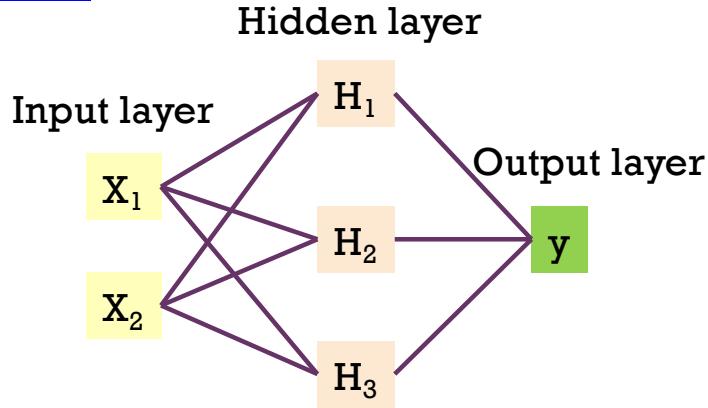


$$\hat{p} = \frac{1}{1 + e^{-\text{logit}(\hat{p})}}$$



Neural Networks (universal approximator)

Non-linear relationship



$$\log\left(\frac{\hat{p}}{1-\hat{p}}\right) = \hat{w}_0 + \hat{w}_1 H_1 + \hat{w}_2 H_2 + \hat{w}_3 H_3$$

$$H_1 = \tanh(\hat{w}_{10} + \hat{w}_{11} x_1 + \hat{w}_{12} x_2)$$

$$H_2 = \tanh(\hat{w}_{20} + \hat{w}_{21} x_1 + \hat{w}_{22} x_2)$$

$$H_3 = \tanh(\hat{w}_{30} + \hat{w}_{31} x_1 + \hat{w}_{32} x_2)$$

Stop when?

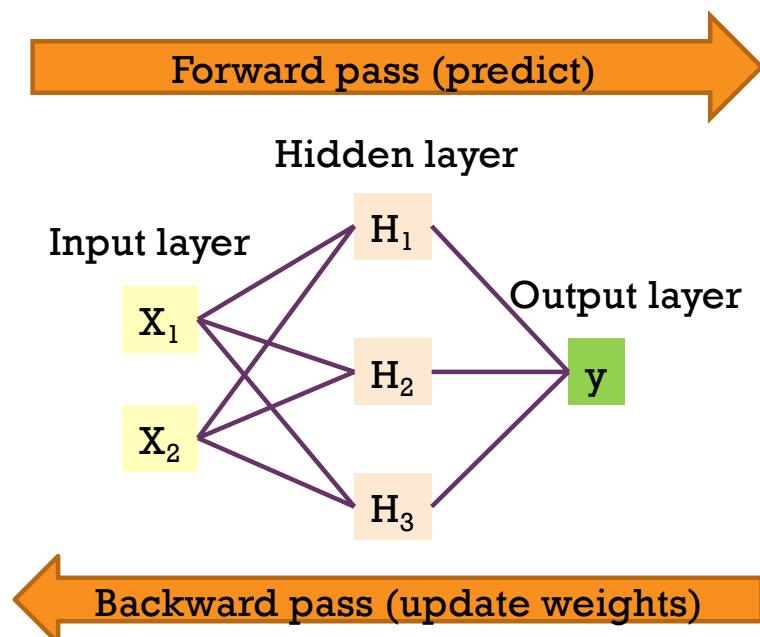
- Converge (no change in loss)
- Max epochs

Important Params:

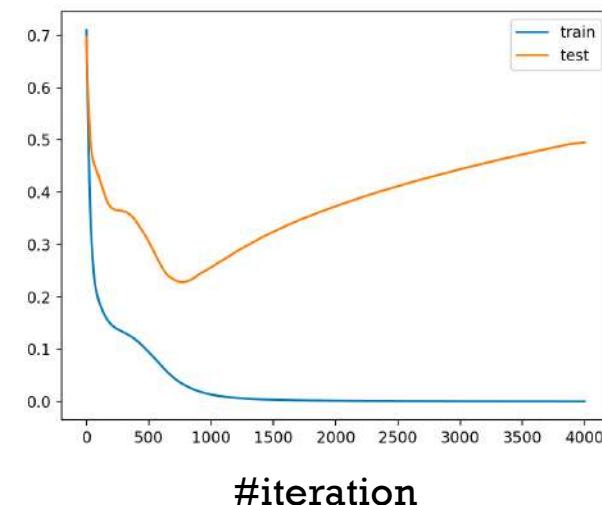
- #hidden units, #hidden layers
- Learning rate, Momentum, decay
- Seed number, etc.



Neural Networks (cont.): Training Non-linear relationship



Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes
35	50,000	Female	Nontaburi	Yes
32	35,000	Male	Bangkok	No



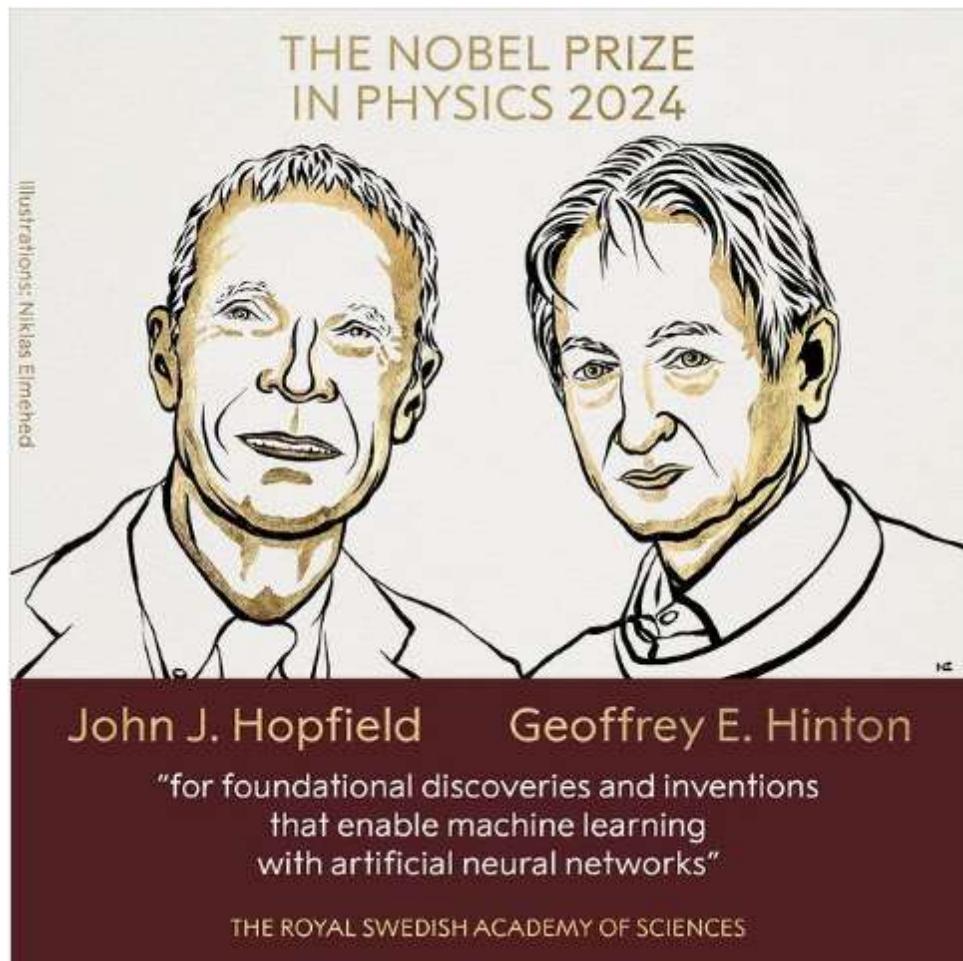
$$\log\left(\frac{\hat{p}}{1-\hat{p}}\right) = \hat{w}_0 + \hat{w}_1 H_1 + \hat{w}_2 H_2 + \hat{w}_3 H_3$$

$$H_1 = \tanh(\hat{w}_{10} + \hat{w}_{11} x_1 + \hat{w}_{12} x_2)$$

$$H_2 = \tanh(\hat{w}_{20} + \hat{w}_{21} x_1 + \hat{w}_{22} x_2)$$

$$H_3 = \tanh(\hat{w}_{30} + \hat{w}_{31} x_1 + \hat{w}_{32} x_2)$$

Nobel Prize to Backpropagation's Inventors



8 October 2024

The Royal Swedish Academy of Sciences has decided to award the Nobel Prize in Physics 2024 to

John J. Hopfield
Princeton University, NJ, USA

Geoffrey E. Hinton
University of Toronto, Canada

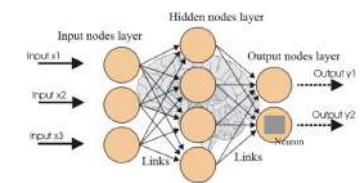
"for foundational discoveries and inventions that enable machine learning with artificial neural networks"

They trained artificial neural networks using physics

John Hopfield invented a network that uses a method for saving and recreating patterns. We can imagine the nodes as pixels.

The *Hopfield network* utilises physics that describes a material's characteristics due to its atomic spin – a property that makes each atom a tiny magnet.

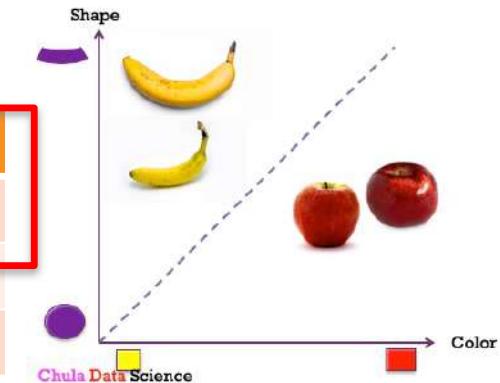
Geoffrey Hinton used the Hopfield network as the foundation for a new network that uses a different method: the *Boltzmann machine*. This can learn to recognise characteristic elements in a given type of data.





Handcrafted features

Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes
35	50,000	Female	Nontaburi	Yes
32	35,000	Male	Bangkok	No



13

Can we still tell the features (columns)?



shutterstock.com - 451802557



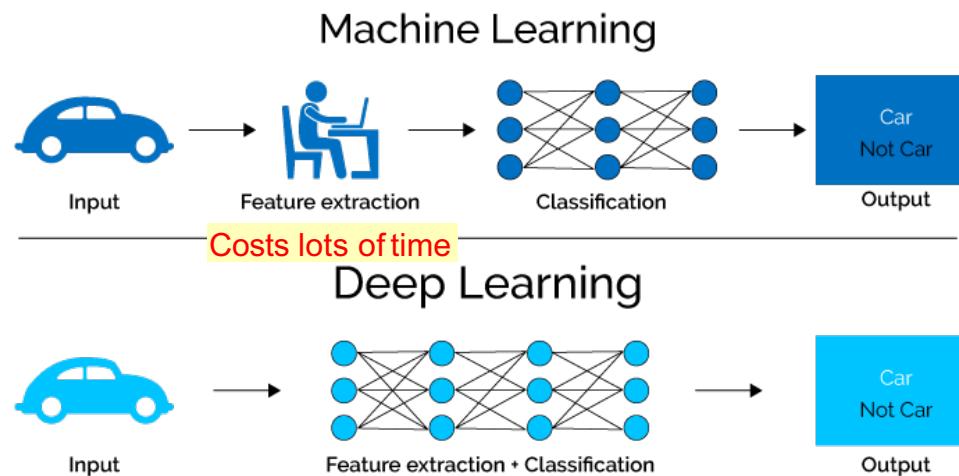
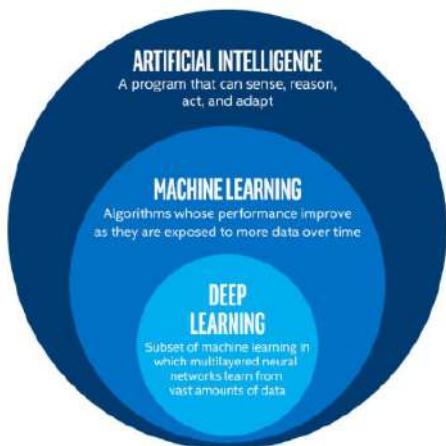
What is Deep Learning (DL)?



Part of the machine learning field of learning representations of data. Exceptional effective at learning patterns.



Utilizes learning algorithms that derive meaning out of data by using a hierarchy of multiple layers that mimic the neural networks of our brain.

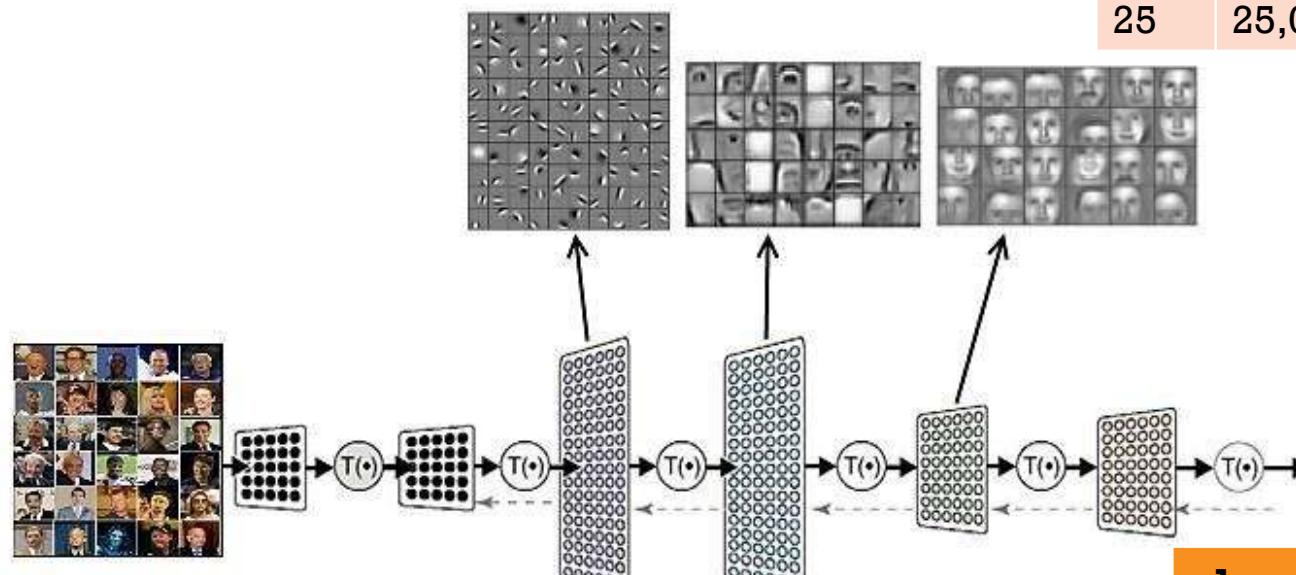




Deep Learning – Basics (cont.)

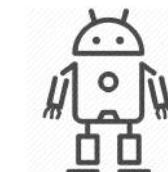
What did it learn?

A deep neural network consists of a **hierarchy of layers**, whereby each layer **transforms the input data** into more abstract representations (e.g., edge -> nose -> face). The output layer combines those features to make predictions.



Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes

x1	x2	x3	x4	Corona
0.7	0.2	-0.5	-0.1	Yes



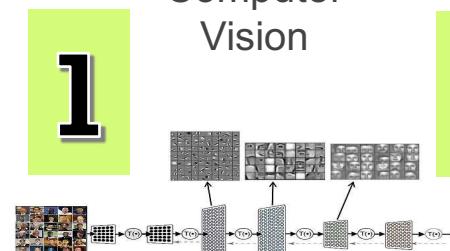
Deep Learning Application



Speech
Recognition



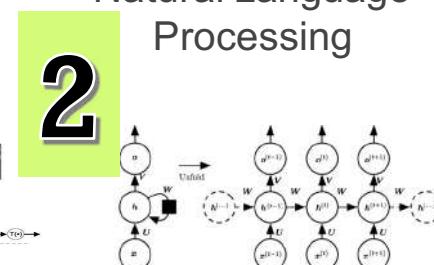
Computer
Vision



CNN

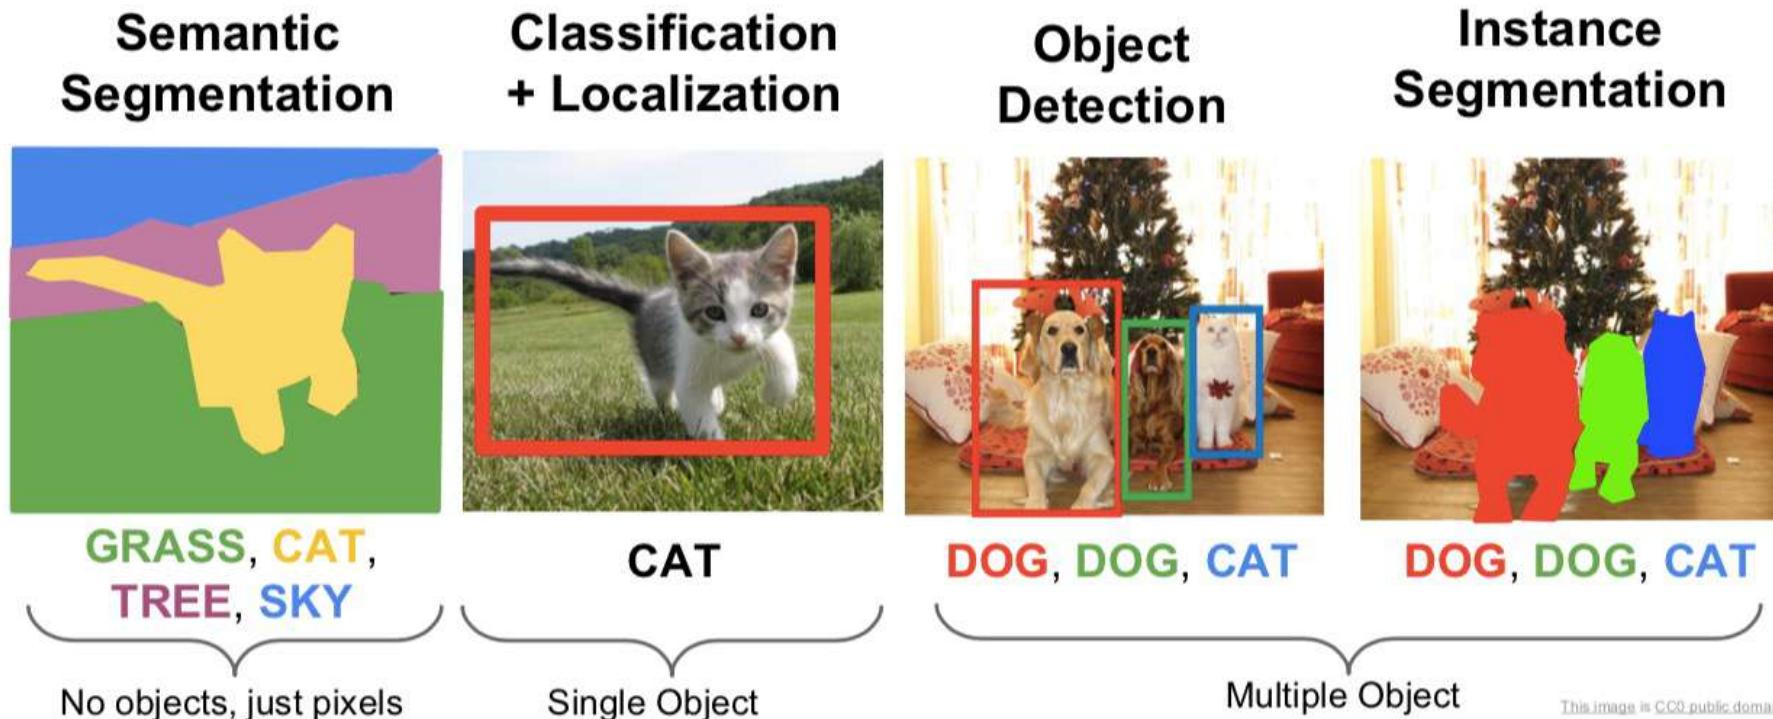


Natural Language
Processing



RNN

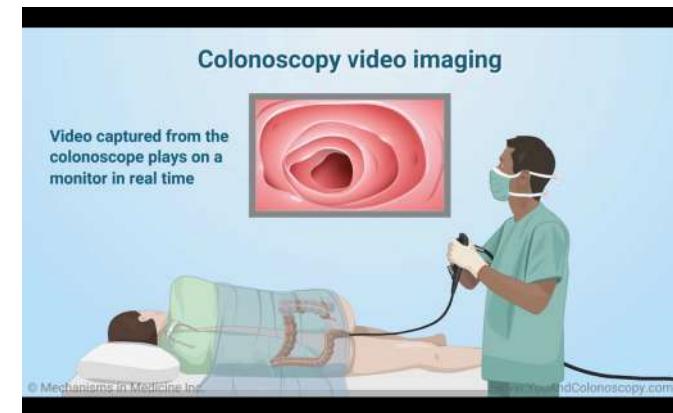
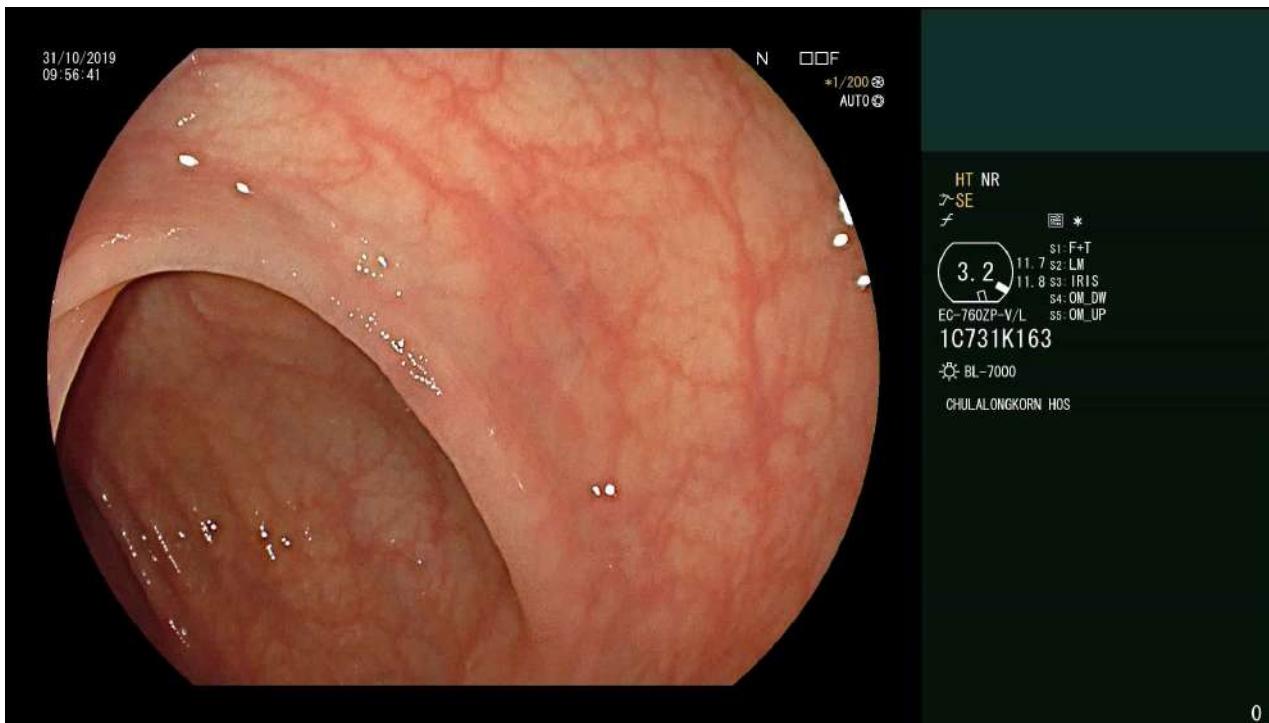
Type of image tasks



Face Recognition



Smart Medical Diagnosis



+

Convolutional Neural Networks (CNN)



Outline

- Introduction
- Basic Image Processing
- Overview of CNN
- Image Processing Tasks with SOTA
- Case Study in Polyp Detection



www.image-net.org

The **ImageNet** project is a large visual database designed for use in visual object recognition software research. Over **14M** URLs of images have been hand-annotated by ImageNet to indicate what objects are pictured on **22K** categories.

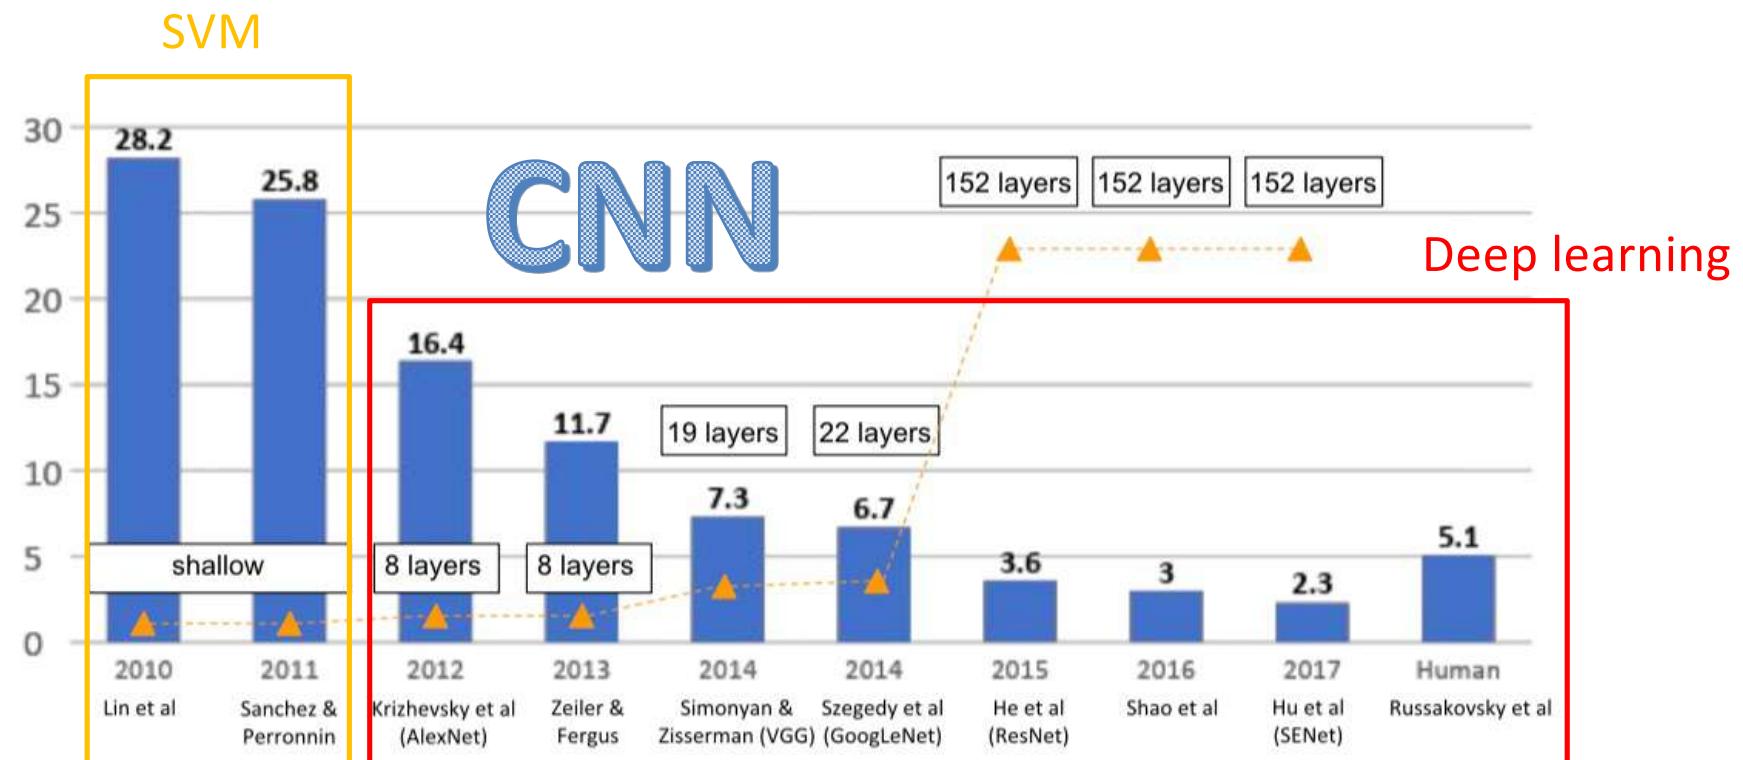


The Image Classification Challenge:
1,000 object classes
1,431,167 images
ImageNet 2017 is the last challenge.

ILSVRC

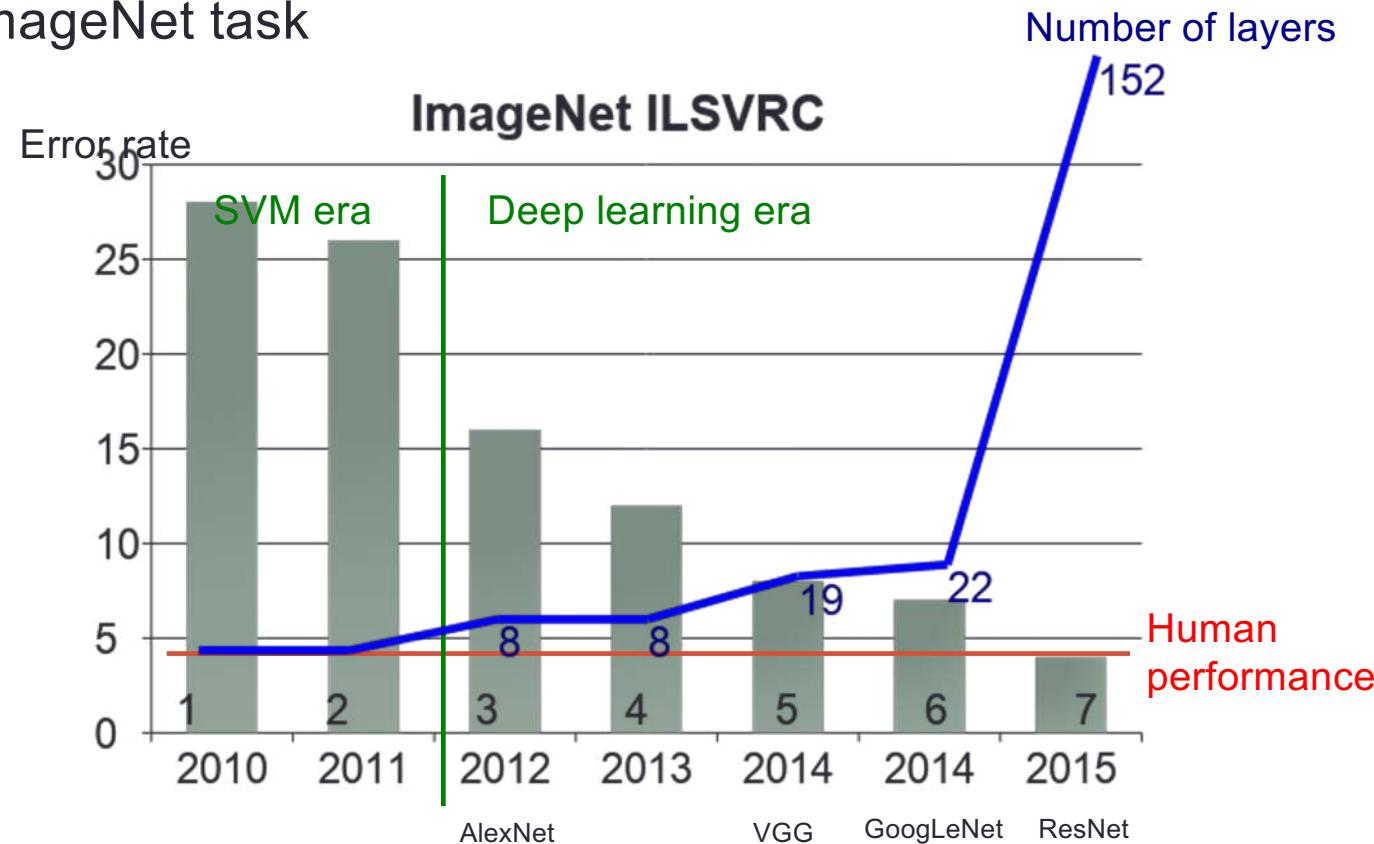


ImageNet Large Scale Visual Recognition Challenge (ILSVRC winners): From SVM to Deep Learning



Wider and deeper networks (Beyond Human)

- ImageNet task

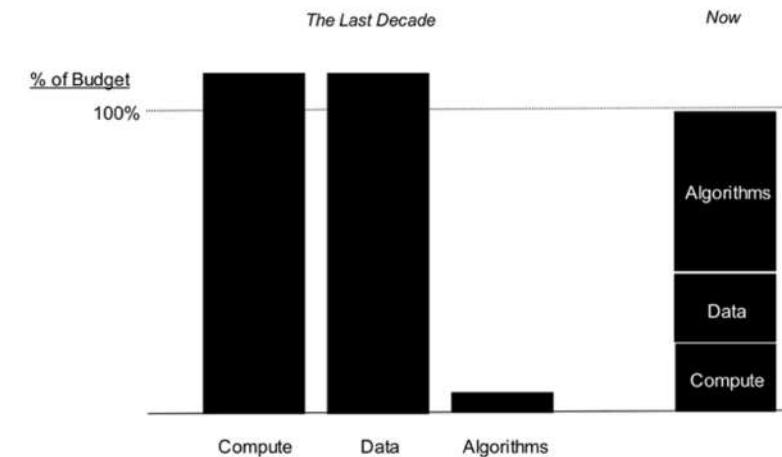


Based on the slide of Aj.Ekapol

Net Large Scale Visual Recognition Challenge, 2014 <https://arxiv.org/abs/1409.0575>

Why now

- Neural Networks has been around since 1990s
- **Big data** – DNN can take advantage of large amounts of data better than other models
- **GPU** – Enable training bigger models possible
- **Deep** – Easier to avoid bad local minima when the model is large



Based on the slide of Aj.Ekapol [/www.kdnuggets.com/2017/06/practical-guide-machine-learning-understand-differentiate-apply.html](http://www.kdnuggets.com/2017/06/practical-guide-machine-learning-understand-differentiate-apply.html)

Application 1: Basic Image Processing

<https://softwaredevelopmentperestroika.wordpress.com/2014/02/11/image-processing-with-python-numpy-scipy-image-convolution/>

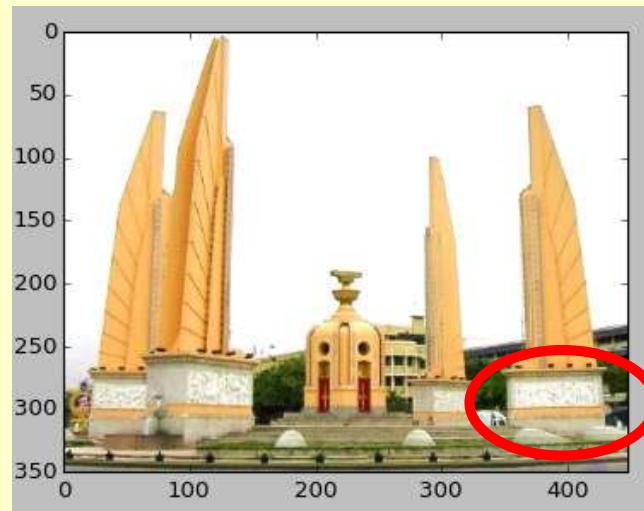
การแสดงรูปภาพใน python ด้วย matplotlib

```
import matplotlib.pyplot as plt  
import matplotlib.image as mpimg
```

เพื่อความง่ายโปรแกรมที่จะเขียน
จากนี้ไปใช้กับไฟล์ png เท่านั้น

```
image = mpimg.imread('monument.png')  
plt.imshow(image)  
plt.show()
```

image เป็นอาร์เรย์ที่แต่ละช่อง
มีค่า 0 ถึง 1 แทนความเข้มของสี



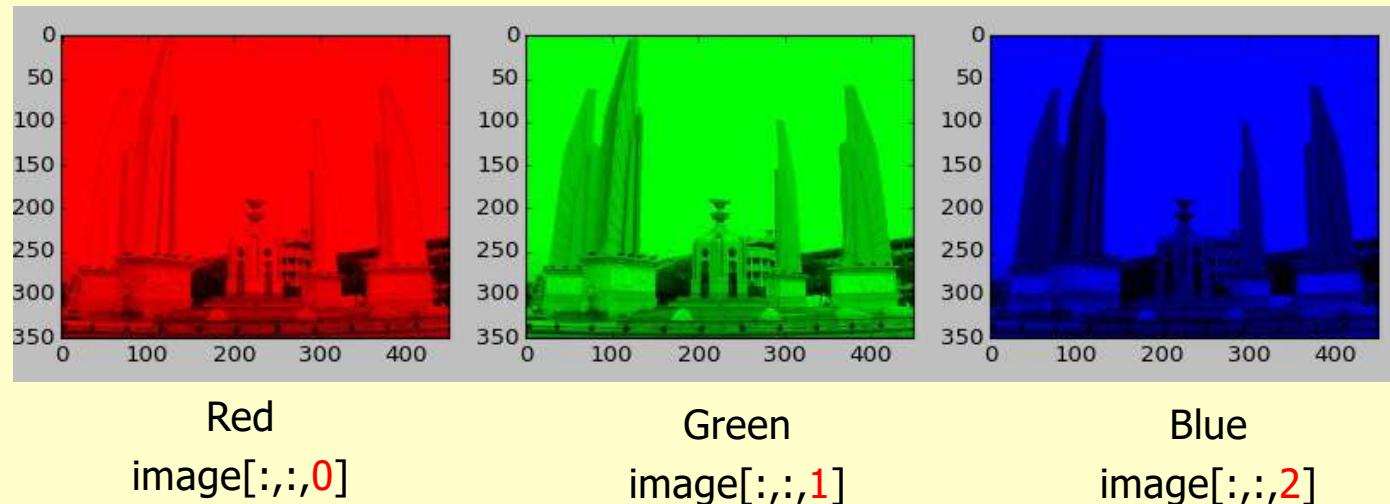
1 pixel = 1 จุดสี มี 3 สีอยู่

0.98762 | 0.72549 | 0.35294



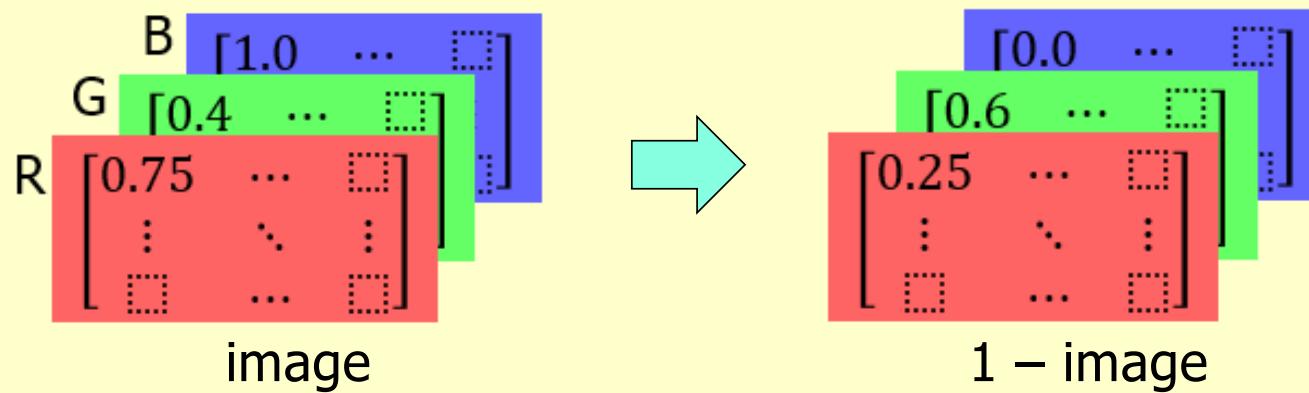
It is an array of R,G,B (3 channels)

```
>>> image.shape  
(350, 449, 3)
```



Basic Image Processing

- `image = mpimg.imread('monument.png')`
- `image.shape → (350, 449, 3)`
- `image = 1 - image` (broadcast & element-wise op.)



How's it changed?

Basic Image Processing (1): Negative Image

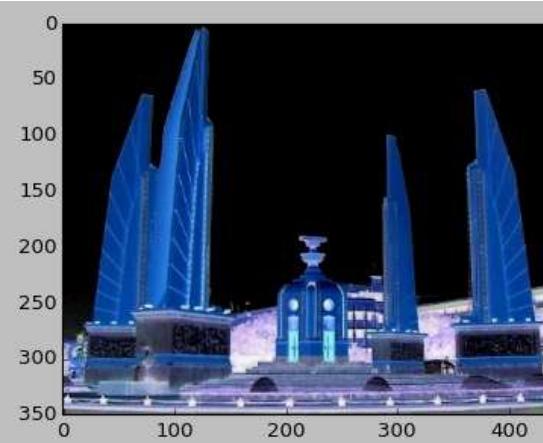
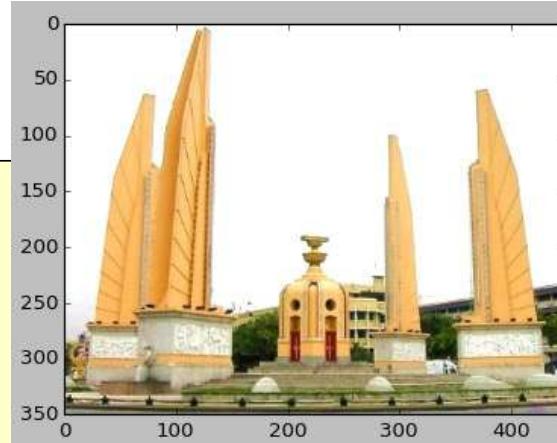
```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

image = mpimg.imread('monument.png')
plt.subplot(1, 2, 1)
plt.imshow(image)

negative = 1 - image
plt.subplot(1, 2, 2)
plt.imshow(negative)

plt.show()
```

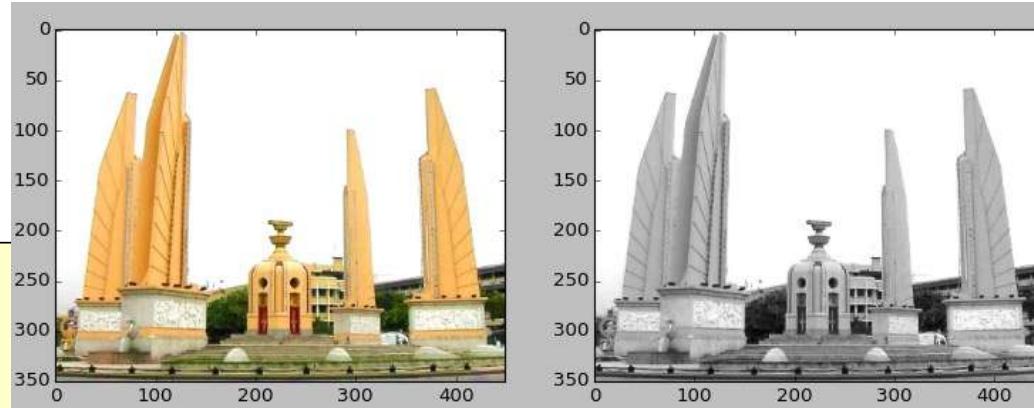
$1 - \text{image}$ = Invert Tone
(White \rightarrow Black, Yellow \rightarrow Blue, ...)
broadcast & element-wise subtract



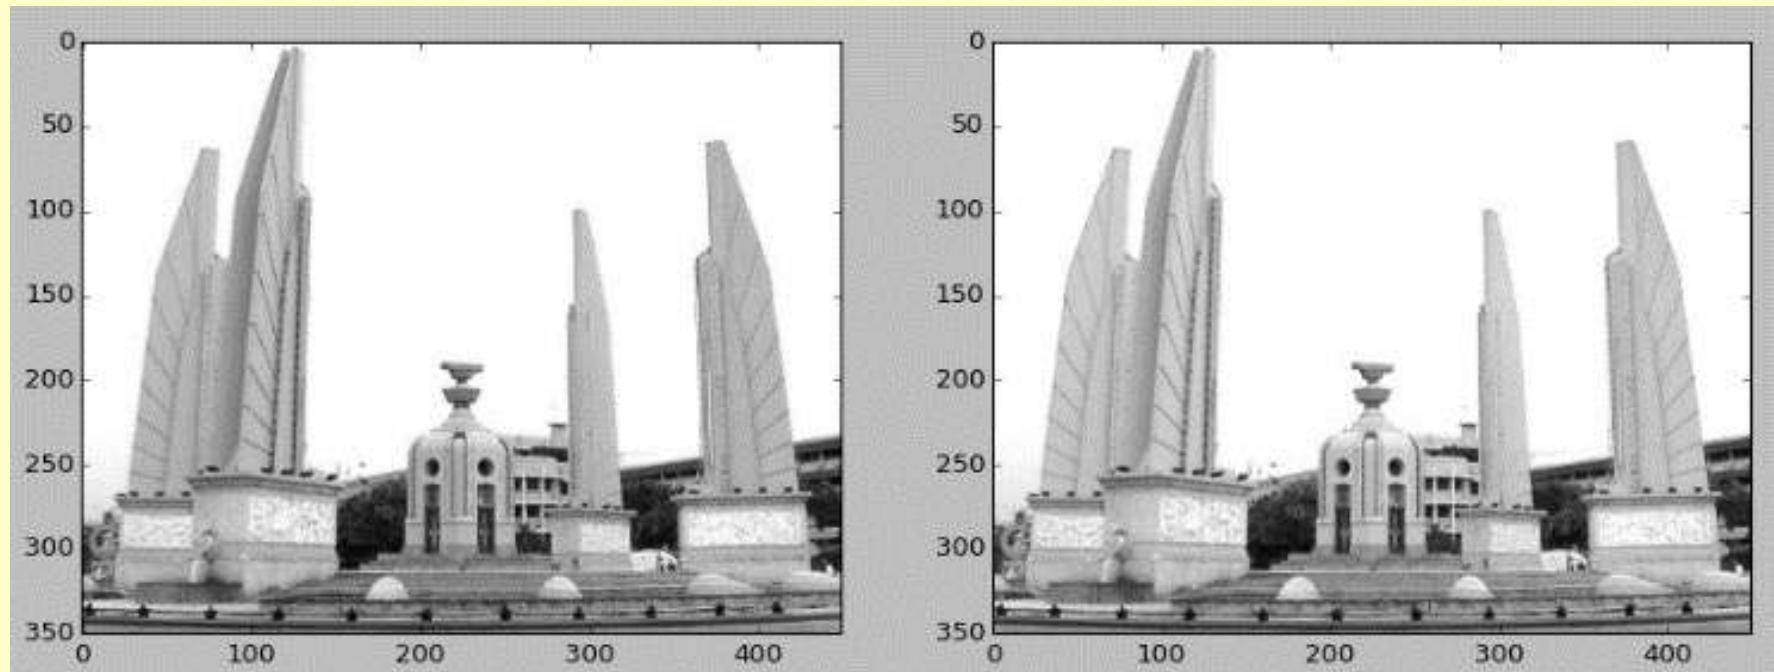
Basic Image Processing (2): Grayscale Image

```
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

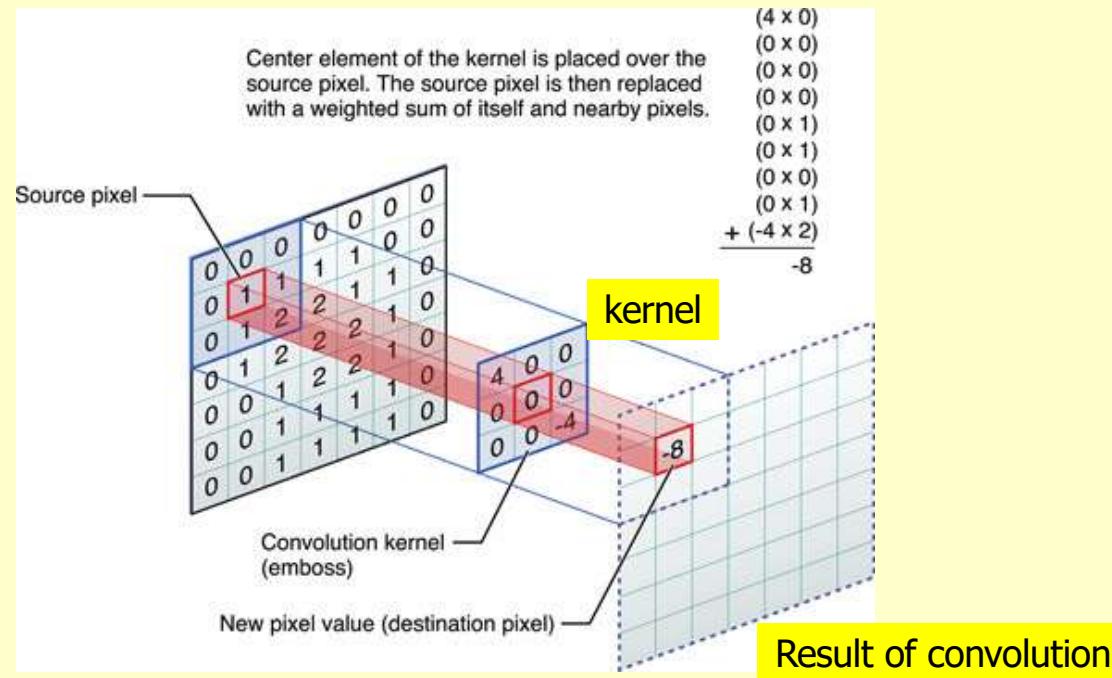
image = mpimg.imread('monument.png')
plt.subplot(1, 2, 1)
plt.imshow(image)
gray = np.ndarray(image.shape)
gray[:, :, 0] = \
gray[:, :, 1] = \
gray[:, :, 2] = (image[:, :, 0]+image[:, :, 1]+image[:, :, 2]) / 3
plt.subplot(1, 2, 2)
plt.imshow(gray)
plt.show()
```



Another Grayscale Image: $0.299*R + 0.587*G + 0.114*B$



Convolution



Changing the kernel (filter) results in different image processing outcomes

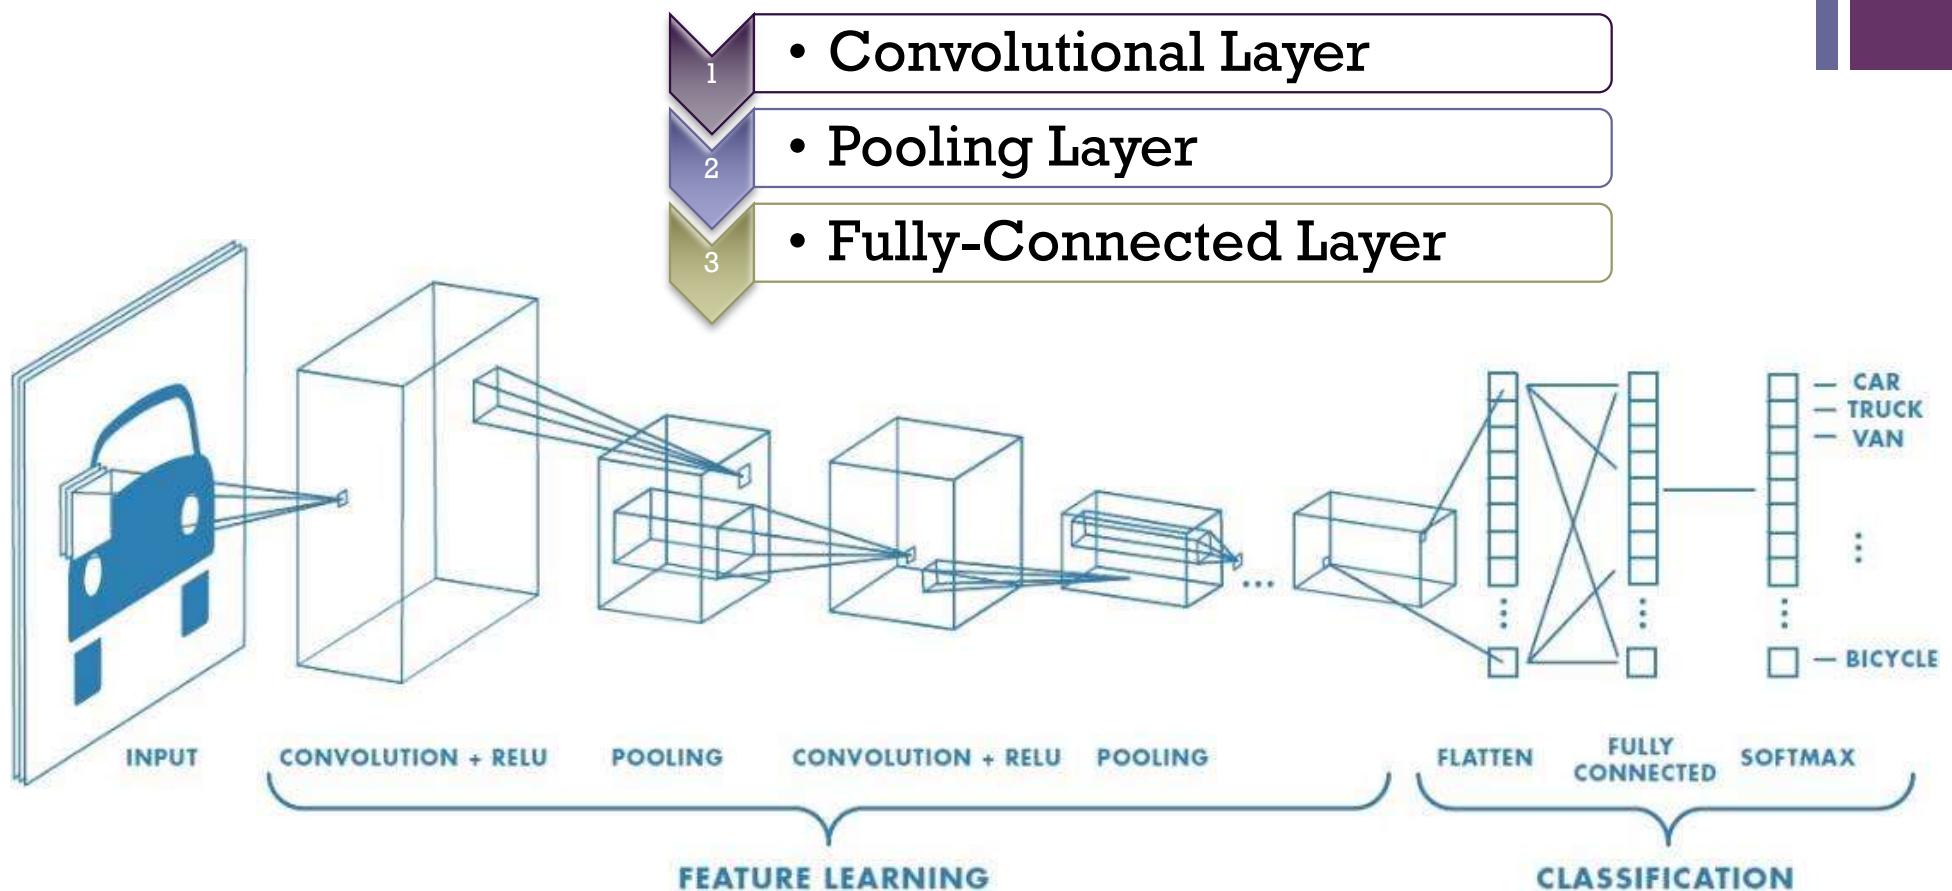
Image Convolution

$$\text{Original} \quad * \quad \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \text{Blur (with a mean filter)}$$

$$\text{Original} \quad * \quad \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 1 & 0 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} = \text{Shifted left By 1 pixel}$$

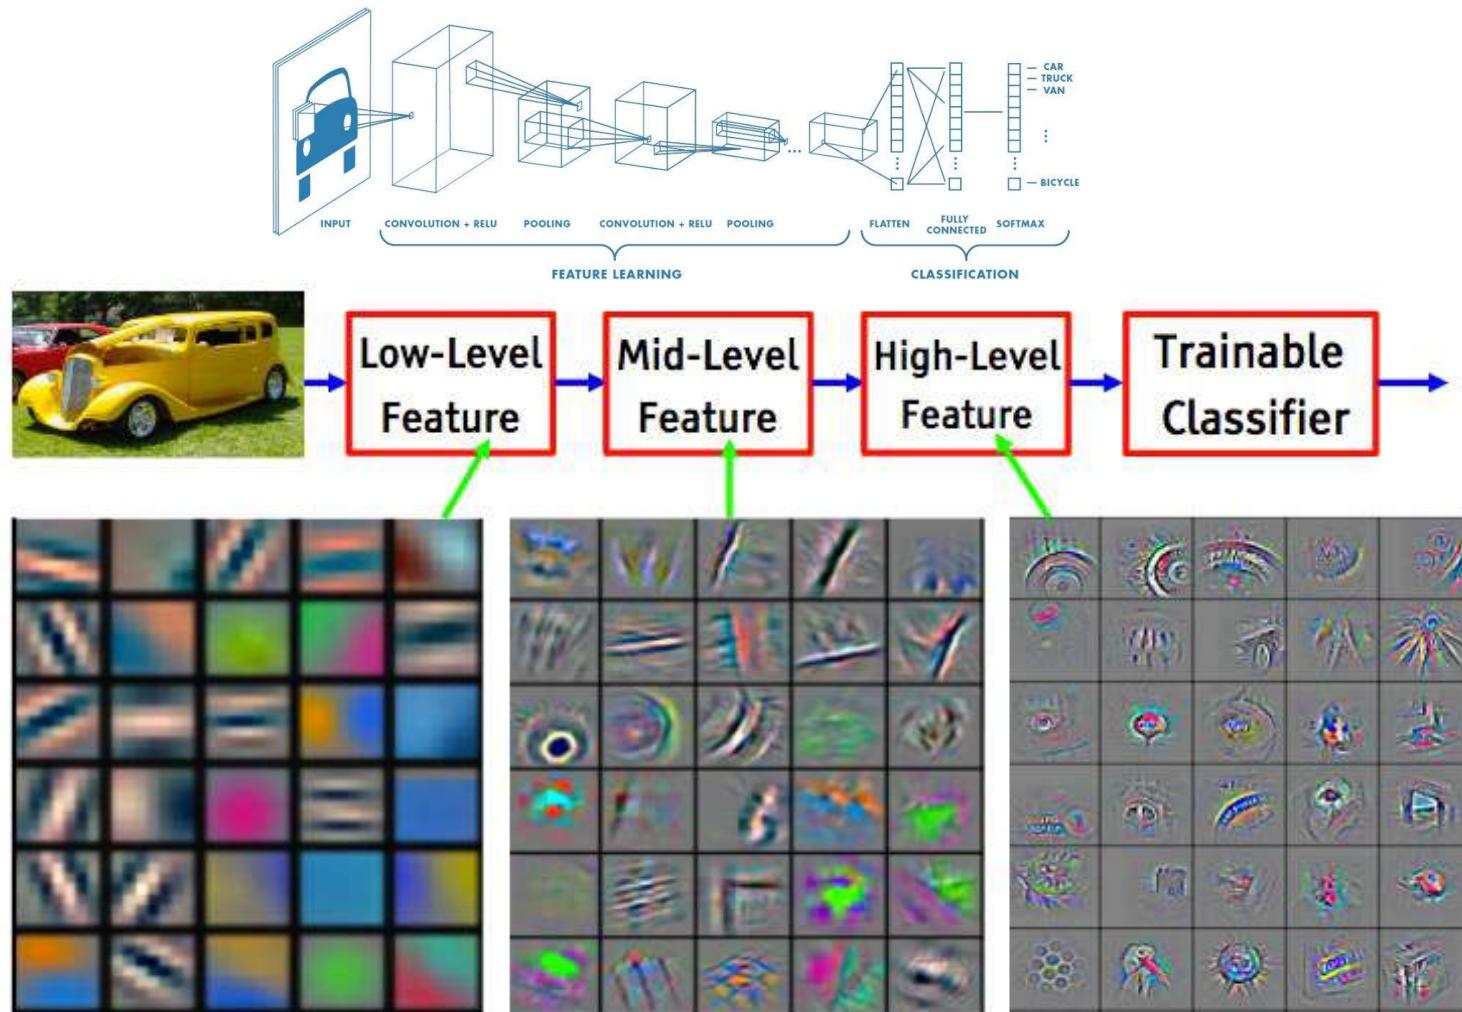


Convolutional Neural Networks (CNN)





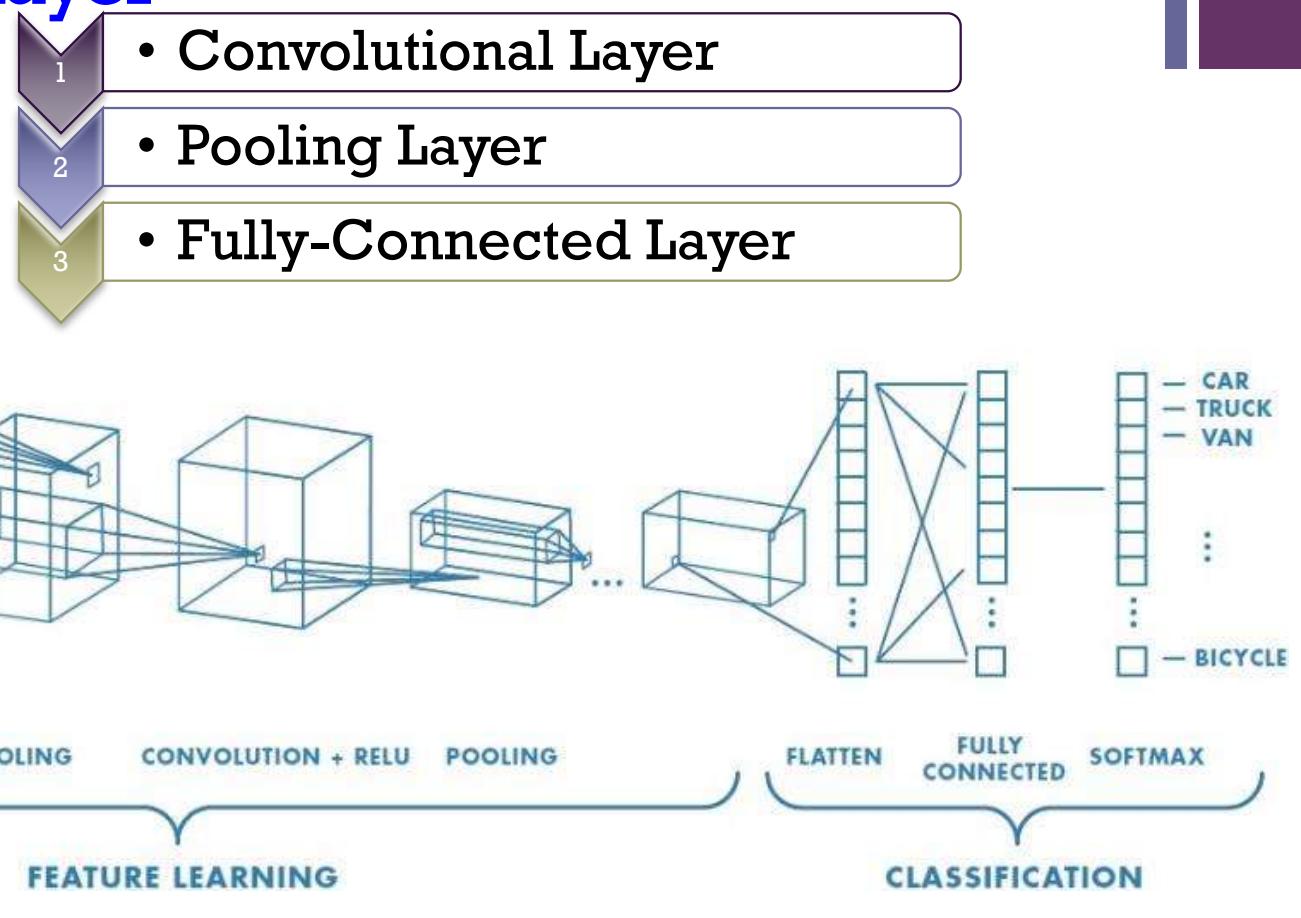
CNN (cont.)





Convolutional Neural Networks (CNN)

1) Convolutional Layer





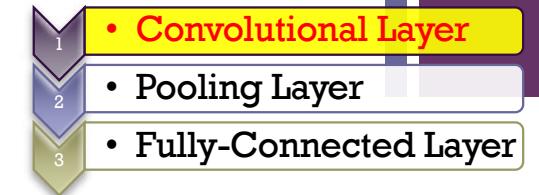
Layer1: Convolutional Layer

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

(3*3 filter)

1	0	1
0	1	0
1	0	1



4		

Convolved Feature

x[::,::,0]	0 0 0 0 0 0 0	0 1 0 2 2 2 0	0 0 0 0 0 0 0	0 2 0 2 2 2 0	0 1 0 0 0 0 0	0 1 0 0 2 1 0	0 0 0 0 0 0 0
x[::,::,1]	0 0 0 0 0 0 0	0 1 0 1 0 1 0	0 0 0 0 0 1 0	0 0 0 2 0 0 0	0 2 0 0 0 0 0	0 0 0 1 0 0 0	0 0 0 0 0 0 0
x[::,::,2]	0 0 0 0 0 0 0	0 0 2 2 0 0 0	0 0 0 0 0 1 0	0 1 1 2 0 2 0	0 2 0 0 0 0 0	0 0 1 1 0 1 0	0 0 0 0 0 0 0

Filter W0 (3x3x3)

w0[::,:,0]	1 0 0	-1 0 0	0 -1 1	-1 -1 0	0 -1 1	1 0 0	
w0[::,:,1]							
w0[::,:,2]							
Bias b0 (1x1x1)	1						

Filter W1 (3x3x3)

w1[::,:,0]	1 -1 0	-1 0 0	0 1 -1	0 1 -1	0 0 0	0 1 1	0 -1 0
w1[::,:,1]							
w1[::,:,2]							

Output Volume (3x3x2)

o[::,:,0]	2 -2 -1	2 -3 -3	2 -1 -2	0 1 -1	-2 4 -1	3 3 0	-2 2 -1
b1[::,:,0]	0						

$$\begin{aligned}
 o(2,2,0) &= \sum x[::,0] \times w[::,0] + \sum x[::,1] \times w[::,1] + \sum x[::,2] \times w[::,2] + b_0 \\
 &= 0 \times 1 + 0 \times 0 + 0 \times 0 + 2 \times (-1) + 1 \times 0 + 0 \times 0 + 0 \times 0 + 0 \times (-1) + 0 \times 1 \\
 &\quad + 0 \times (-1) + 0 \times (-1) + 0 \times 0 + 0 \times (-1) + 0 \times (-1) + 0 \times 1 + 0 \times 1 + 0 \times 0 + 0 \times 1 \\
 &\quad + 0 \times 1 + 0 \times 1 + 0 \times 0 + 0 \times 0 + 1 \times (-1) + 0 \times 1 + 0 \times 1 + 0 \times 1 + 0 \times 1 \\
 &\quad + 1 \\
 &= -2
 \end{aligned}$$

CNNs: 1) Convolutional Layer

Feature Extraction

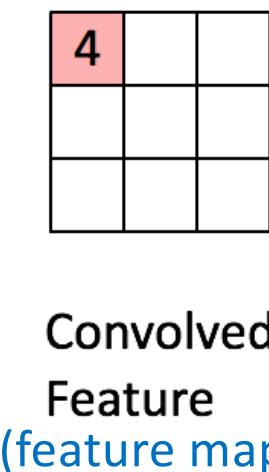
<http://cs231n.github.io/convolutional-networks/>

(3*3 filter)

$$w^T x + b$$

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

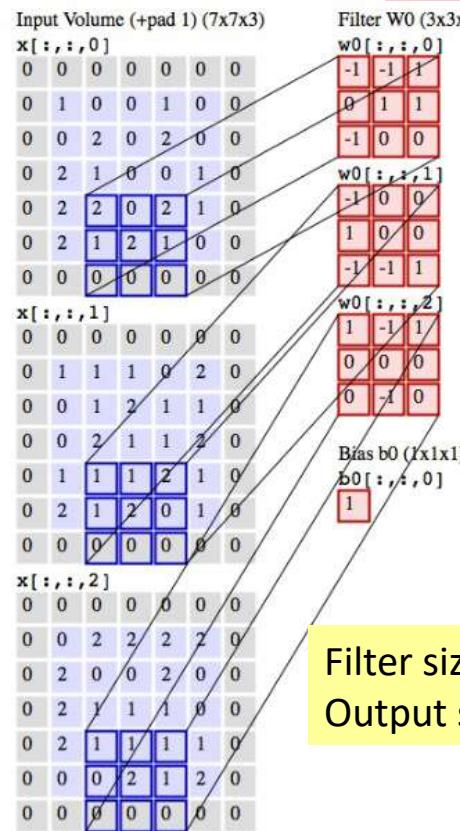


Convolved Feature
(feature map)

Filter size=(3*3), #Filters=1, Stride=1, Padding=0

Output size = $(5 + 2(0) - 3) / 1 + 1 = 3$

1	0	1
0	1	0
1	0	1



Filter size=(3*3*3), #Filters=2, Stride=2, Padding=1
Output size = $(5 + 2(1) - 3) / 2 + 1 = 3$

$$\text{Output size} = (N + 2*P - F) / S + 1$$

(feature map1)

Input Volume (+pad 1) (7x7x3)	Filter W0 (3x3x3)	Filter W1 (3x3x3)	Output Volume (3x3x2)
x[:, :, 0]	w0[:, :, 0]	w1[:, :, 0]	o[:, :, 0]
0 0 0 0 0 0 0	-1 -1 1 0 1 1 -1 0 0	1 0 1 0 0 0 -1 -1 -1	1 -1 -3
0 1 0 0 1 0 0	0 1 1 -1 0 0	-1 -1 -1 0 0 1 -1 -1 1	2 1 -4
0 0 2 0 2 0 0	-1 0 0 1 0 0 -1 -1 1	0 0 1 -1 0 0 0 1 -1	3 5 -4
0 2 1 0 0 1 0	1 0 0 -1 0 0 0 0 0	-1 -1 -1 0 0 1 -1 0 0	2 -10 -6
0 2 2 0 2 1 0	-1 0 0 1 0 0 0 0 0	0 0 1 -1 0 0 0 1 -1	-2 -6 -6
0 2 1 2 1 0 0	1 0 0 -1 0 0 0 0 0	-1 -1 1 0 0 0 0 1 -1	-2 -2 -3
0 0 0 0 0 0 0	-1 -1 1 0 0 0 0 -1 0	0 -1 -1 -1 0 0 0 1 -1	
x[:, :, 1]	w0[:, :, 1]	w1[:, :, 1]	
0 0 0 0 0 0 0	1 -1 1 0 0 0 0 -1 0	0 -1 -1 -1 0 0 0 1 -1	
0 1 1 1 0 2 0	0 0 0 0 0 0 0 -1 0	-1 0 0 0 0 0 0 1 -1	
0 0 1 2 1 1 0	0 0 0 0 0 0 0 -1 0	0 0 0 -1 0 0 0 1 -1	
0 0 2 1 1 2 0	0 0 0 0 0 0 0 -1 0	0 0 0 -1 0 0 0 1 -1	
0 1 1 1 2 1 0	0 0 0 0 0 0 0 -1 0	0 0 0 -1 0 0 0 1 -1	
0 2 1 2 0 1 0	0 0 0 0 0 0 0 -1 0	0 0 0 -1 0 0 0 1 -1	
0 0 0 0 0 0 0	0 0 0 0 0 0 0 -1 0	0 0 0 -1 0 0 0 1 -1	
x[:, :, 2]	b0[:, :, 0]		
0 0 0 0 0 0 0	1		
0 0 2 2 2 2 0			
0 2 0 0 2 0 0			
0 2 1 1 1 0 0			
0 2 1 1 1 1 0			
0 0 0 2 1 2 0			
0 0 0 0 0 0 0			

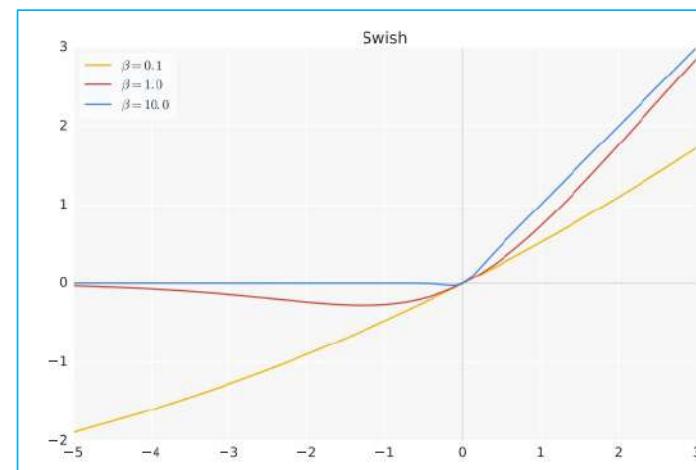
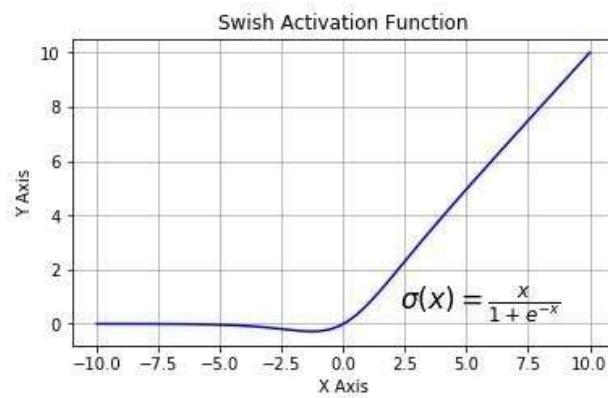
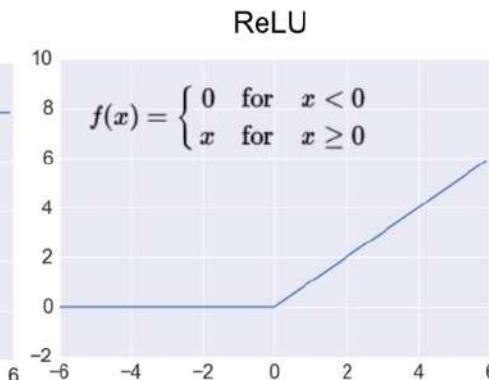
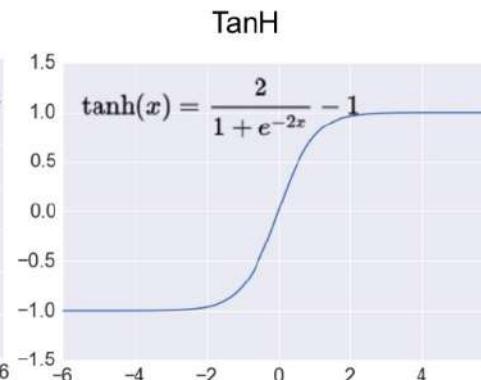
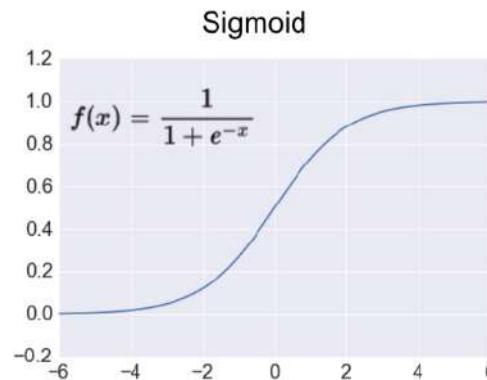
toggle movement

R: $2(-1)+2(1)+2(1)+1(1) = 3$
G: $1(-1)+1(1) = 0$
B: $1(1)+1(-1)+1(1) = 1$
Ans : $4+1 = 5$



CNNs: 1) Convolutional Layer (cont.)

Non-Linear Activation Function

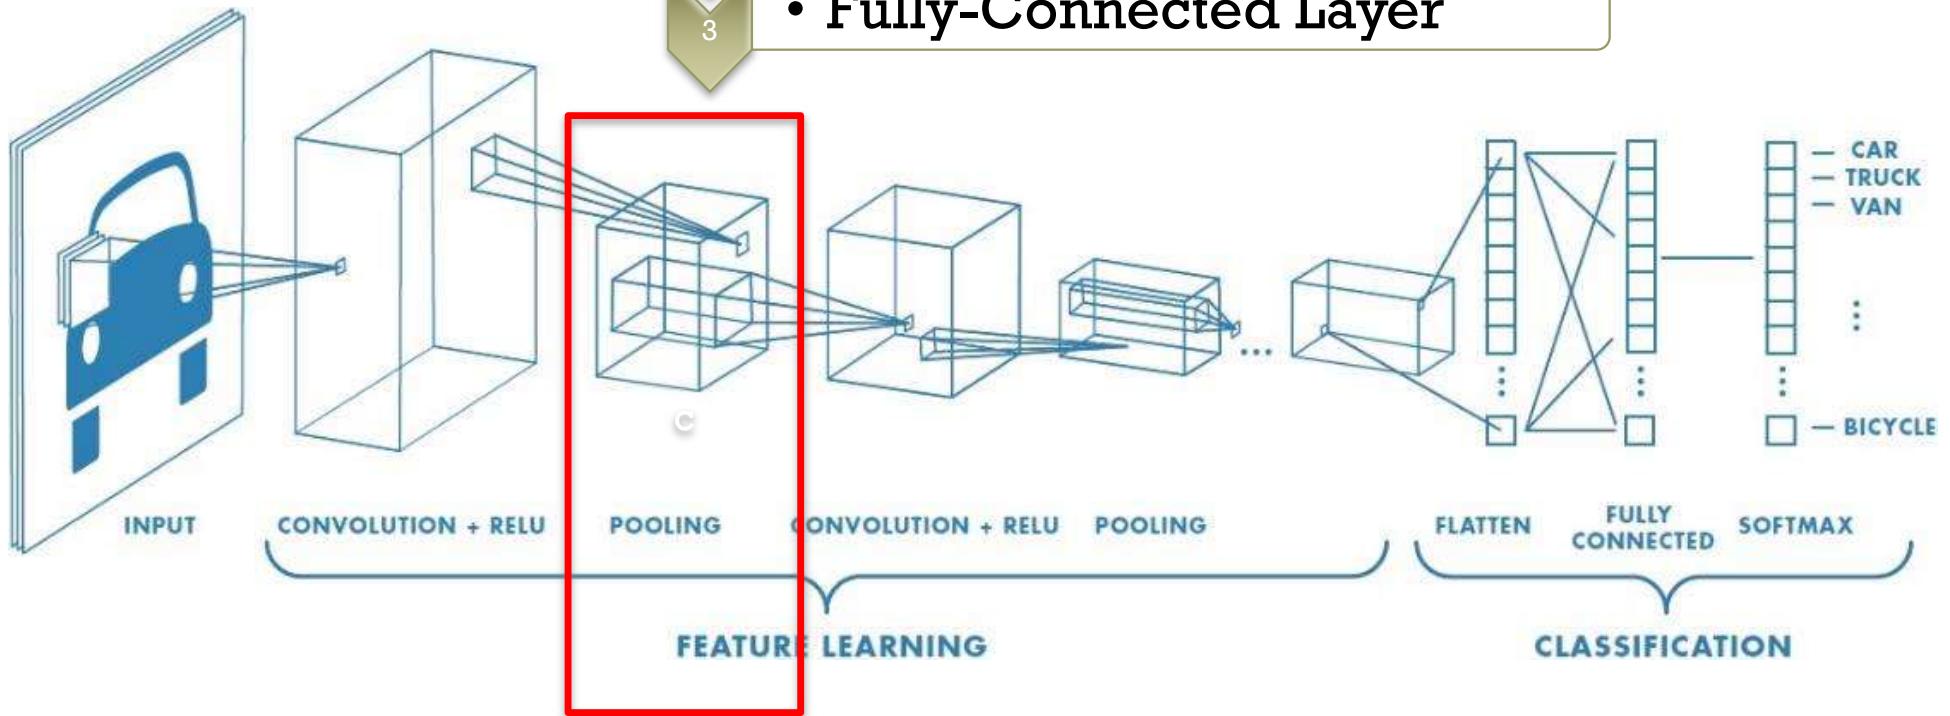




Convolutional Neural Networks (CNN)

2) Pooling Layer

- Convolutional Layer
- Pooling Layer
- Fully-Connected Layer

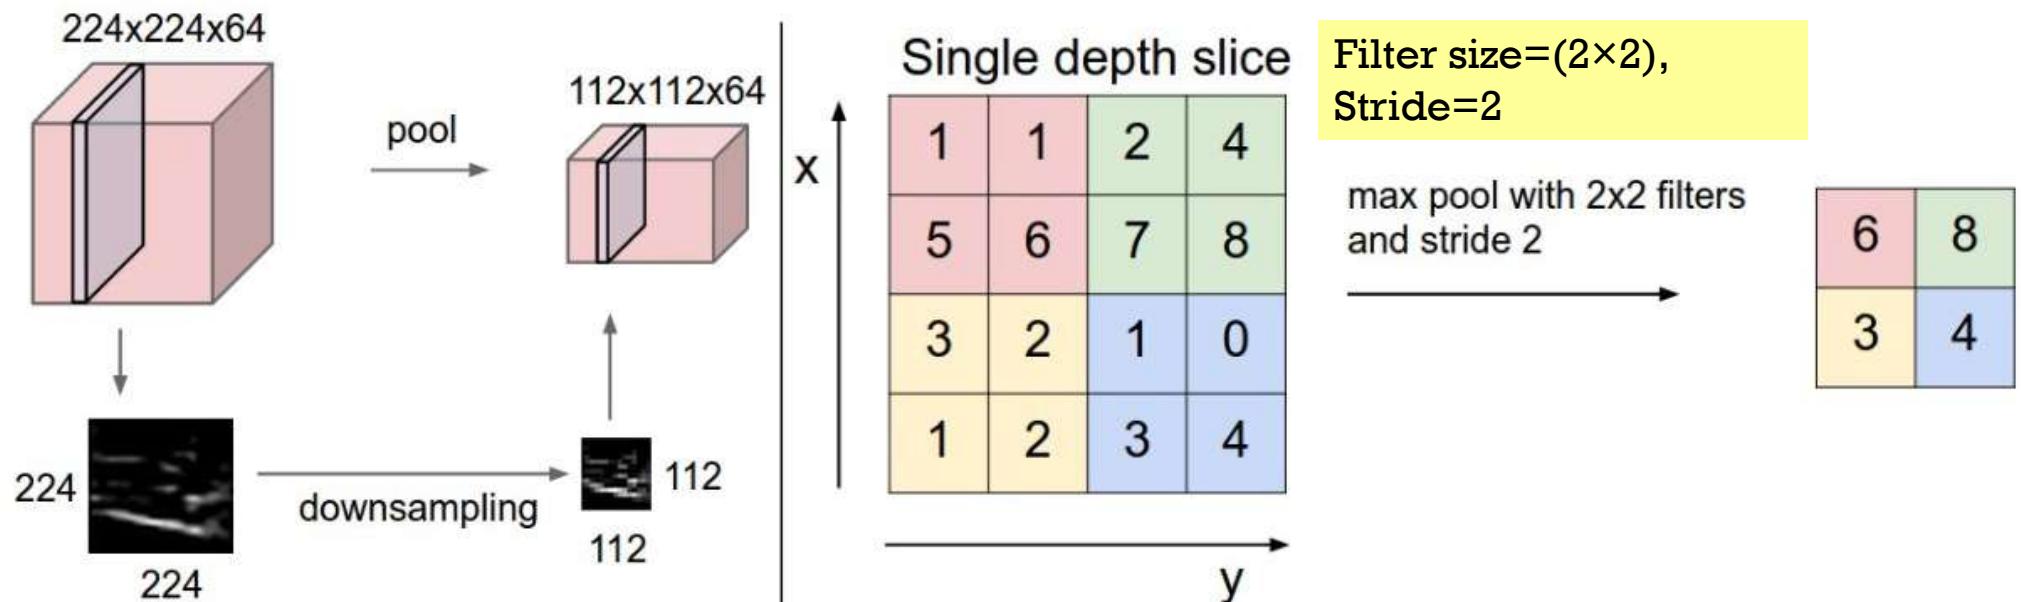


CNNs: 2) Pooling Layer

<http://cs231n.github.io/convolutional-networks/>

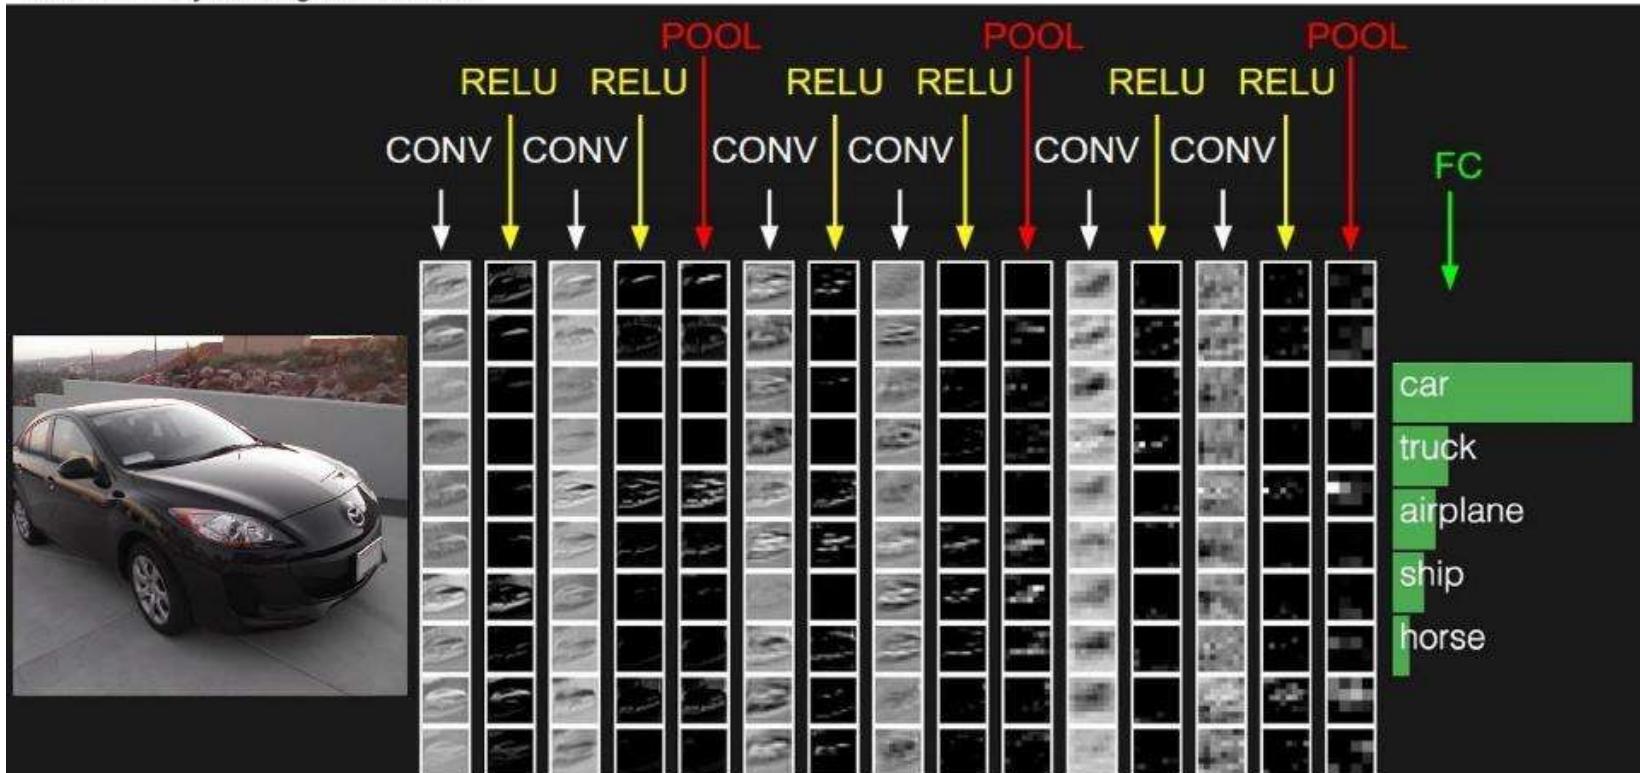
Summary: **2 parameters** for pooling layer
• Filter size, Stride

- Feature selection (reduction); downsampling





two more layers to go: POOL/FC

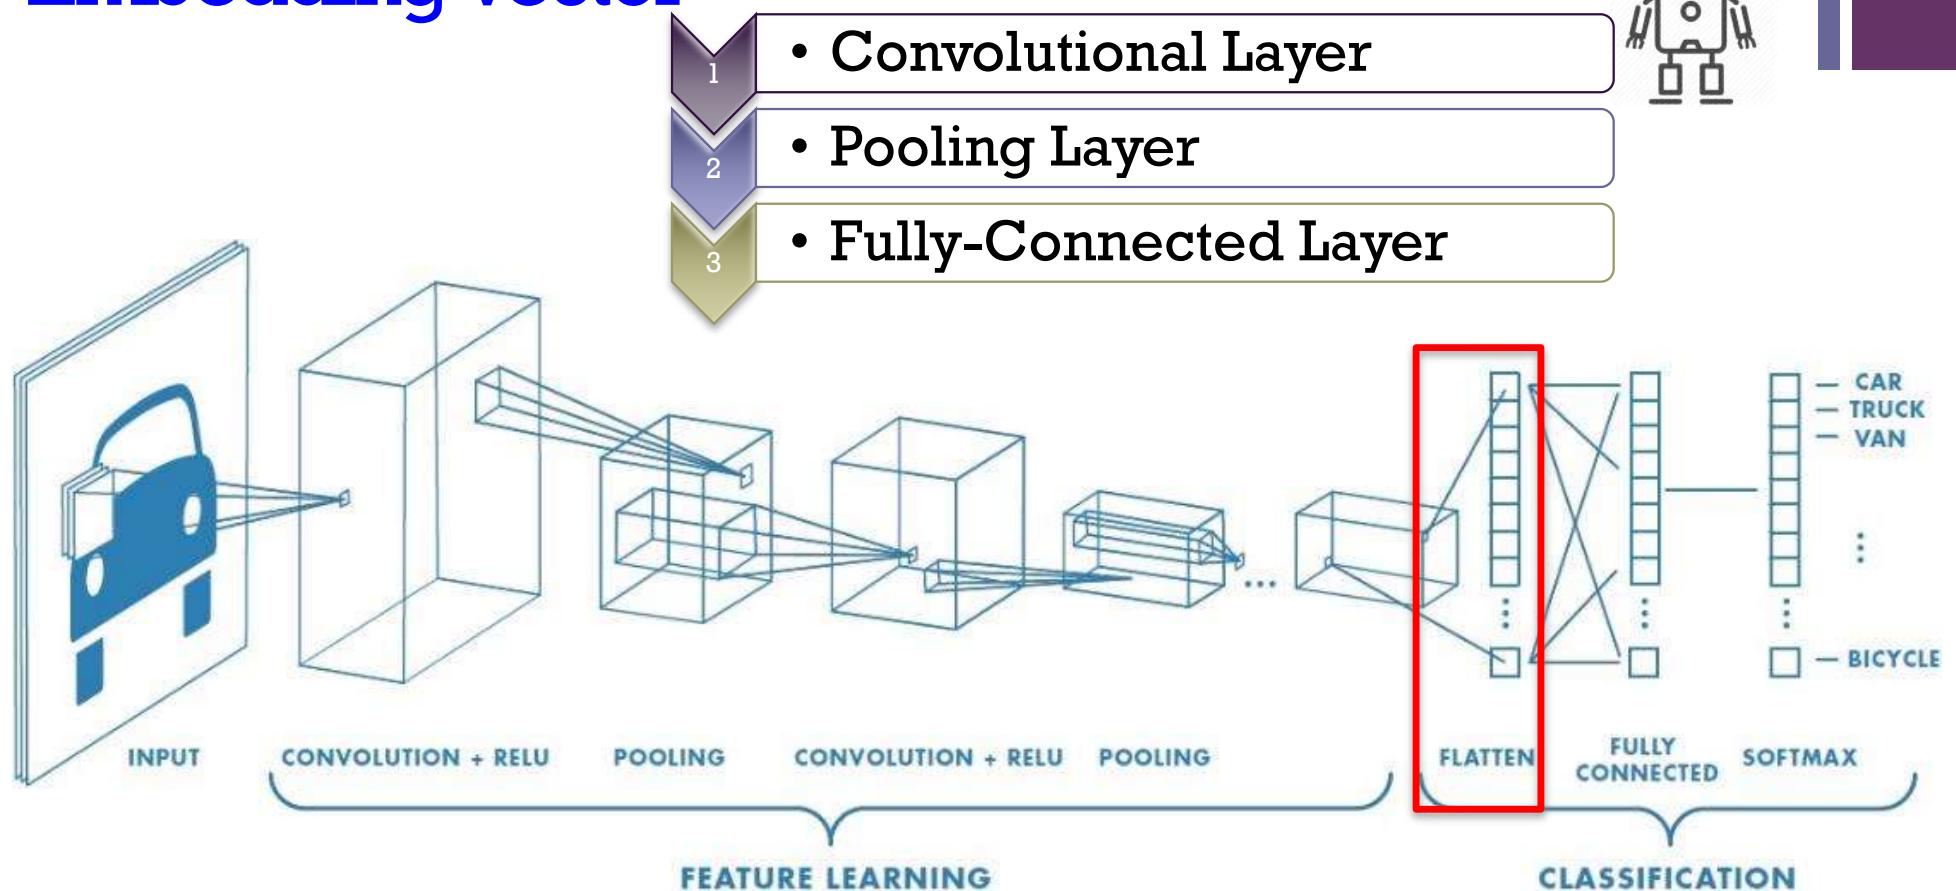




Convolutional Neural Networks

Embedding Vector

x1	x2	x3	x4	Corona
0.7	0.2	-0.5	-0.1	Yes

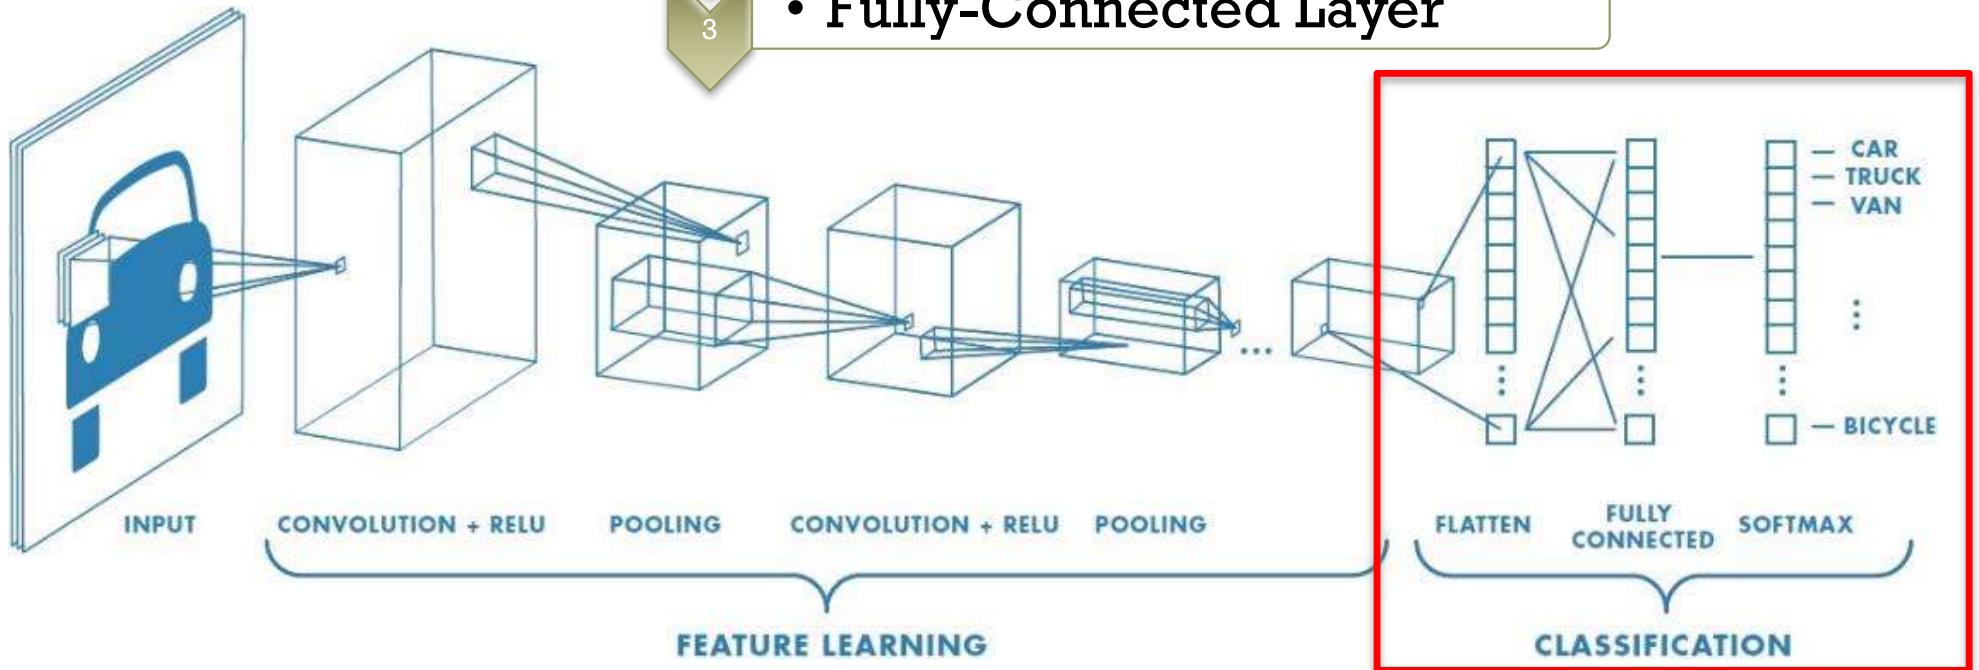




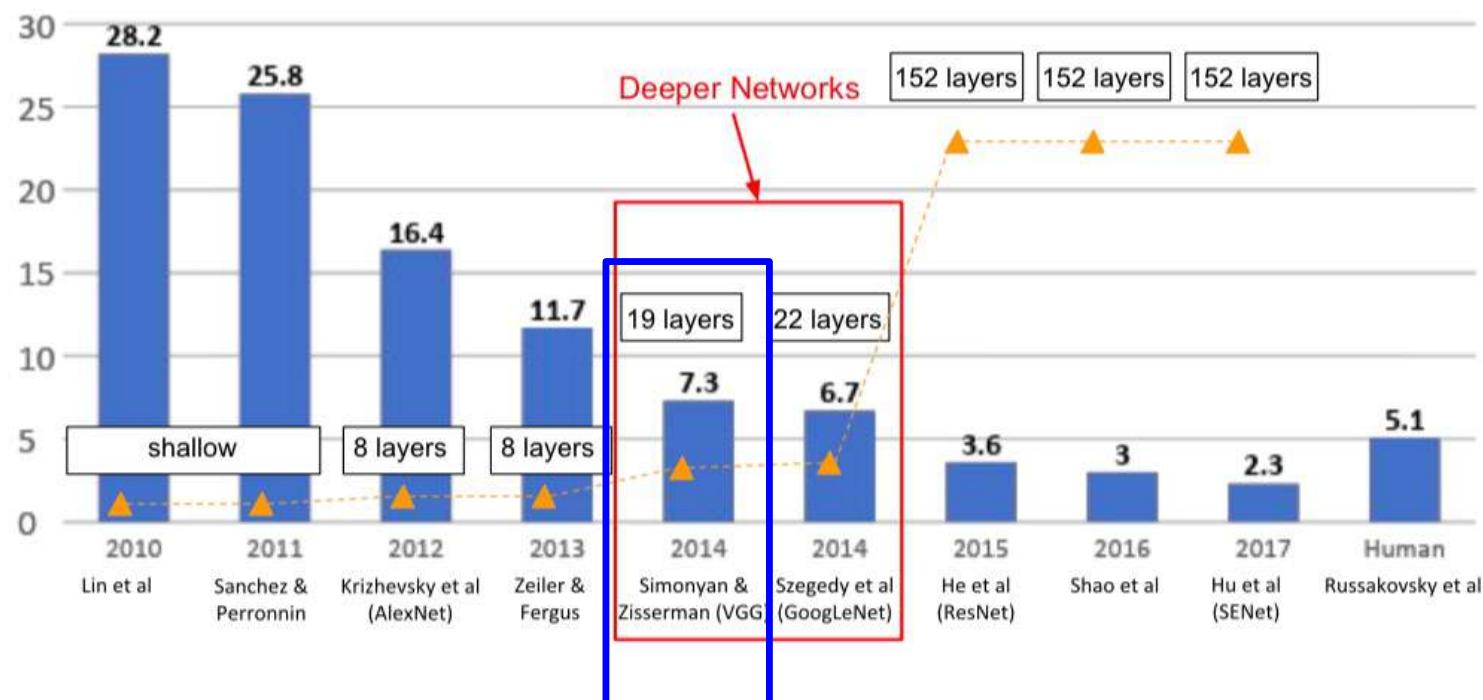
Convolutional Neural Networks (CNN)

3) Fully-Connected Layer (FC) = Neural Networks

- Convolutional Layer
- Pooling Layer
- Fully-Connected Layer



ImageNet Large Scale Visual Recognition Challenge (ILSVRC 2nd runner-up in 2014): VGG19



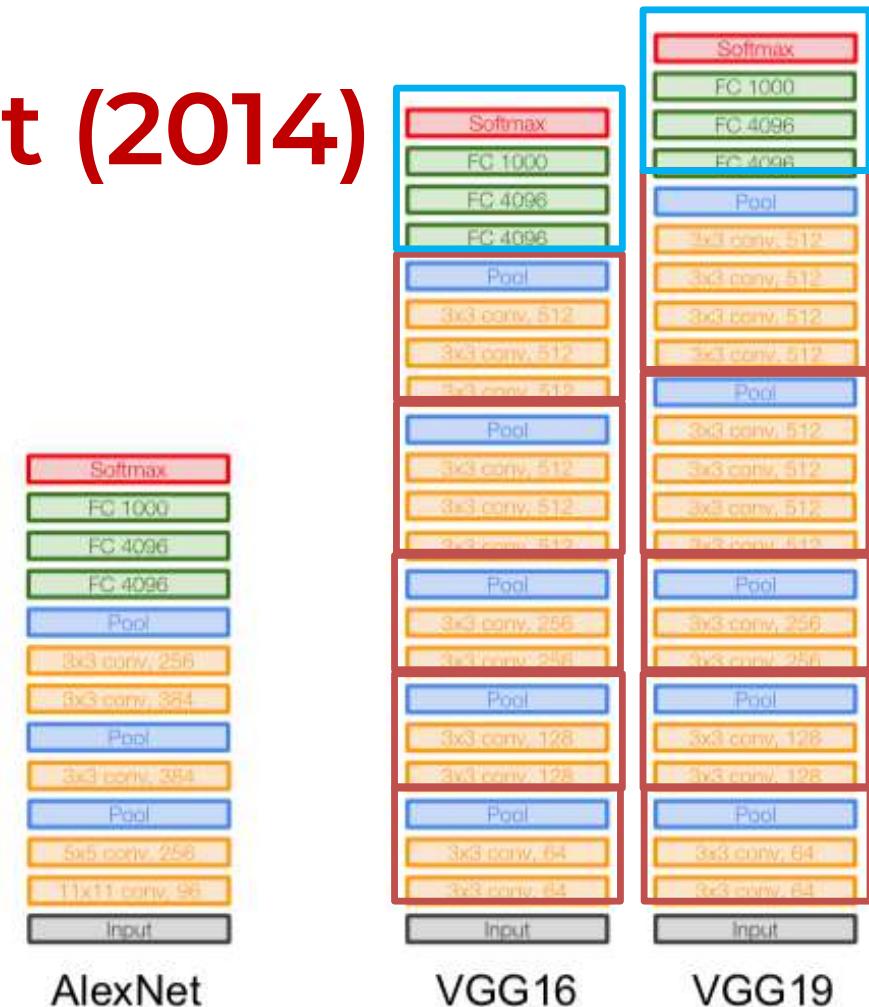
CNN Example: VGGNet (2014)

[Simonyan and Zisserman, 2014]

Small filters, Deeper networks

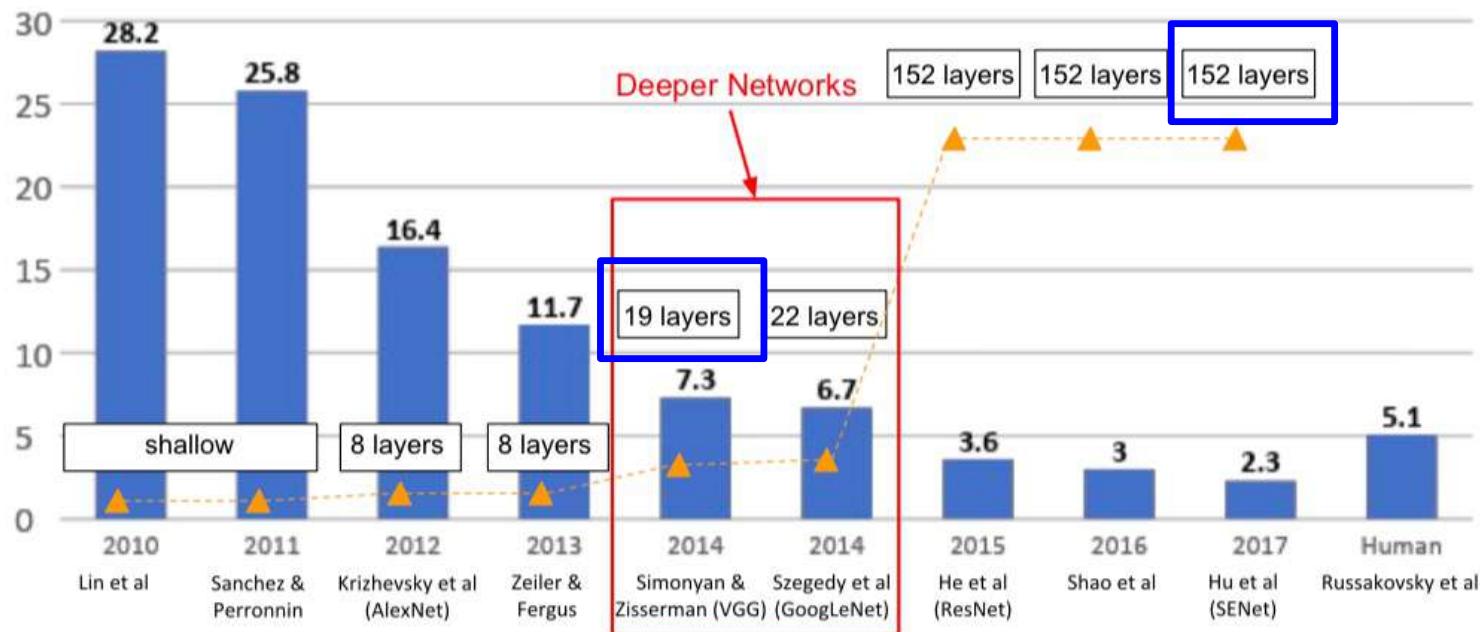
8 layers (AlexNet) → 16-19 layers (VGG16Net)

Only 3x3 CONV Stride 1, pad 1
And 2x2 MAX POOL stride 2



Network	A	B	C	D	E
Number of parameters (Millions)	133	133	134	138	144

ImageNet Large Scale Visual Recognition Challenge (ILSVRC 2nd runner-up in 2014): Deeper than VGG19





- Overview
- Input
- Model
- Sequential
- activations
- applications
 - Overview
 - DenseNet121
 - DenseNet169
 - DenseNet201
 - EfficientNetB0
 - EfficientNetB1
 - EfficientNetB2
 - EfficientNetB3
 - EfficientNetB4
 - EfficientNetB5
 - EfficientNetB6
 - EfficientNetB7
 - InceptionResNetV2
 - InceptionV3
 - MobileNet
 - MobileNetV2
 - MobileNetV3Large
 - MobileNetV3Small
 - NASNetLarge
 - NASNetMobile

Modules

[densenet](#) module: DenseNet models for Keras.

[efficientnet](#) module: EfficientNet models for Keras.

[imagenet_utils](#) module: Utilities for ImageNet data preprocessing & prediction decoding.

[inception_resnet_v2](#) module: Inception-ResNet V2 model for Keras.

[inception_v3](#) module: Inception V3 model for Keras.

[mobilenet](#) module: MobileNet v1 models for Keras.

[mobilenet_v2](#) module: MobileNet v2 models for Keras.

[nasnet](#) module: NASNet-A models for Keras.

[resnet](#) module: ResNet models for Keras.

[resnet50](#) module: Public API for tf.keras.applications.resnet50 namespace.

[resnet_v2](#) module: ResNet v2 models for Keras.

[vgg16](#) module: VGG16 model for Keras.

[vgg19](#) module: VGG19 model for Keras.

Model 1

- VGG19 (random initialized weights) + 2 Dense layers + Output layer

```
1 base_model = VGG19(weights=None, include_top=False, input_shape=(224, 224, 3))
2
3 for layer in base_model.layers:
4     layer.trainable = True
5
6 x = base_model.output
7 x = Flatten()(x)
8 x = Dense(1024)(x)
9 x = Dropout(0.5)(x)
10 x = Dense(512)(x)
11 x = Dropout(0.5)(x)
12 output = Dense(num_class, activation='softmax')(x)
13
```

[!\[\]\(1b72a119a678e7a0a5c908017deea8ba_img.jpg\) Search Docs](#)[Package Reference](#)[torchvision.datasets](#)[torchvision.io](#)[torchvision.models](#)[torchvision.ops](#)[torchvision.transforms](#)[torchvision.utils](#)[PyTorch Libraries](#)[PyTorch](#)[torchaudio](#)[torchtext](#)[torchvision](#)[TorchElastic](#)[TorchServe](#)[PyTorch on XLA Devices](#)

TORCHVISION.MODELS

The models subpackage contains definitions of models for addressing different tasks, including: image classification, pixelwise semantic segmentation, object detection, instance segmentation, person keypoint detection and video classification.

Classification

The models subpackage contains definitions for the following model architectures for image classification:

- [AlexNet](#)
- [VGG](#)
- [ResNet](#)
- [SqueezeNet](#)
- [DenseNet](#)
- [Inception v3](#)
- [GoogLeNet](#)
- [ShuffleNet v2](#)
- [MobileNetV2](#)
- [MobileNetV3](#)
- [ResNeXt](#)
- [Wide ResNet](#)
- [MNASNet](#)

We provide pre-trained models, using the PyTorch `torch.utils.model_zoo`. These can be constructed by passing `pretrained=True`:

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)
alexnet = models.alexnet(pretrained=True)
squeezenet = models.squeeze1_0(pretrained=True)
vgg16 = models.vgg16(pretrained=True)
densenet = models.densenet161(pretrained=True)
inception = models.inception_v3(pretrained=True)
googlenet = models.googlenet(pretrained=True)
shufflenet = models.shufflenet_v2_x1_0(pretrained=True)
mobilenet_v2 = models.mobilenet_v2(pretrained=True)
mobilenet_v3_large = models.mobilenet_v3_large(pretrained=True)
mobilenet_v3_small = models.mobilenet_v3_small(pretrained=True)
resnext50_32x4d = models.resnext50_32x4d(pretrained=True)
wide_resnet50_2 = models.wide_resnet50_2(pretrained=True)
mnasnet = models.mnasnet1_0(pretrained=True)
```

<https://pytorch.org/vision/0.9/models.html>



v1.1.0 ▾

[pytorch-transformers](#)[Star 131,473](#)[Search docs](#)

NOTES

[Installation](#)[Quickstart](#)[Pretrained models](#)[Examples](#)[Notebooks](#)[Loading Google AI or OpenAI pre-trained weights or PyTorch dump](#)[Serialization best-practices](#)[Converting Tensorflow Checkpoints](#)[Migrating from pytorch-pretrained-bert](#)[BERTology](#)

ⓘ You are viewing legacy docs. Go to [latest documentation](#) instead.

[Docs » Pytorch-Transformers](#)

Pytorch-Transformers

PyTorch-Transformers is a library of state-of-the-art pre-trained models for Natural Language Processing (NLP).

The library currently contains PyTorch implementations, pre-trained model weights, usage scripts and conversion utilities for the following models:

1. [BERT](#) (from Google) released with the paper [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) by Jacob Devlin, Ming-Wei Lee and Kristina Toutanova.
2. [GPT](#) (from OpenAI) released with the paper [Improving Language Understanding by Generative Pre-Training](#) by Alec Radford, Karthik Narasimhan, Tim Salim Sutskever.
3. [GPT-2](#) (from OpenAI) released with the paper [Language Models are Unsupervised Multitask Learners](#) by Alec Radford*, Jeffrey Wu*, Rewon Child, David Luan and Ilya Sutskever**.
4. [Transformer-XL](#) (from Google/CMU) released with the paper [Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context](#) by Zihang Dai*, Zhen Yang, Jaime Carbonell, Quoc V. Le, Ruslan Salakhutdinov.
5. [XLNet](#) (from Google/CMU) released with the paper [XLNet: Generalized Autoregressive Pretraining for Language Understanding](#) by Zhilin Yang*, Zihang Dai, Jaime Carbonell, Ruslan Salakhutdinov, Quoc V. Le.
6. [XLM](#) (from Facebook) released together with the paper [Cross-lingual Language Model Pretraining](#) by Guillaume Lample and Alexis Conneau.

Notes

- [Installation](#)



Image Classification Tasks

Classification



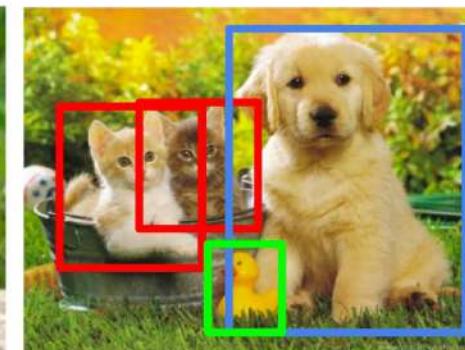
CAT

Classification + Localization



CAT

Object Detection



CAT, DOG, DUCK

Instance Segmentation



CAT, DOG, DUCK

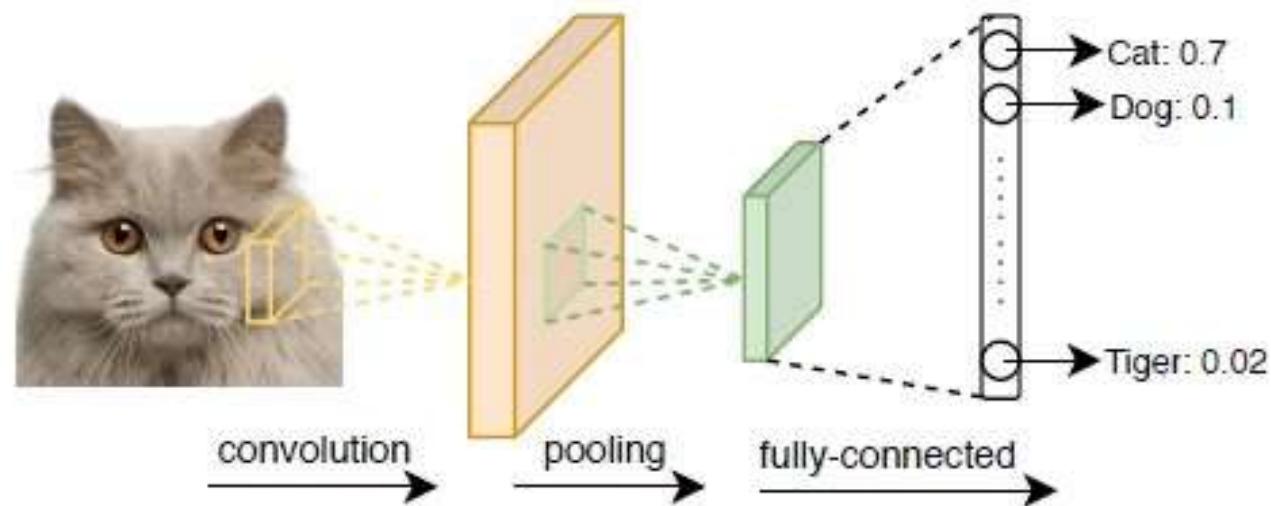
Single object

Multiple objects

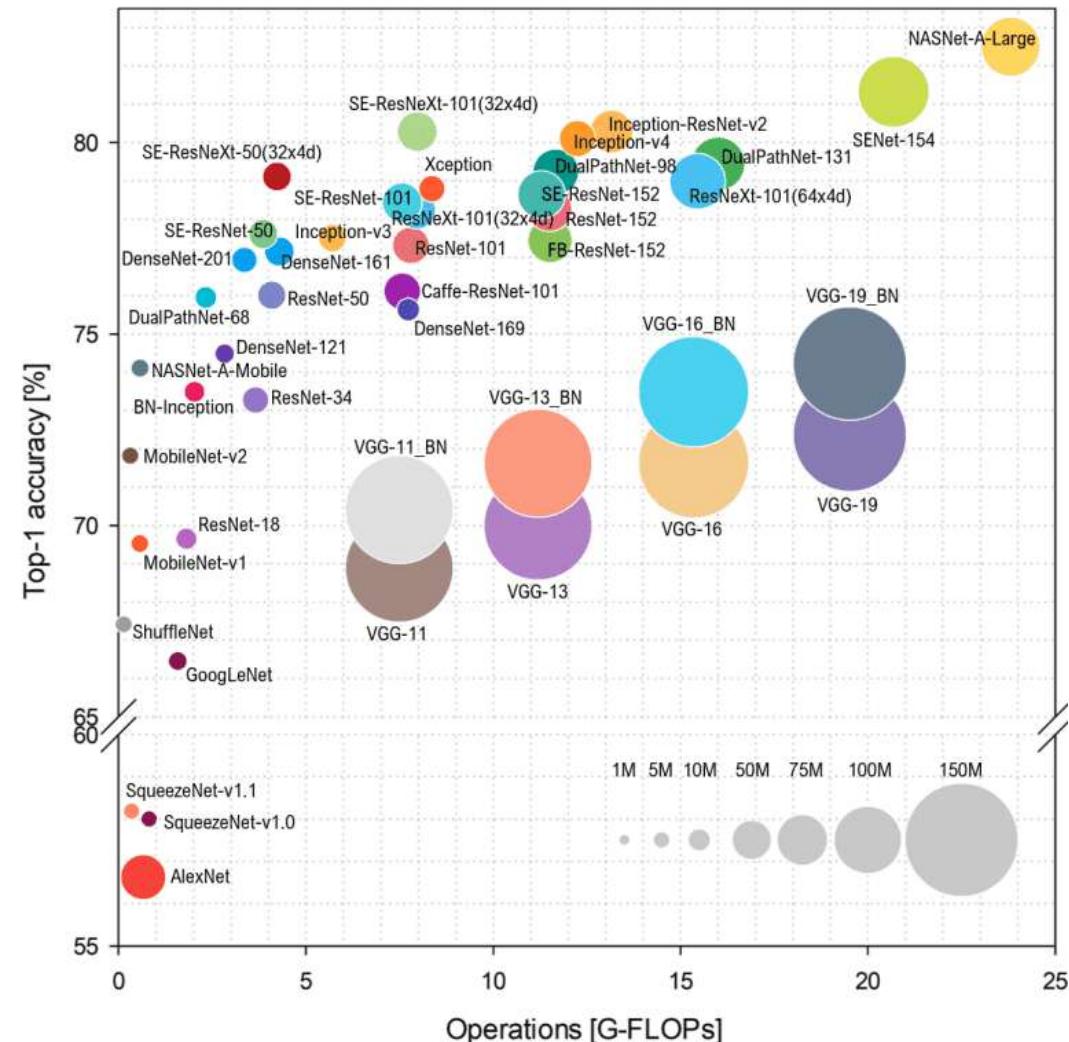
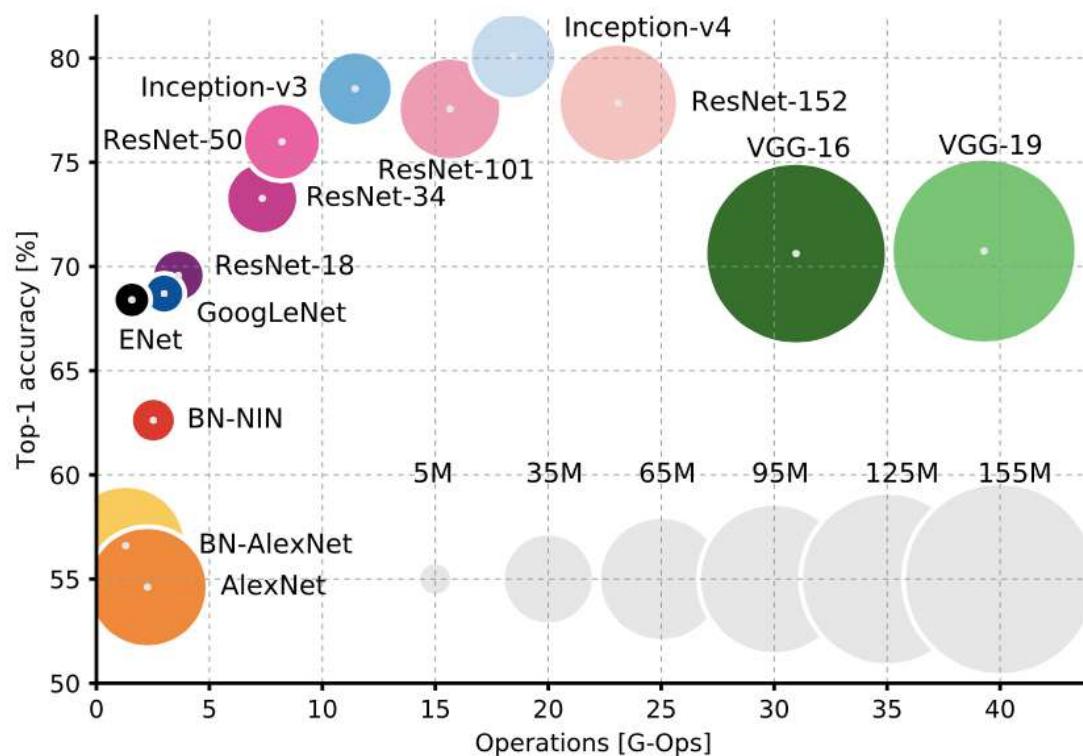


Image Classification (recap)

Convolutional Neural Network



SOTA of Image Classification



https://blog.csdn.net/qq_34216467/article/details/83061692

<https://theaisummer.com/cnn-architectures/>



Common Image Classification Models

Model	Year	Team / Organization
ResNet	2015	Microsoft Research
MobileNet	2017	Google
EfficientNet	2019	Google
Vision Transformer (ViT)	2020	Google Research
Swin Transformer	2021	Microsoft Research
ConvNeXt	2022	Meta AI

Image Classification on ImageNet



<https://paperswithcode.com/sota/image-classification-on-imagenet>



Image classification dashboard (2025)

■ CNN, Transformer (ViT, VLM)

- 2021–2022 → CNNs (EfficientNetV2) were strong, but Transformers started dominating.
- 2022–2023 → Vision–Language models (VLM) like CoCa + giant ViTs (EVA, Model Soups) pushed past 90%.
- 2024 → Specialized ViT variants (DaViT, Florence-CoSwin) consolidated.
- 2025 → Large-scale foundation models (DINOv3) + fine-tuned CoCa remain at the frontier.

■ Note:

- CoCa / CLIP → **Vision–Language (image + text)**, trained with image–text pairs, suitable for cross-modal tasks such as captioning and retrieval
- DINOv3 → **Vision-only (image)**, self-supervised ViT, backbone for computer vision

<https://paperswithcode.com/sota/image-classification-on-imagenet>

EfficientNet (May, 2019) → EffNetV2 (2021)

EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks

Mingxing Tan¹ Quoc V. Le¹

Abstract

Convolutional Neural Networks (ConvNets) are commonly developed at a fixed resource budget, and then scaled up for better accuracy if more resources are available. In this paper, we systematically study model scaling and identify that carefully balancing network depth, width, and resolution can lead to better performance. Based on this observation, we propose a new scaling method that uniformly scales all dimensions of depth/width/resolution using a simple yet highly effective *compound coefficient*. We demonstrate the effectiveness of this method on scaling up MobileNets and ResNet.

To go even further, we use neural architecture

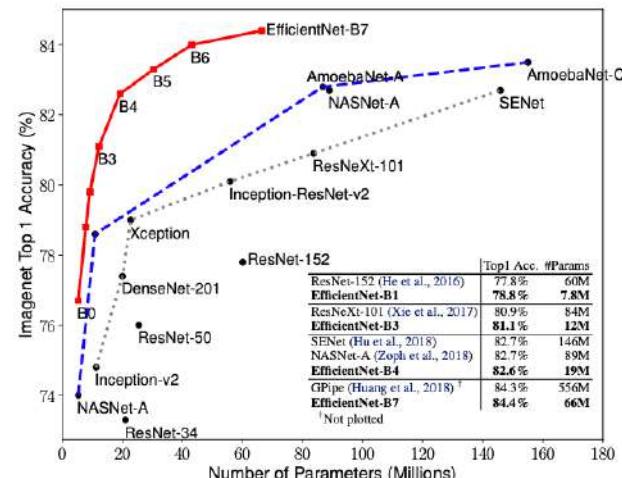
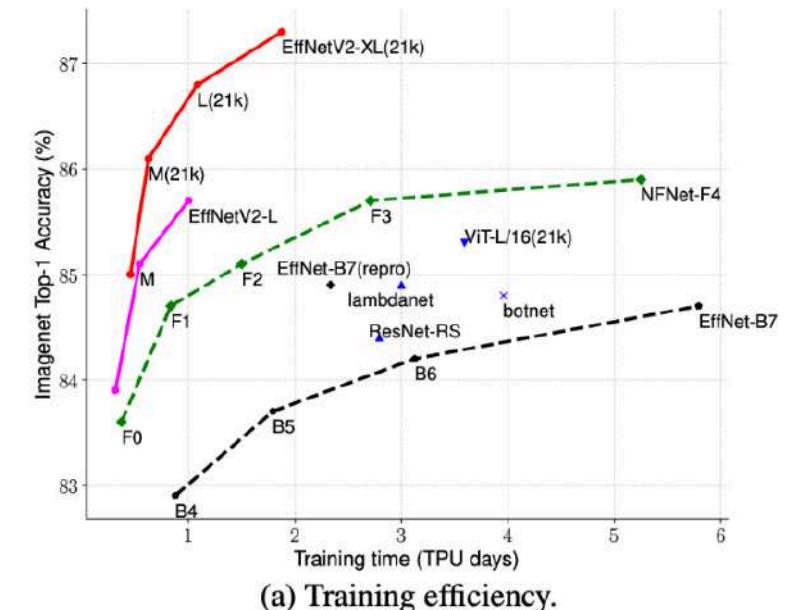


Figure 1: Model Size vs. ImageNet Accuracy. All numbers are



(a) Training efficiency.

	EfficientNet (2019)	ResNet-RS (2021)	DeiT/ViT (2021)	EfficientNetV2 (ours)
Top-1 Acc.	84.3%	84.0%	83.1%	83.9%
Parameters	43M	164M	86M	24M

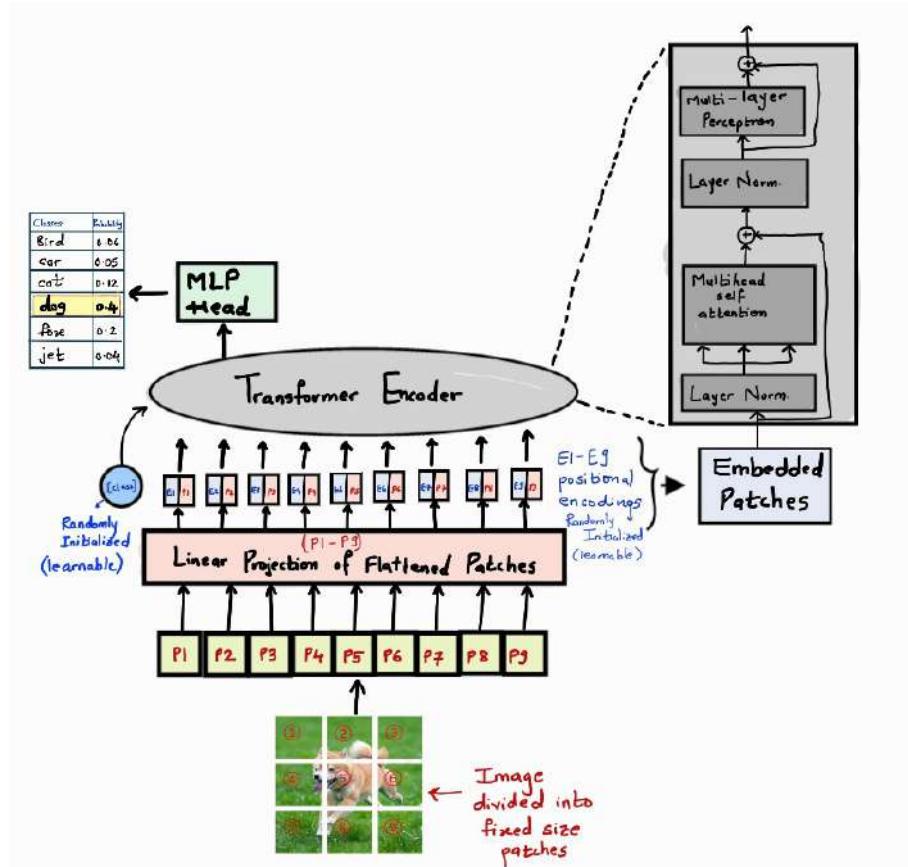
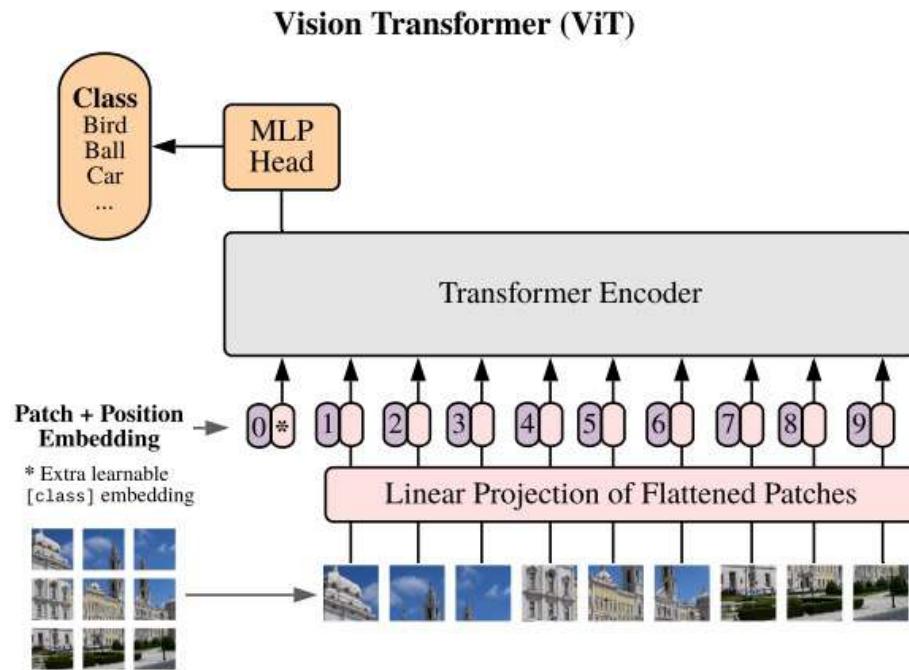
(b) Parameter efficiency.



Vision Transformer (ViT) (ICLR2021)

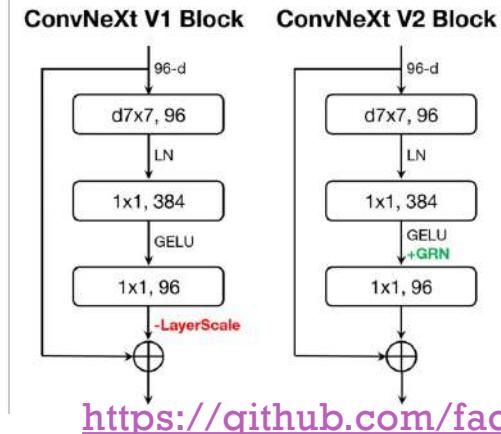
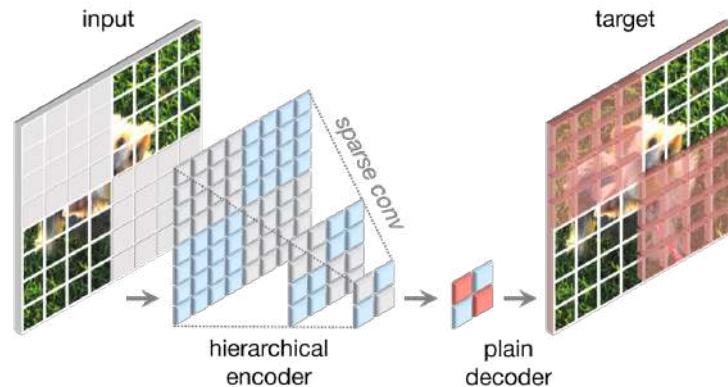
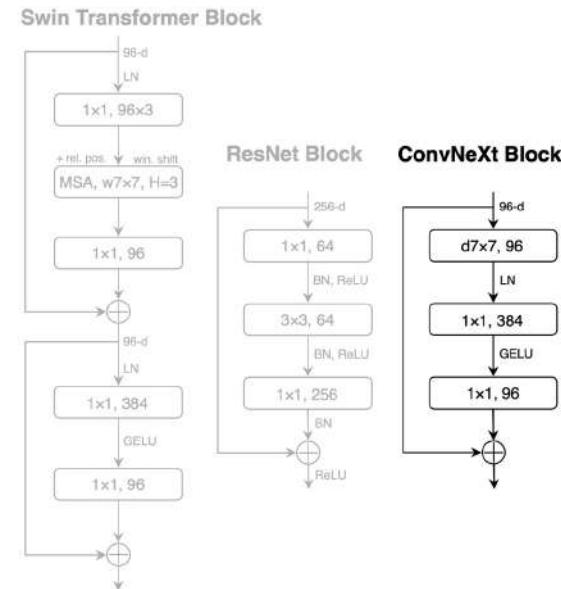
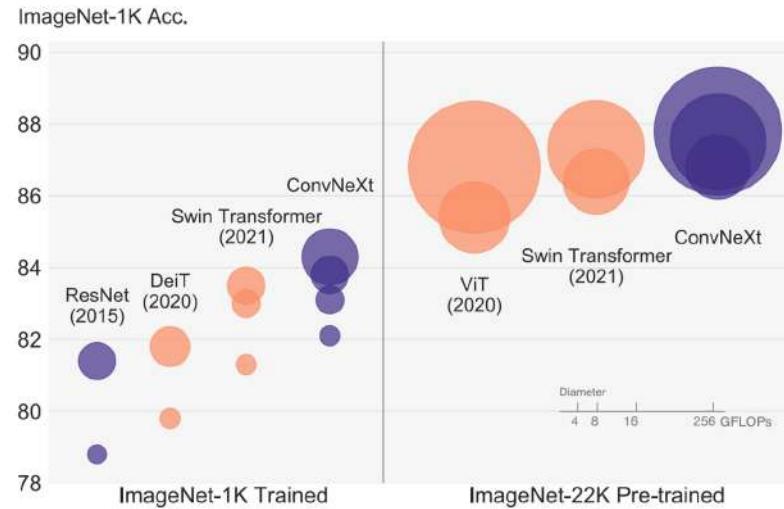


Hugging Face



<https://medium.com/machine-intelligence-and-deep-learning-lab/vit-vision-transformer-cc56c8071a20>

+ ConvNeXt (CVPR2022) → ConvNeXt V2 (2023)



<https://github.com/facebookresearch/ConvNeXt-V2>

<https://github.com/facebookresearch/ConvNeXt>

+ DINOv3: Self-supervised learning for vision at unprecedented scale

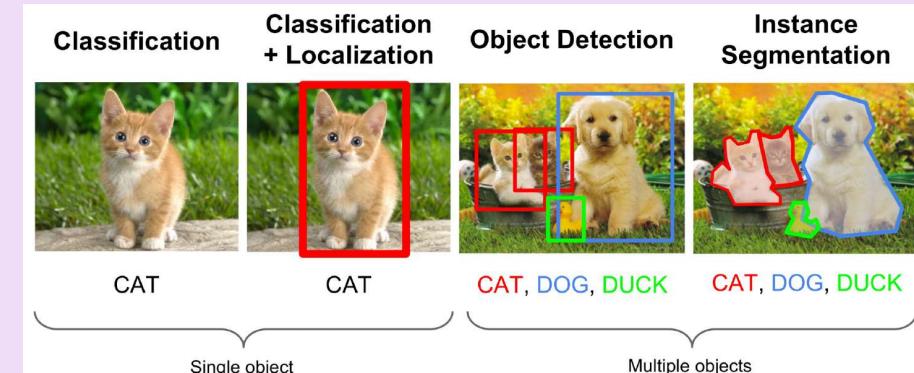
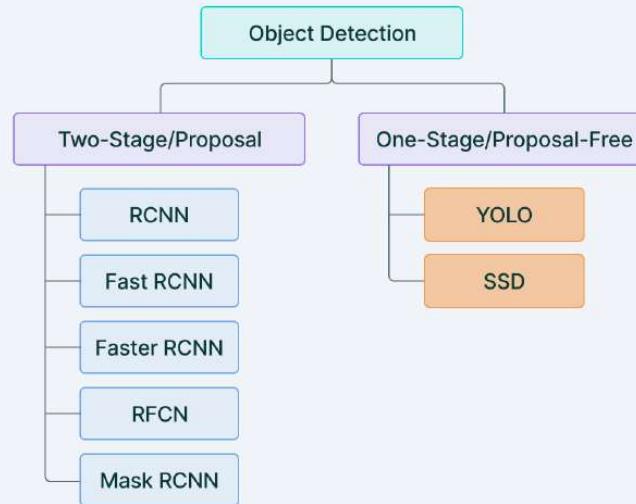
August 14, 2025 • 11 minute read

- DINOv3: Self-Supervised Vision Transformer (Meta AI, 2025)
- Vision-only (**encoder-only**) backbone (not multimodal, unlike CLIP/CoCa).
- Self-supervised learning → no human labels needed.
- **Massive scale: up to 7B parameters, trained on 1.7B images.**
- Key innovations:
 - Gram Anchoring → stabilizes feature learning.
 - Axial RoPE → improved positional encoding.
 - Resolution robustness → better handling of varied image sizes.
- Strengths:
 - Provides general-purpose visual representations.
 - **Excels in downstream tasks (classification, detection, segmentation).**
 - Foundation model role for vision, similar to LLMs in NLP.

[Link](#)

SOTA of Object Detection

One and two stage detectors



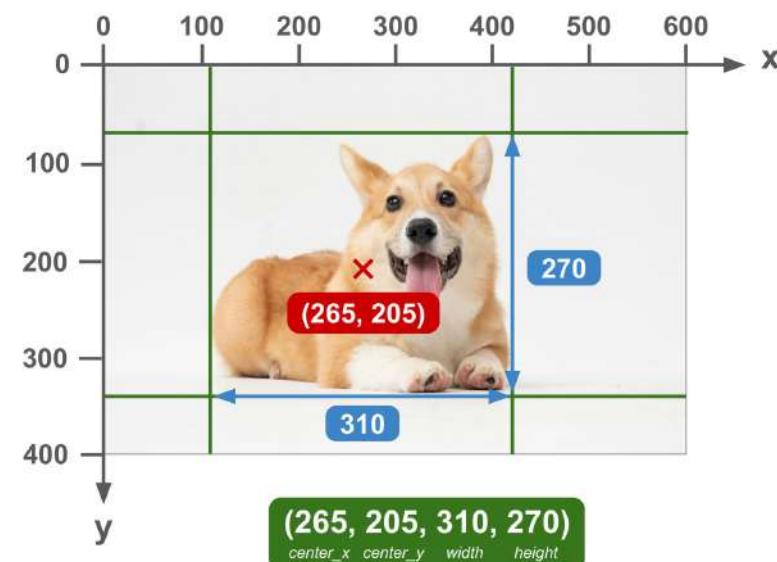
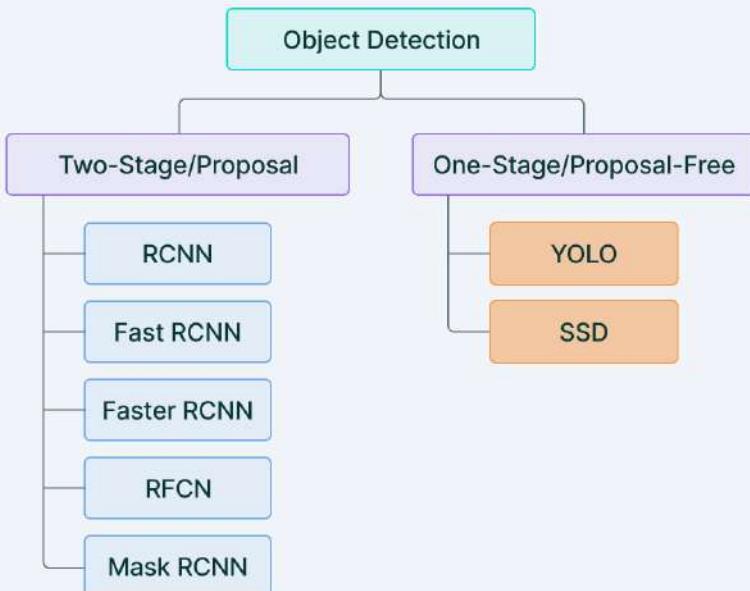
V7 Labs

https://medium.com/@jonathan_hui/real-time-object-detection-with-yolo-yolov2-28b1b93e2088



Object Detection

One and two stage detectors



corner coordinate	units	
	absolute	normalized
(left, top, right, bottom)		Pascal VOC
(left, top, width, height)		COCO
(center_x, center_y, width, height)		CreateML
		YOLO

<https://dragoneye.ai/blog/a-guide-to-bounding-box-formats/>



YOLO Open Images in New York



YOLO Open Images in New York



Joseph Redmon
6.95K subscribers

Subscribe

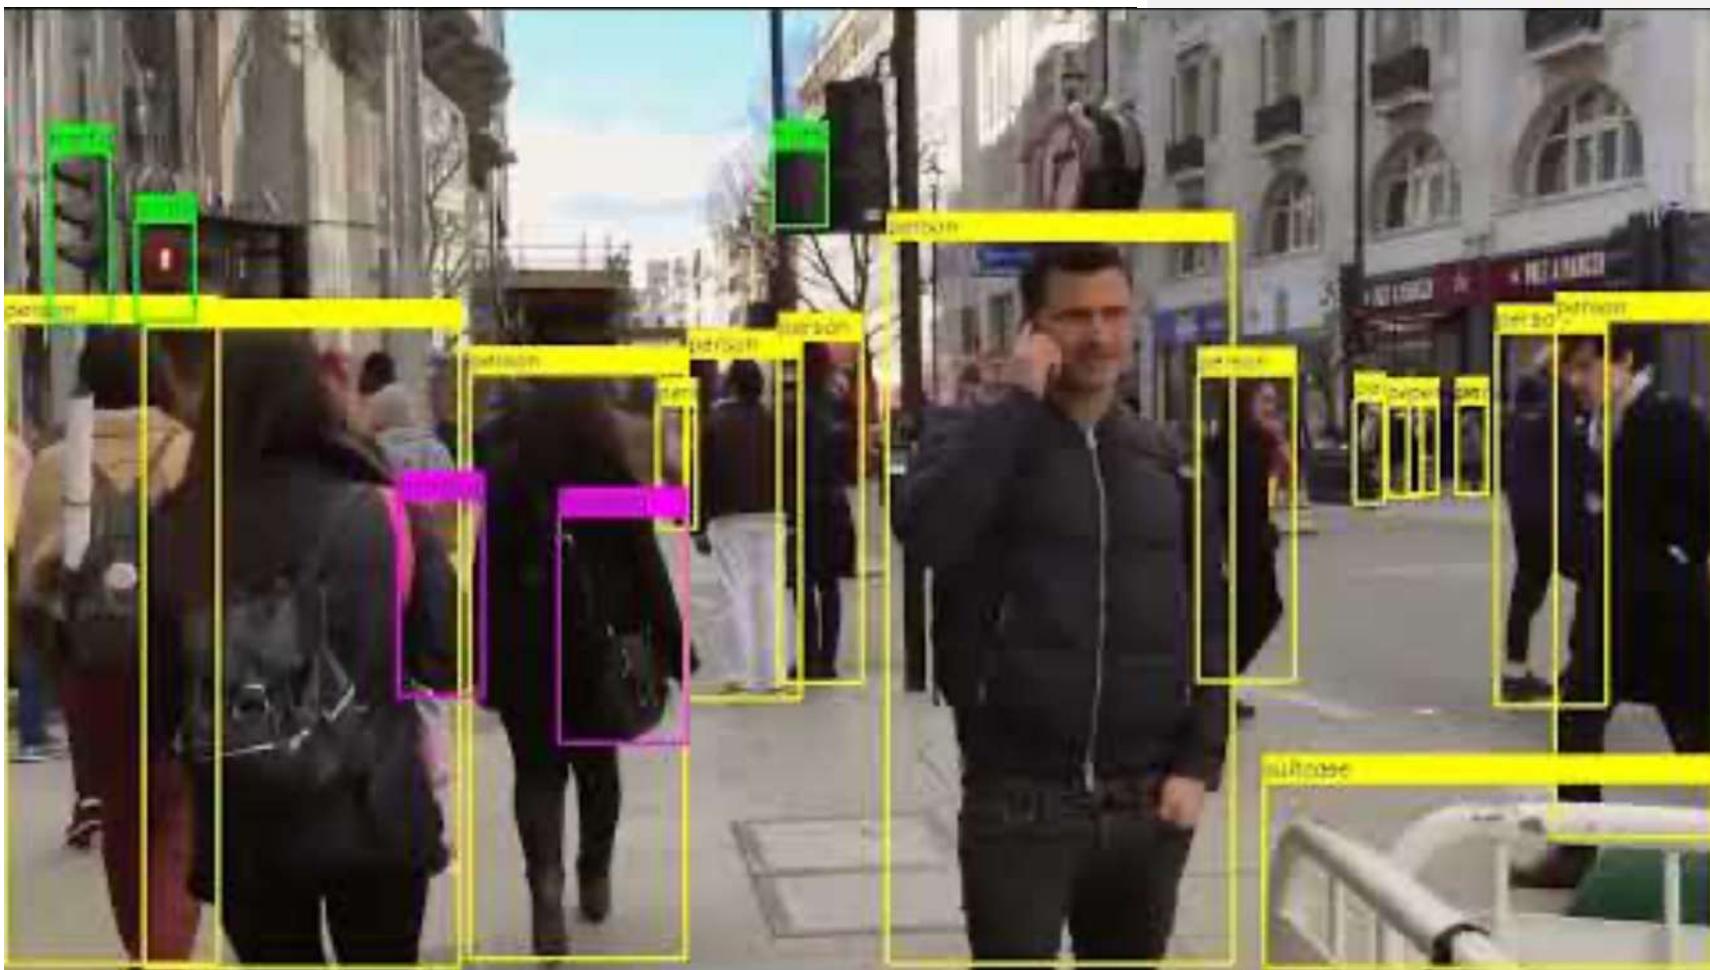
13,911 views Nov 13, 2017

No description has been added to this video.

64



YOLO v5



Example of YOLO v5 detection on video file

 Luiz doleron
127 subscribers

Subscribe

5,781 views Jan 23, 2022

Video example of YOLOv5 performing object detection.

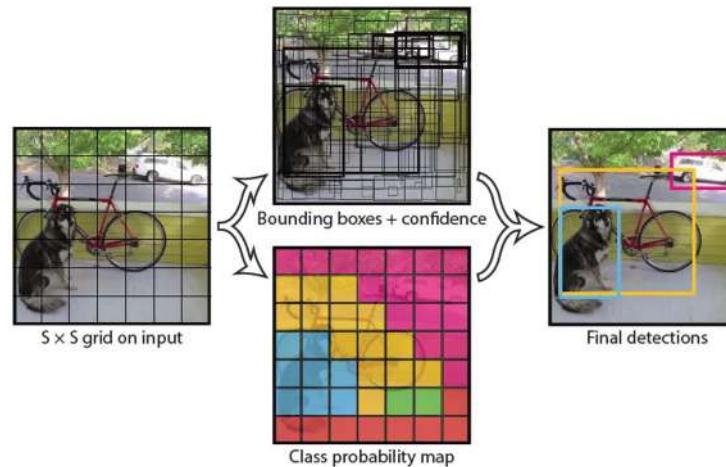
- Check out the explanation at [/ detecting-objects-with-yolov5-opencv-python...](#)
- Github repository is here: <https://github.com/doleron/yolov5-ope...>



What is YOLO?

You Only Look Once (YOLO) is a one-stage, end-to-end object detection framework that predicts bounding boxes and class probabilities in a single forward pass of a neural network.

Unlike earlier object detection approaches that adapted image classifiers into multi-step detection pipelines, YOLO treats object detection as a single regression problem.



<https://medium.com/data-science/yolov1-you-only-look-once-object-detection-e1f3ffec8a89>



Home People Dataset Tasks Evaluate

info@cocodataset.org

News

- We are pleased to announce the LVIS 2021 Challenge and Workshop to be held at ICCV.
- Please note that there will not be a COCO 2021 Challenge, instead, we encourage people to participate in the LVIS 2021 Challenge.
- We have partnered with the team behind the open-source tool **FiftyOne** to make it easier to download, visualize, and evaluate COCO
- **FiftyOne** is an open-source tool facilitating visualization and access to COCO data resources and serves as an evaluation tool for model analysis on COCO.

What is COCO?



COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- ✓ Object segmentation
- ✓ Recognition in context
- ✓ Superpixel stuff segmentation
- ✓ 330K images (>200K labeled)
- ✓ 1.5 million object instances
- ✓ 80 object categories
- ✓ 91 stuff categories
- ✓ 5 captions per image
- ✓ 250,000 people with keypoints

<https://cocodataset.org/#home>

Collaborators

Tsung-Yi Lin Google Brain
Genevieve Patterson MSR, Trash TV
Matteo R. Ronchi Caltech
Yin Cui Google
Michael Maire TTI-Chicago
Serge Belongie Cornell Tech
Lubomir Bourdev WaveOne, Inc.
Ross Girshick FAIR
James Hays Georgia Tech
Pietro Perona Caltech
Deva Ramanan CMU
Larry Zitnick FAIR
Piotr Dollár FAIR

Sponsors



ultralytics/cfg/datasets/coco8.yaml

Classes
names:

79 classes

0: person
1: bicycle
2: car
3: motorcycle
4: airplane
5: bus
6: train
7: truck
8: boat
9: traffic light
10: fire hydrant
11: stop sign
12: parking meter
13: bench
14: bird
15: cat
16: dog
17: horse
18: sheep
19: cow
20: elephant
21: bear
22: zebra
23: giraffe

<https://docs.ultralytics.com/datasets/detect/#ultralytics-yolo-format>



YOLO's output



1. Resize image.
2. Run convolutional network.
3. Non-max suppression.

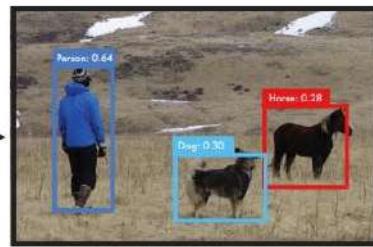
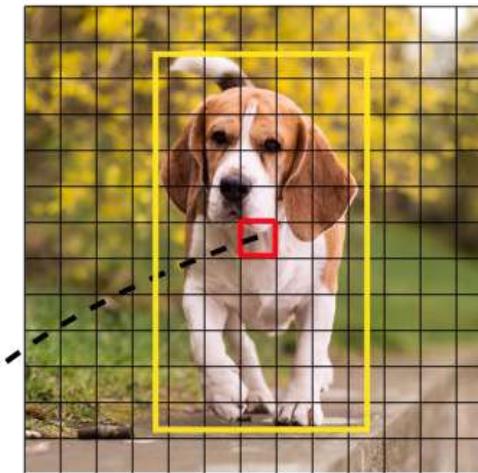


Image Grid. The Red Grid is responsible for detecting the dog



68

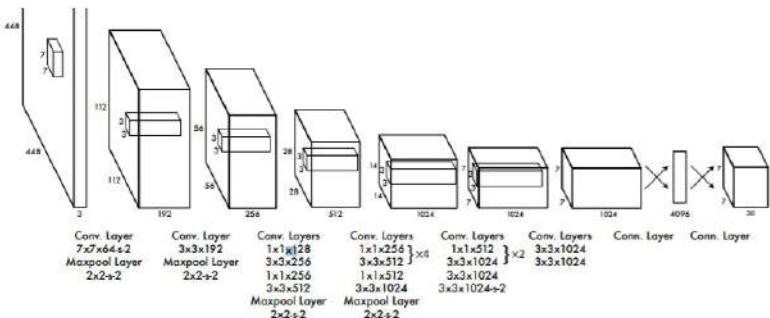
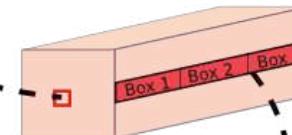
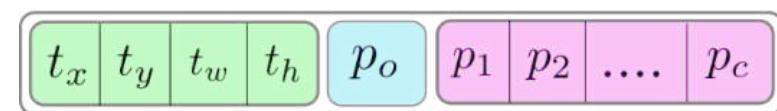


Figure 3: The Architecture. Our detection network has 24 convolutional layers followed by 2 fully connected layers. Alternating 1×1 convolutional layers reduce the features space from preceding layers. We pretrain the convolutional layers on the ImageNet classification task at half the resolution (224×224 input image) and then double the resolution for detection.

Prediction Feature Map



Attributes of a bounding box

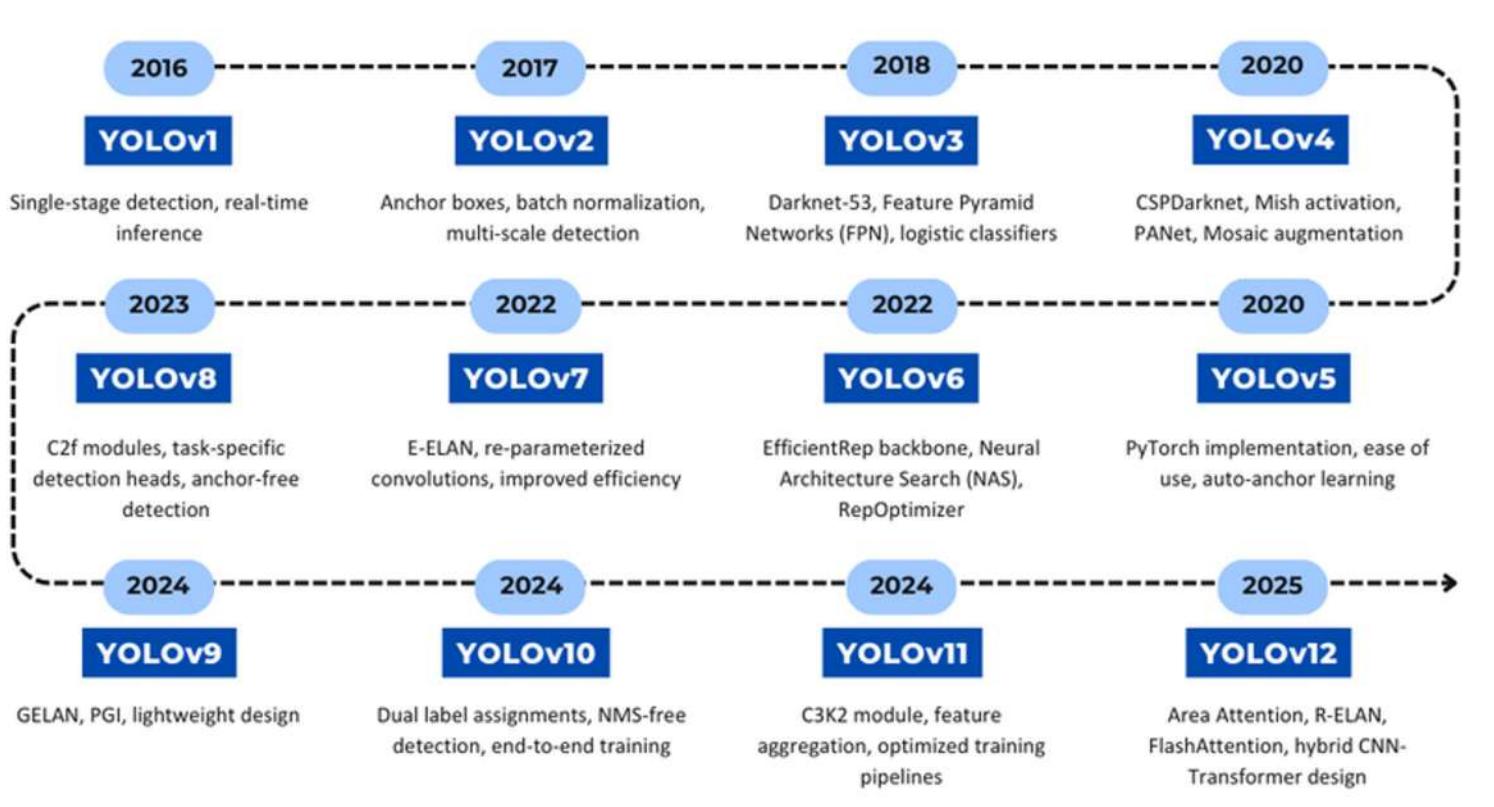


<https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/>

x B

<https://arxiv.org/abs/2504.11995>

The Evolution of YOLO

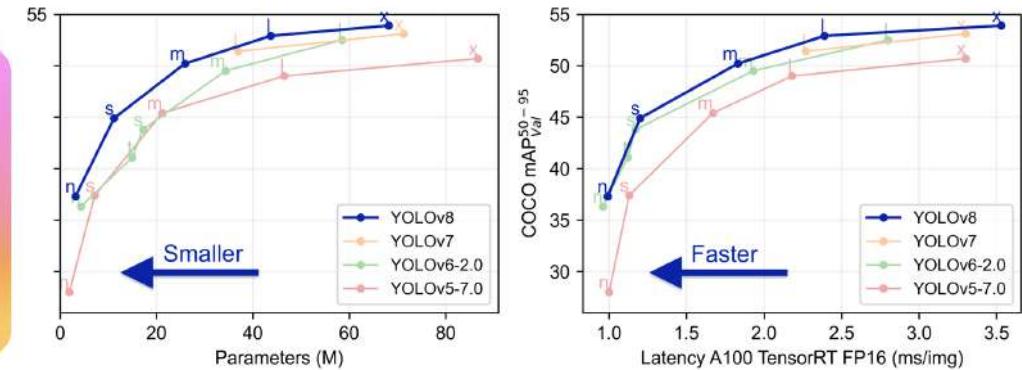


Ultralytics YOLO → production-ready, practical framework (YOLOv5 → YOLOv8 → YOLOv11).
Academic YOLO → new research architectures (YOLOv9, v10, v11, v12).

+ Ultralytics YOLO

YOLOv5 (2020) → YOLOv8 (2023) → Ultralytics YOLO 11 (2024)

70



Ultralytics YOLOv8 is a cutting-edge, state-of-the-art (SOTA) model that builds upon the success of previous YOLO versions and introduces new features and improvements to further boost performance and flexibility. YOLOv8 is designed to be fast, accurate, and easy to use, making it an excellent choice for a wide range of object detection and tracking, instance segmentation, image classification and pose estimation tasks.

We hope that the resources here will help you get the most out of YOLOv8. Please browse the YOLOv8 [Docs](#) for details, raise an issue on [GitHub](#) for support, and join our [Discord](#) community for questions and discussions!

To request an Enterprise License please complete the form at [Ultralytics Licensing](#).

Models

YOLOv8 Detect, Segment and Pose models pretrained on the COCO dataset are available here, as well as YOLOv8 Classify models pretrained on the ImageNet dataset. Track mode is available for all Detect, Segment and Pose models.



All Models download automatically from the latest Ultralytics [release](#) on first use.



YOLO Models (v9 – v12) Timeline

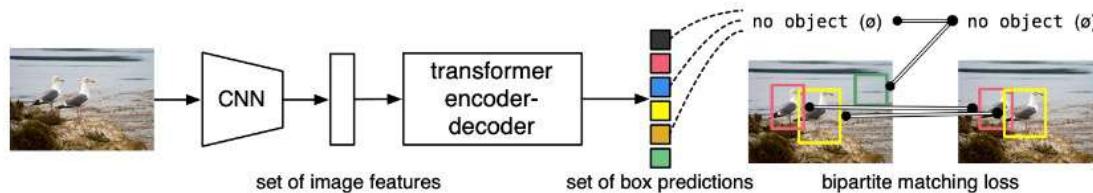
Model	Release Timeframe	Team / Organization	Key Highlight
YOLOv12	Feb 2025	Google Research + Tsinghua University (collaboration)	Attention-centric architecture with high accuracy and real-time speed
YOLOv11	Late 2024	Tencent Youtu Lab	Enhanced architecture (C3k2, SPPF, C2PSA), versatile across detection, segmentation, pose
YOLOv10	Mid 2024	Tsinghua University + Huawei Noah's Ark Lab	End-to-end design without NMS, major speed and efficiency gains
YOLOv9	Early 2024	Wong Kin Yiu (original YOLO contributor) + Community researchers	PGI and GELAN for better gradients and efficiency
Ultralytics YOLO11	2024–2025	Ultralytics	Practical SOTA release, supports detection, segmentation, pose, classification, track

+ DETR (Detection Transformer) [Meta, 2020]

DETR: End-to-End Object Detection with Transformers

Support Ukraine

PyTorch training code and pretrained models for DETR (DEtection TTransformer). We replace the full complex hand-crafted object detection pipeline with a Transformer, and match Faster R-CNN with a ResNet-50, obtaining 42 AP on COCO using half the computation power (FLOPs) and the same number of parameters. Inference in 50 lines of PyTorch.



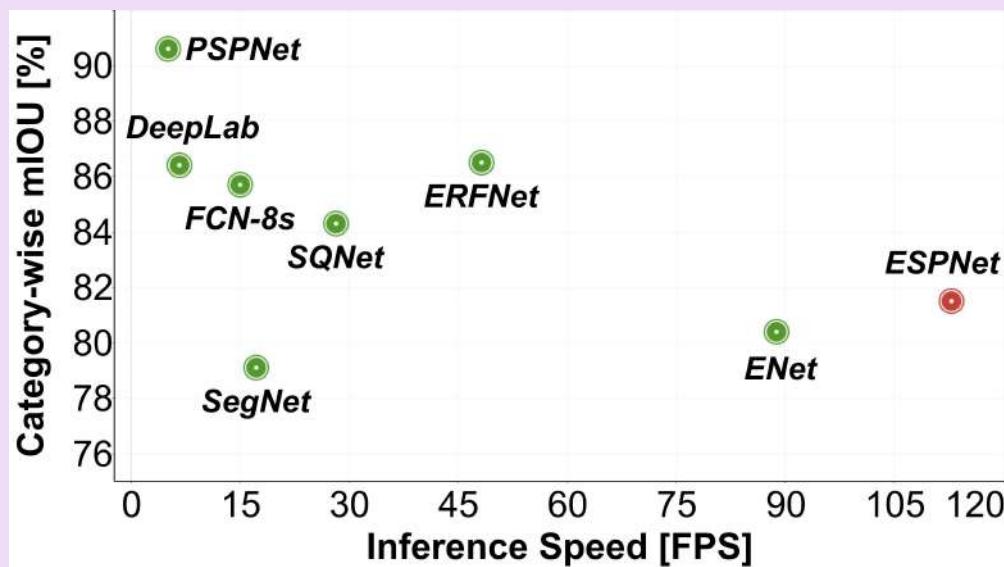
- 1. Deformable DETR (2020)
■ Team: Microsoft Research Asia
- 2. Conditional DETR (2021)
■ Team: Chinese Academy of Sciences
- 3. Anchor DETR (2021)
■ Team: SenseTime and Shanghai Jiao Tong University
- 4. Efficient DETR (2021)
■ Team: UC Berkeley and Google Research
- 5. DINO (DETR with Improved Non-Autoregressive Decoding) (2022)
■ Team: Meta AI (formerly Facebook AI Research)
- 6. RT-DETR (Real-Time DETR) (2024)
■ Team: Alibaba DAMO Academy



Top Object Detection Models (2025)

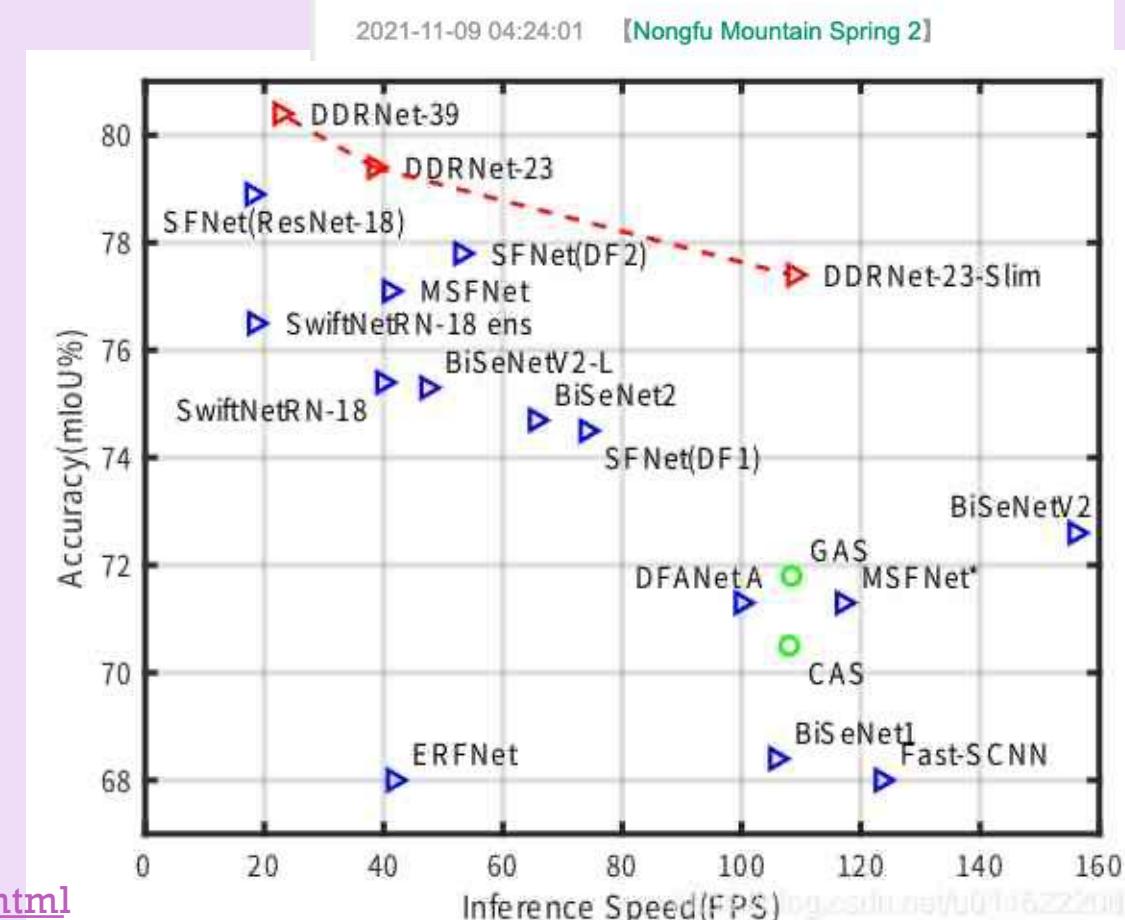
Model	Team / Source	Notable Metric (COCO)	Strengths
RF-DETR	Roboflow	73.6 AP ₅₀ / 54.7 AP_{50:95} (Medium)	Real-time DETR; best mAP at low latency
YOLOv12	YOLO authors (YOLO Series)	—	Latest YOLO, attention-enhanced
YOLO11 (Ultralytics)	Ultralytics	Up to ~50 AP ₅₀ / ~48.6 AP_{50:95}	Multi-task, production-ready

SOTA of Semantic Segmentation



- **Real-time semantic segmentation**
- BiSeNet V1, V2
- DDRNet

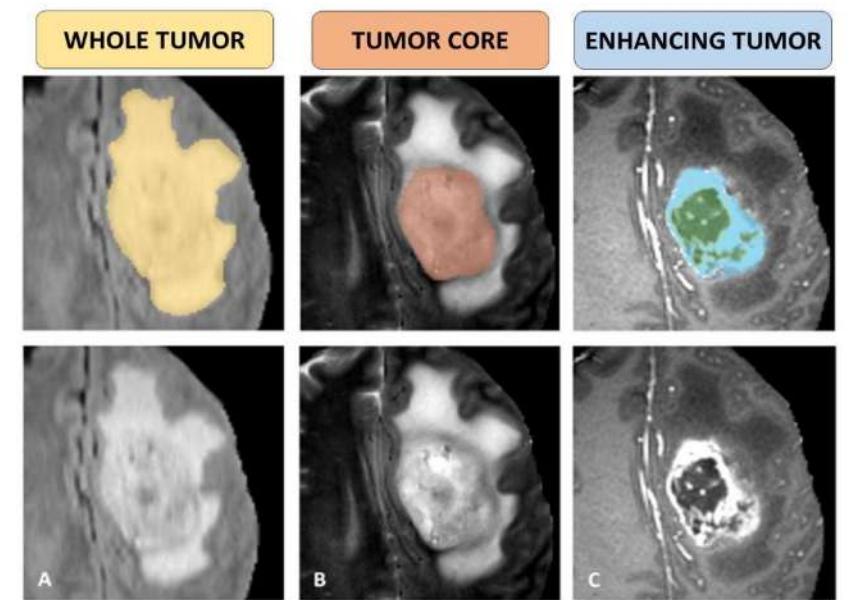
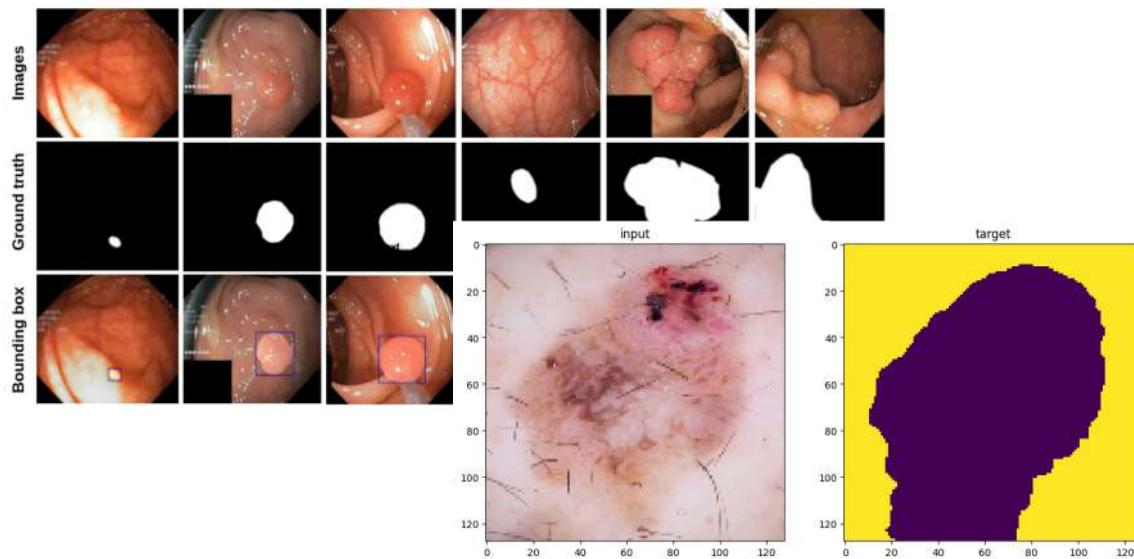
<https://chowdera.com/2021/11/20211109042359120o.html>





Semantic Segmentation

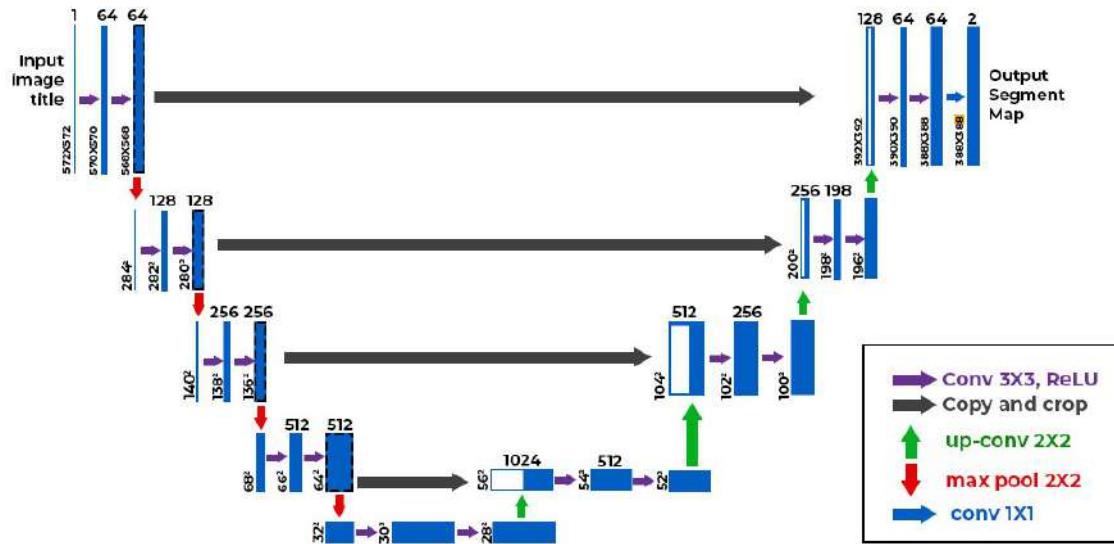
- Semantic segmentation = pixel-wise classification
- Each pixel gets a class label (road, car, polyp, tumor, sky, background).



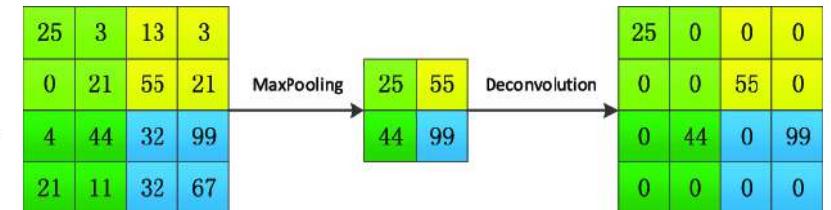


UNet: Encoder-Decoder Network: Encoder, Decoder, Skip Connections

■ U-Net



Deconvolutional layer



$$\mathcal{L} = \lambda_1 \cdot \text{CE} + \lambda_2 \cdot \text{Dice} ; \text{CE (pixel loss) \& Dice (region loss; semantic)}$$

Architecture of the U-net for a given input image. The blue boxes correspond to feature maps blocks with their denoted shapes. The white boxes correspond to the copied and cropped feature maps.

Source: [O. Ronneberger et al. \(2015\)](#)

+ Recent interesting segmentation networks

- **1. ContextFormer (2025)**
- Team: Abid, Mehta, Wu, Dharejo, Timofte
- Architecture: Hybrid CNN + Vision Transformer
- Strengths: Combines high accuracy with real-time efficiency
- Applications: Urban scene understanding, robotics, autonomous driving

- **2. Segment Anything Model (SAM) [Meta, 2023] → SAM V2 [Meta, 2024]**
- Explanation: SAM is a highly versatile model designed to segment any object within an image with minimal user input. It uses a prompt-based approach, where the user provides a prompt (e.g., a point or bounding box), and the model generates a mask for the object. SAM is particularly noted for its generalizability across different segmentation tasks.

- **3. Mask2Former [Meta, 2022]**
- Explanation: Mask2Former is a unified model for instance, semantic, and panoptic segmentation tasks. It improves upon the original MaskFormer by incorporating transformers into the architecture, enabling it to handle different segmentation tasks within a single framework efficiently.

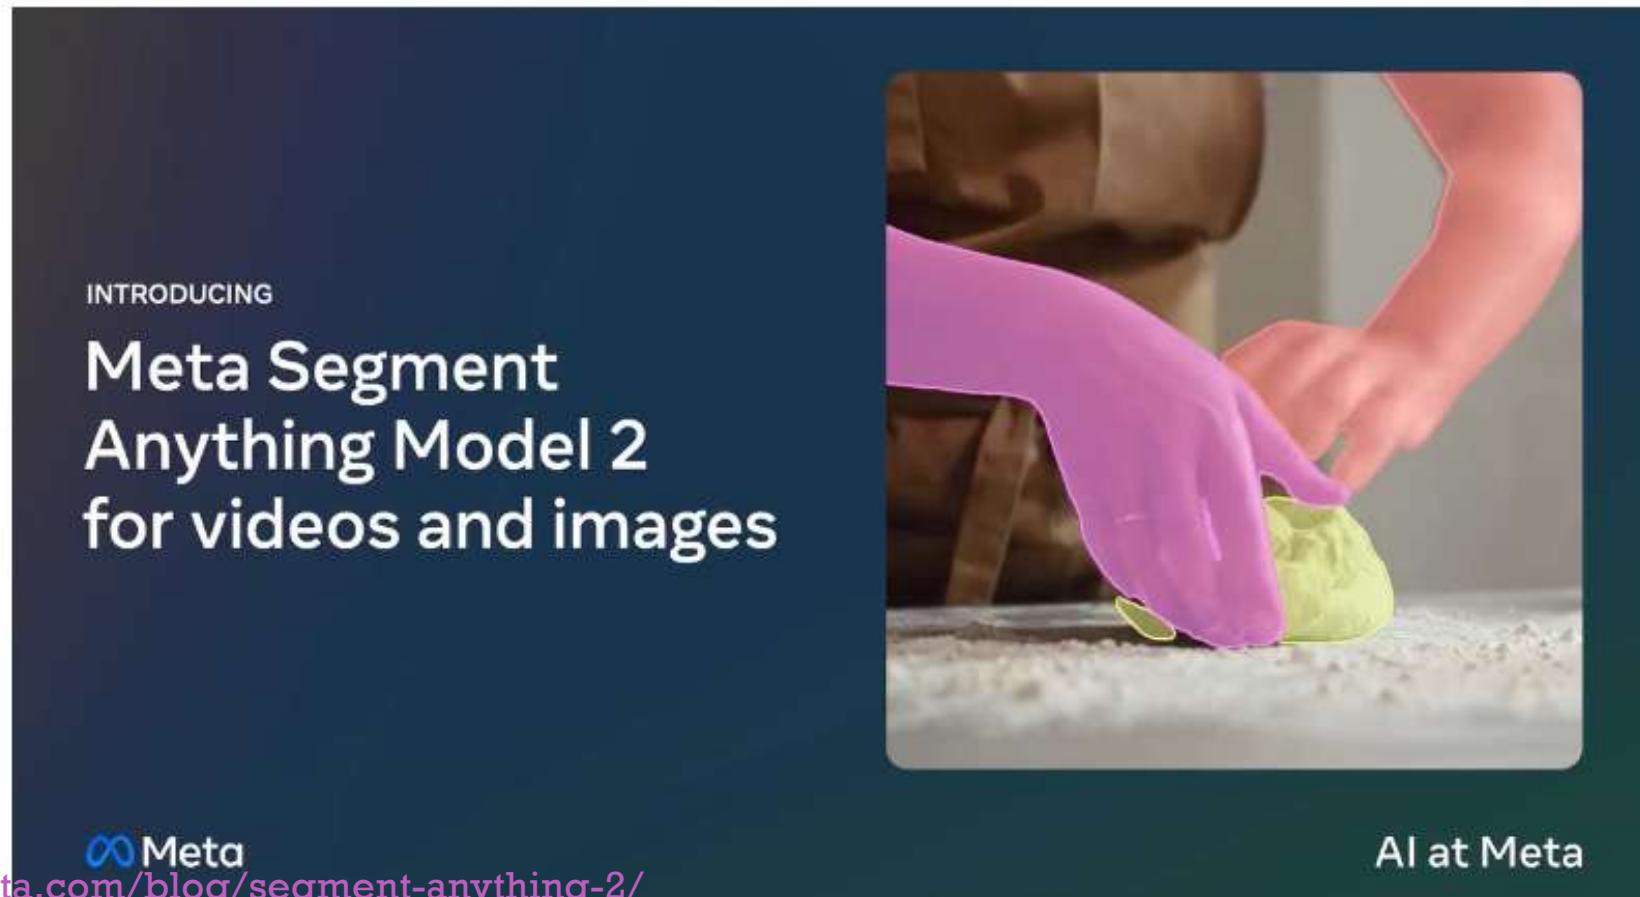
- **4. Swin Transformer V2 [Microsoft, 2022]**
- Explanation: Swin Transformer V2 builds on the hierarchical vision transformer design introduced in the original Swin Transformer. It enhances performance in both segmentation and classification tasks and has achieved SOTA results on benchmarks such as COCO and ADE20K.

- **5. SegFormer [NVIDIA, 2021]**
- Explanation: SegFormer is a transformer-based model that is designed to be both efficient and accurate for semantic segmentation tasks. It integrates transformers with lightweight MLP decoders, enabling high performance while maintaining computational efficiency.

- **6. DeepLabV3+ [Google, 2018]**
- Explanation: DeepLabV3+ is an extension of the DeepLab series, which combines spatial pyramid pooling and atrous convolutions. The "+" version introduces a decoder module to improve object boundary delineation, making it a strong performer in semantic segmentation.

Introducing SAM 2: The next generation of Meta Segment Anything Model for videos and images

July 29, 2024 • ① 15 minute read



INTRODUCING

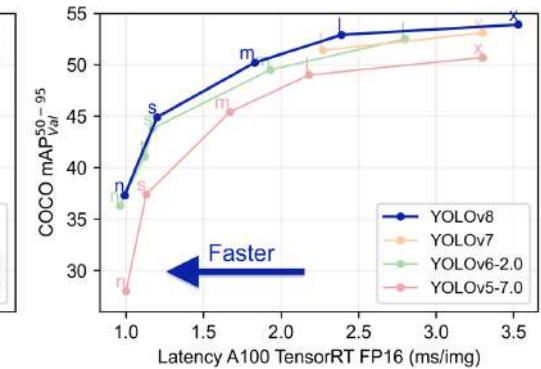
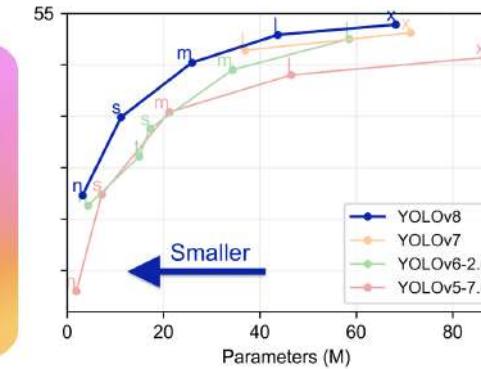
Meta Segment Anything Model 2 for videos and images

AI at Meta

<https://ai.meta.com/blog/segment-anything-2/>

+ Ultralytics's YOLO (extra segmentation model)

YOLOv5 (2020) → YOLOv8 (2023) → Ultralytics YOLO 11 (2024)



English | 简体中文

Ultralytics CI passing | codecov 87% | DOI 10.5281/zenodo.7347926 | docker pulls 22k

[Run on Gradient](#) | [Open in Colab](#) | [Open in Kaggle](#)

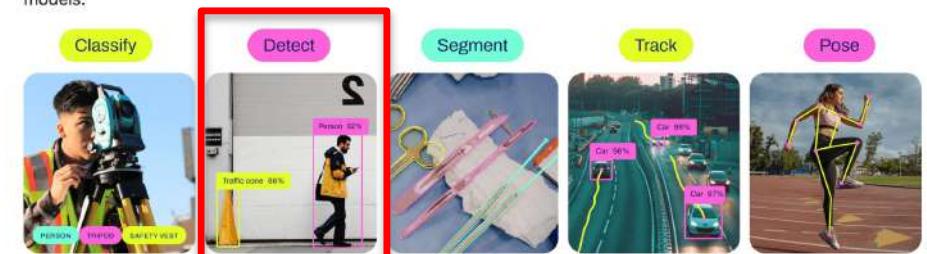
Ultralytics YOLOv8 is a cutting-edge, state-of-the-art (SOTA) model that builds upon the success of previous YOLO versions and introduces new features and improvements to further boost performance and flexibility. YOLOv8 is designed to be fast, accurate, and easy to use, making it an excellent choice for a wide range of object detection and tracking, instance segmentation, image classification and pose estimation tasks.

We hope that the resources here will help you get the most out of YOLOv8. Please browse the YOLOv8 [Docs](#) for details, raise an issue on [GitHub](#) for support, and join our [Discord](#) community for questions and discussions!

To request an Enterprise License please complete the form at [Ultralytics Licensing](#).

Models

YOLOv8 Detect, Segment and Pose models pretrained on the COCO dataset are available here, as well as YOLOv8 Classify models pretrained on the ImageNet dataset. Track mode is available for all Detect, Segment and Pose models.



All Models download automatically from the latest Ultralytics [release](#) on first use.



YOLOv8-Seg: Object Detection → Segment

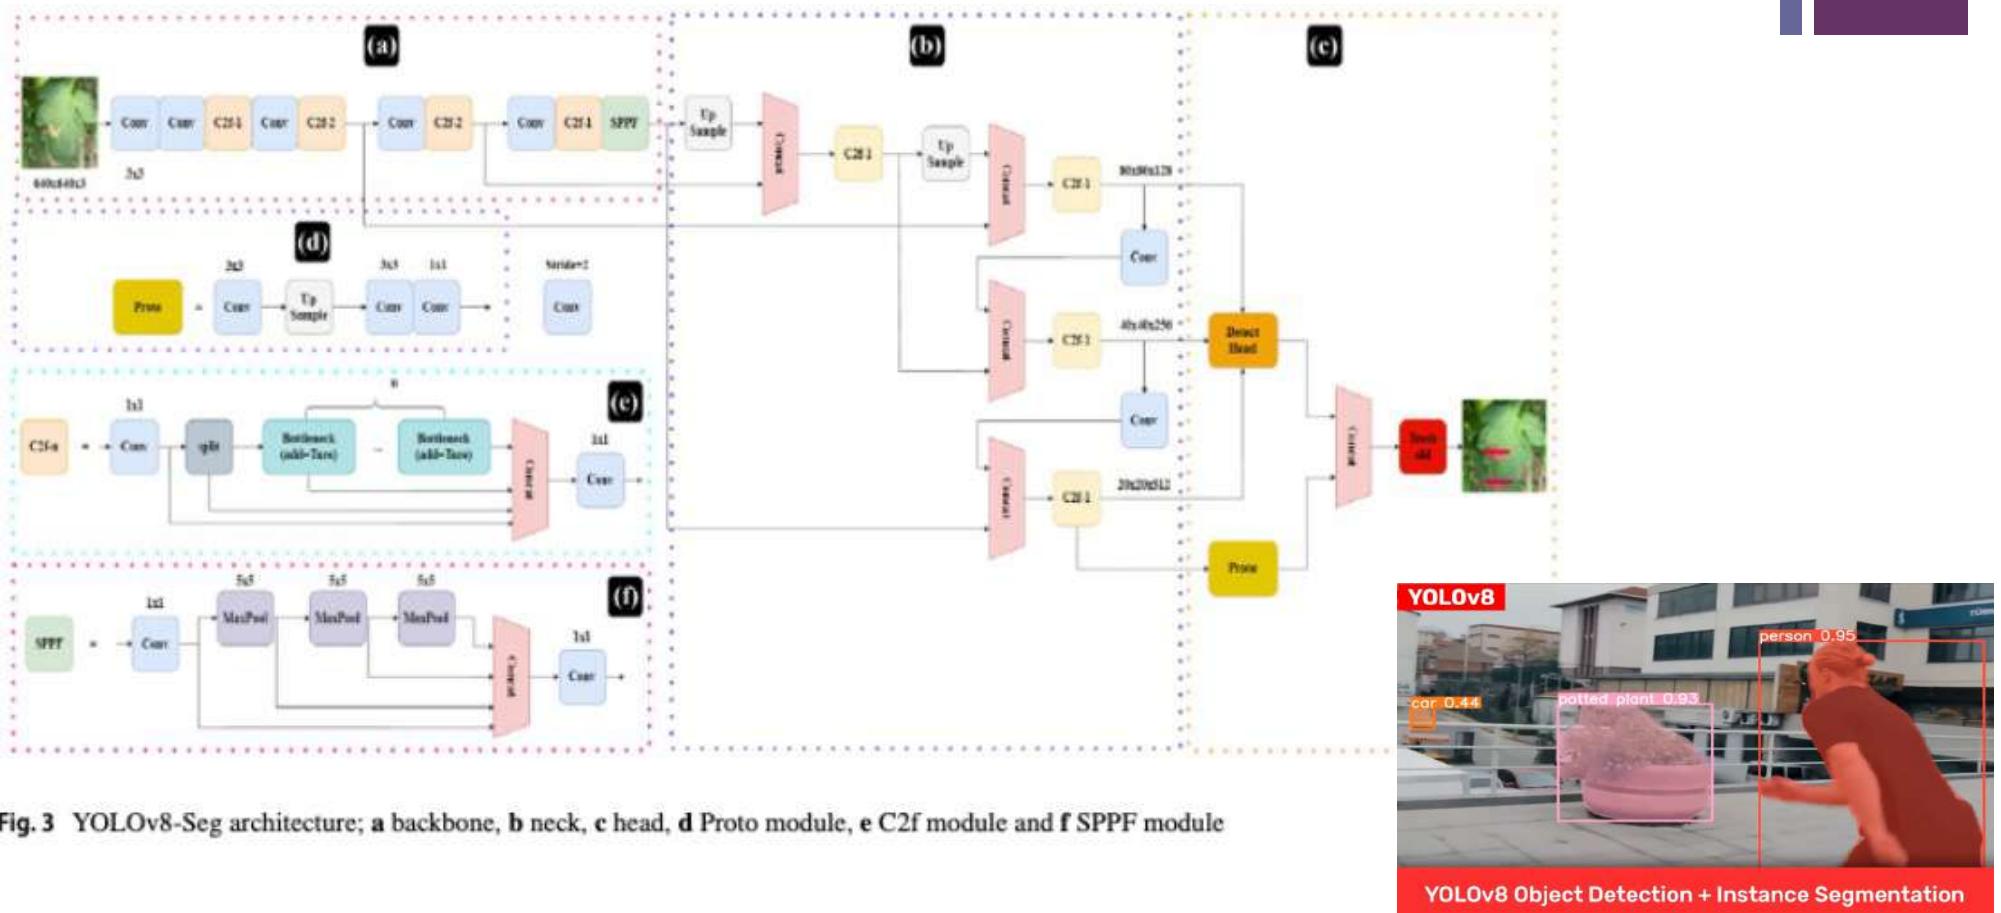
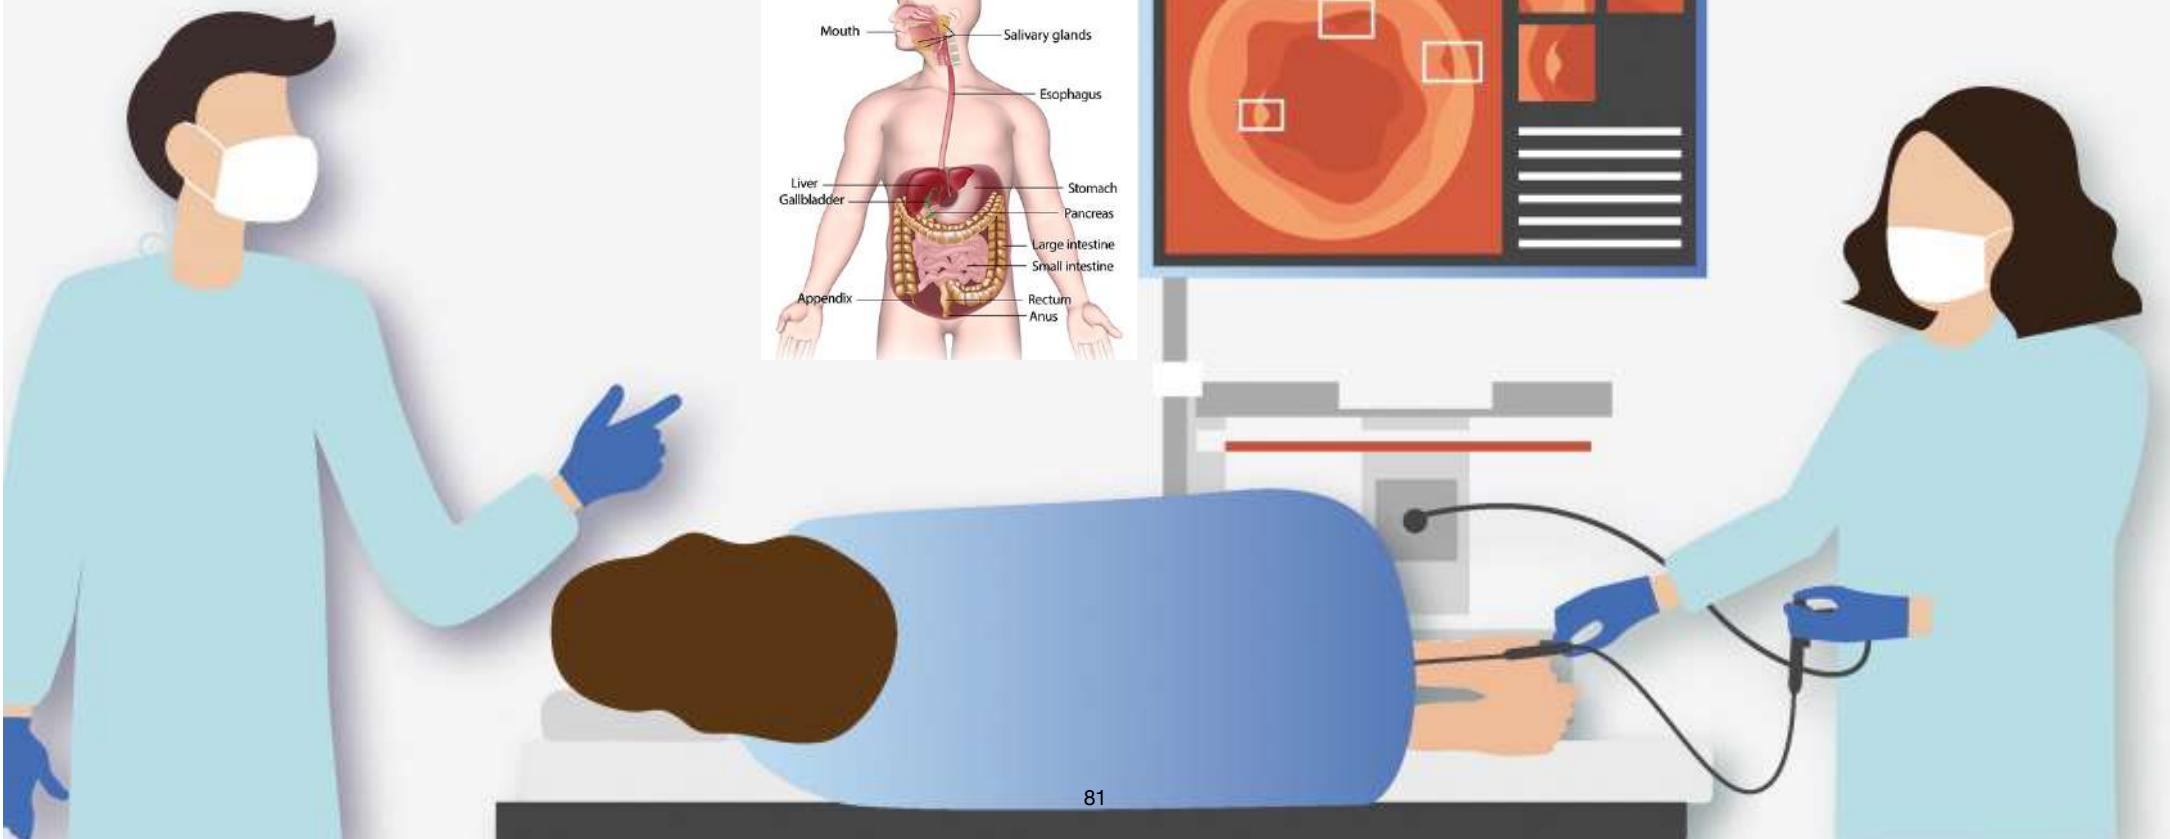


Fig. 3 YOLOv8-Seg architecture; **a** backbone, **b** neck, **c** head, **d** Proto module, **e** C2f module and **f** SPPF module

[Link](#)

Real-Time Colonic Polyp Detection

An AI-assisted solution that aims to improve colonic polyp detection in real-time. It is compatible with all standard colonoscope systems.

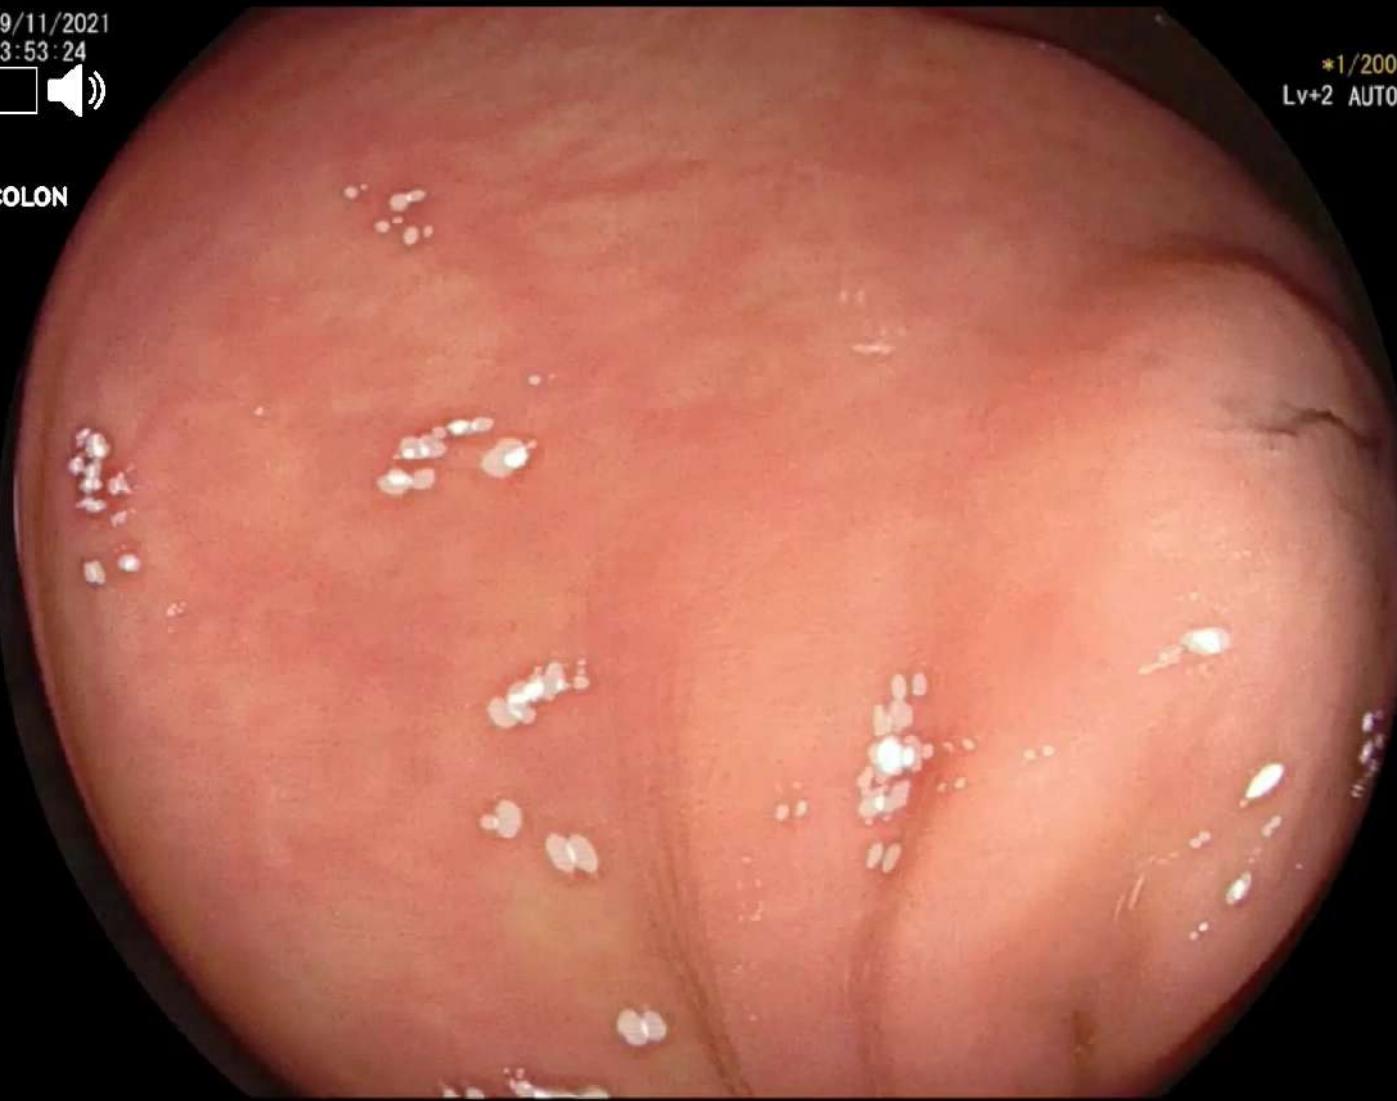


09/11/2021

13:53:24

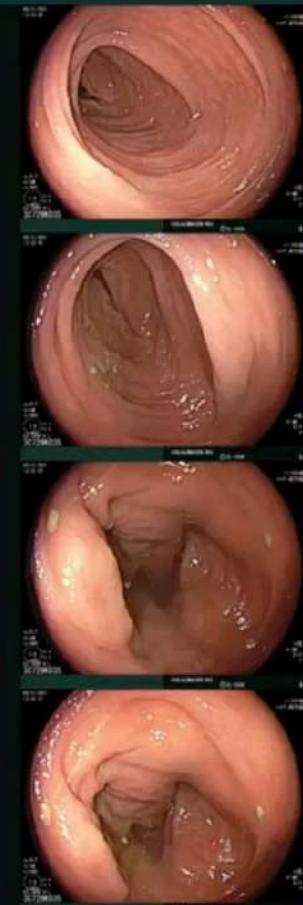


COLON



*1/200
Lv+2 AUTO

HT NR
SE
3.8 12.0 s1: F+T
12.0 s2: LM
s3: CAD
s4:
EC-760R-V/L
3C729K035
BL-7000
CHULALONGKORN HOS



NBI

Name:

Sex: Age:

D.O.B.:

18/05/2021

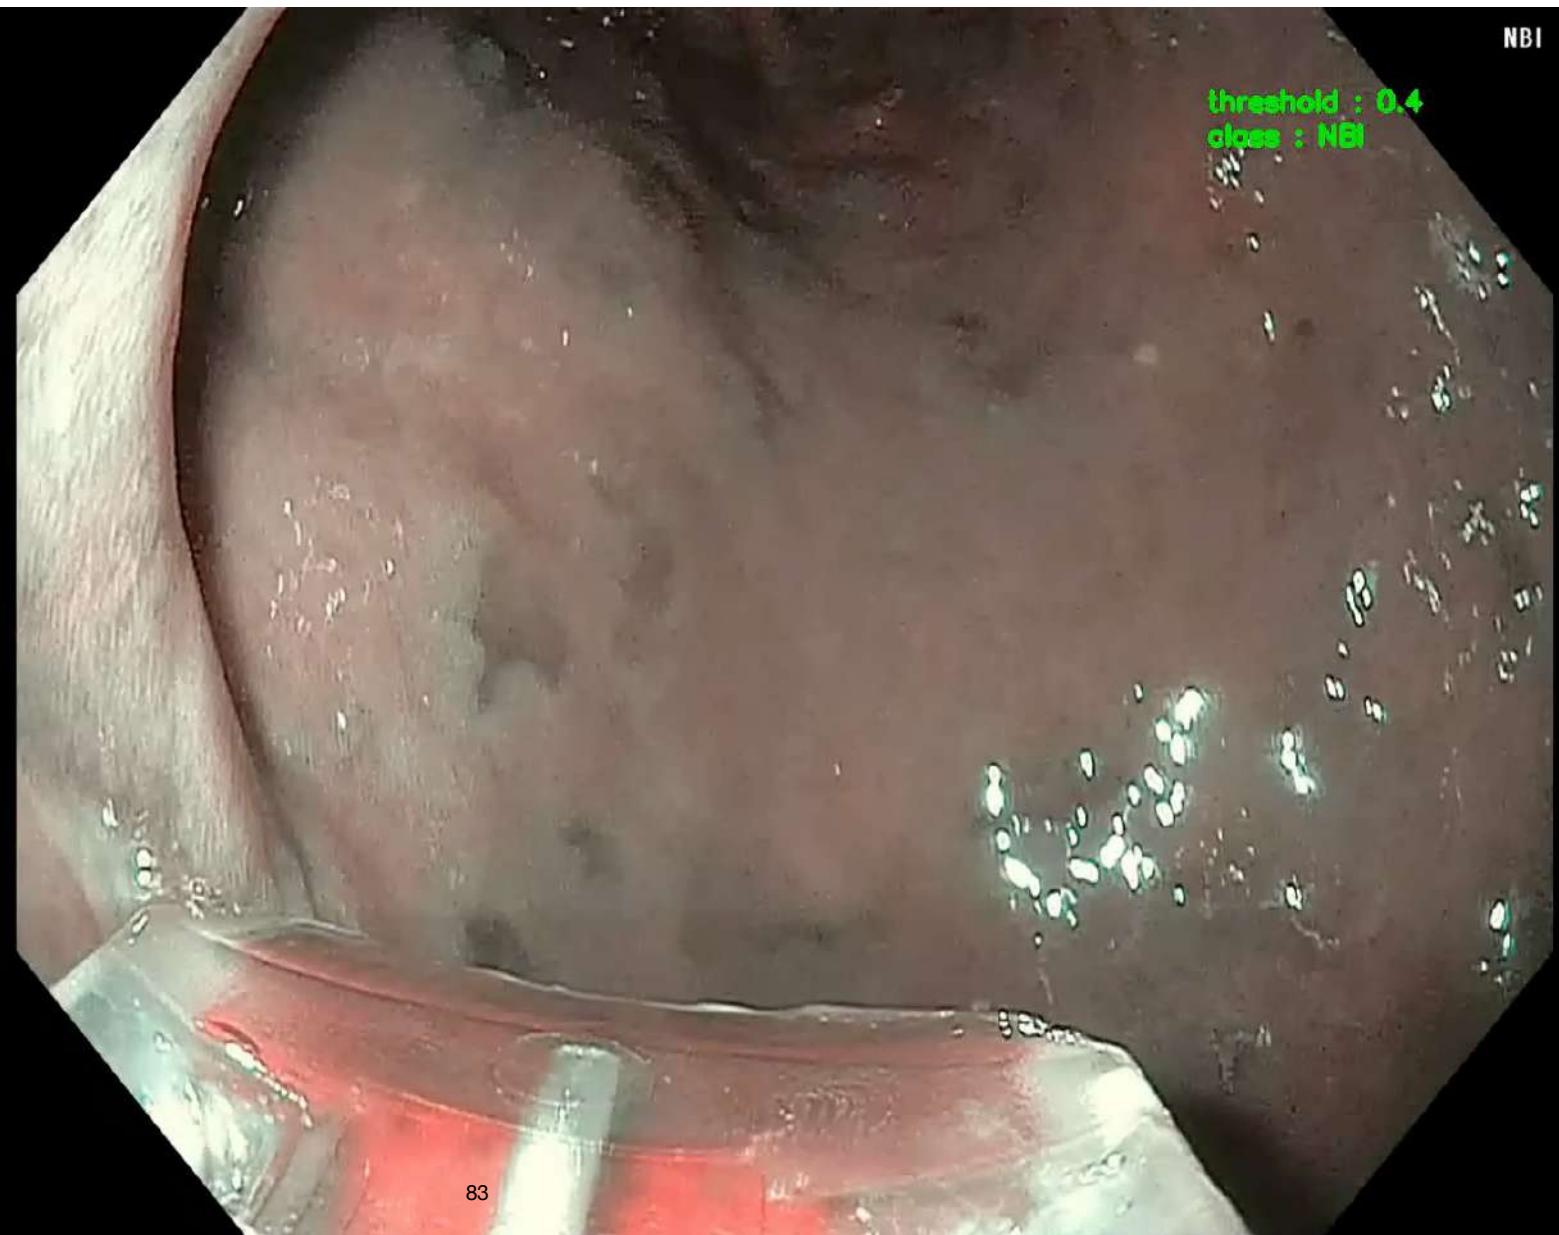
08:29:36

threshold : 0.4
close : NBI

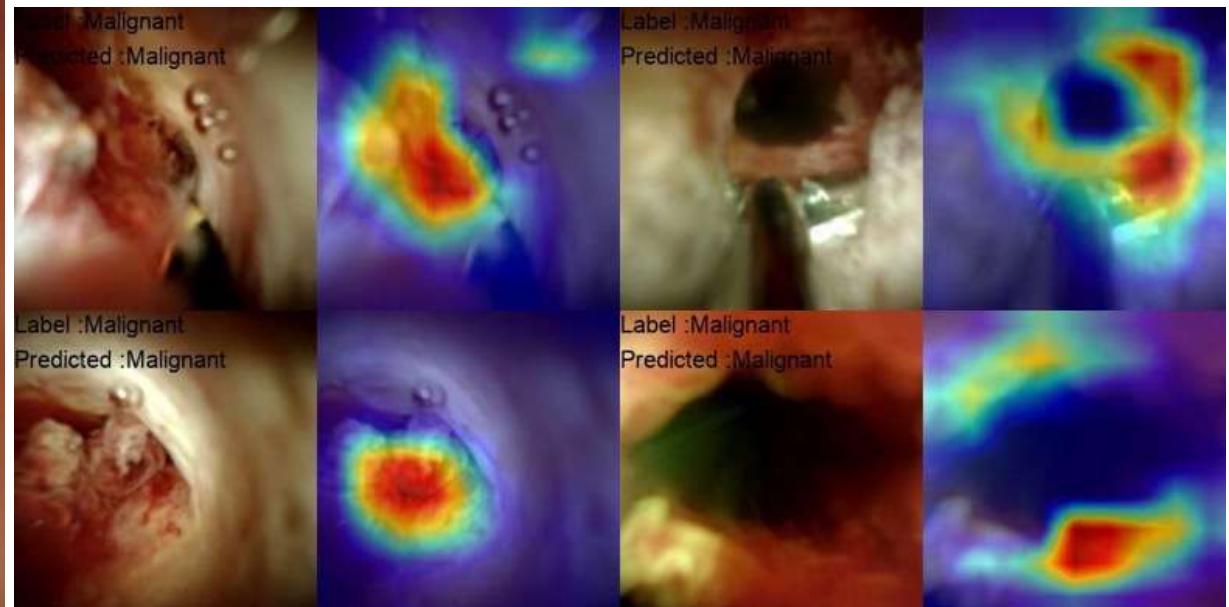
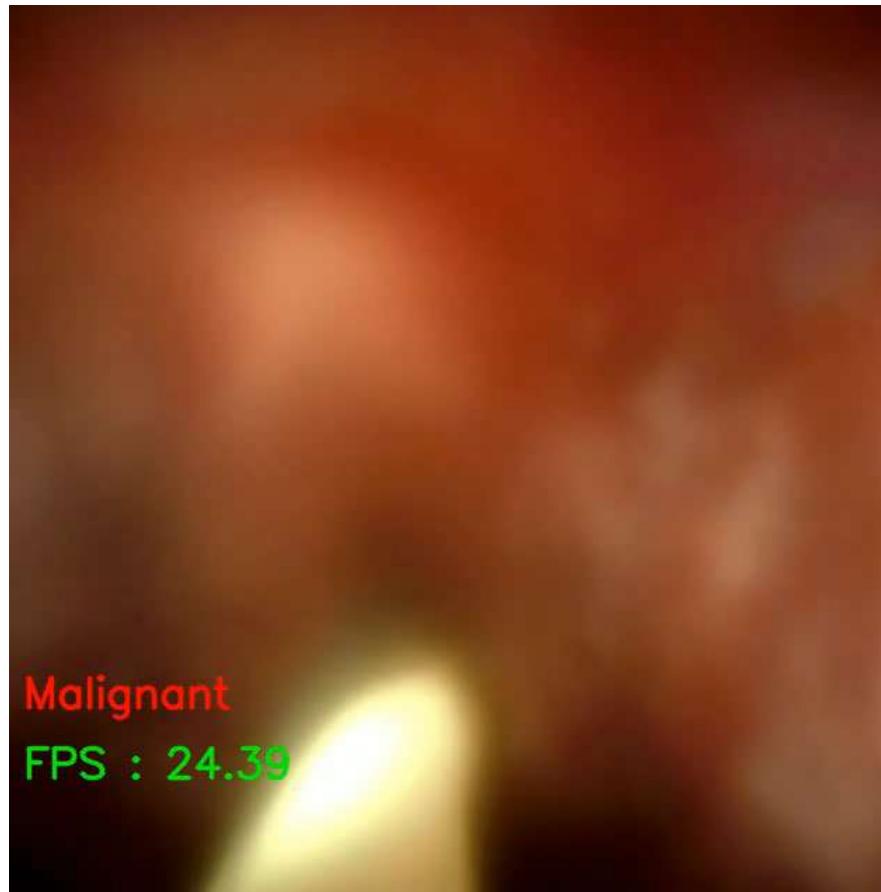
■■□/---(0/1)

Eh:B8 Cm:1

Comment:



Cholangioscopy



+

Appendix